

Machine Learning

Liyao Tang

June 6, 2019

Contents

1	Math	1
1.1	Convolution	1
1.2	Linear Algebra	1
1.2.1	Essence	1
1.2.2	Interchanging Coordinates	12
1.3	Calculus	14
1.3.1	Integral	14
1.4	Probability Theory	16
1.4.1	Introduction	16
1.4.2	Expectations and Covariances	18
1.4.3	Transformations of Random Variables	19
1.4.4	Gaussian Distribution	20
1.4.5	Bayesian Interpretation of Probability	22
2	Introduction	26
2.1	General Concern	26
2.1.1	Types of Learning	26
2.2	Decision Theory	27
2.2.1	Overview	27
2.3	Information Theory	33
2.3.1	33
2.4	Recommended Practice	33
2.4.1	Data	33
2.4.2	Dataset	35
2.4.3	Orthogonalization Practice	37
2.4.4	Tunning Hyperparameters	38
2.5	Model Analysis	40
2.5.1	Measurements of Problem	40
2.5.2	Improving Model	43
2.5.3	Evaluating Hypothesis	46
2.5.4	Skewed classes	46
2.6	Supervised Learning	46
2.7	Linear Regression	47
2.8	Bayesian Regression	48
2.9	Logistic Regression (Classification)	50
2.10	Latent Variable Analysis	54
2.10.1	Principal Component Analysis (PCA)	54
2.10.2	Independent Component Analysis (ICA)	57
2.10.3	t-SNE	57
2.10.4	Anomaly Detection	57
2.10.5	Recommender System	58
2.11	Large Scale Machine Learning	59

2.11.1 Online Learning	59
2.11.2 Map-reduce	60
3 Linear Regression	61
4 Linear Classification	62
5 Kernel Methods	63
6 Graphical Models	64
7 Mixture Models and EM	65
8 Approximate Inference	66
9 Sampling Methods	67
10 Continuous Latent Variable	68
11 Sequential Data	69
11.1 Markov Model	69
11.1.1 Markov Chain	69
11.1.2 Hidden Markov Model	70
11.2 Linear Dynamic System	76
12 Deep Learning	78
12.1 Interview of Fame	78
12.1.1 Geoffrey Hinton	78
12.1.2 Pieter Abbeel	79
12.1.3 Ian Goodfellow	80
12.1.4 Yoshua Bengio	80
12.1.5 Yuanqing Lin	81
12.1.6 Andrej Karpathy	81
12.1.7 Ruslan Salakhutdinov	82
12.1.8 Research	82
12.2 Basic Neural Network	88
12.2.1 Advantages	88
12.2.2 Problem	89
12.2.3 Learning	91
12.3 Operations & Layers Structure	91
12.3.1 Operations in Network	91
12.3.2 Operations on Network	93
12.3.3 Cost	98
12.3.4 Layers	99
12.3.5 Blocks	106
12.4 Architectures	107
12.4.1 Convolutional Networks	107
12.4.2 Recurrent Neural Network	109
12.4.3 Encoder-Decoder Architecture	110
12.4.4 Generative Network	111
12.5 Computer Vision	111
12.5.1 Objects Detection	111
12.5.2 Object Tracking	117
12.5.3 Face Recognition	128
12.5.4 Recognition at a Distance	130

12.5.5	Image Style Transfer	133
12.5.6	Point Cloud Data Processing	134
12.6	Natural Language Processing	137
12.6.1	Language Representation	137
12.6.2	Language Modeling	140
12.6.3	Name-Entity Recognition	141
12.6.4	Sentiment Classification	141
12.6.5	Neural Machine Translation	142
12.6.6	Speech Recognition	143
12.6.7	Machine Reading Comprehension	144
12.6.8	Image Caption	144
12.6.9	Referring Segmentation	145
12.7	Special Learning	147
12.7.1	Transfer Learning	147
12.7.2	Multi-task Learning	149
12.7.3	K-shot Learning	150
12.8	Generative Adversarial Nets	150
12.8.1	Temporal Generative Adversarial Nets	150

List of Figures

11.1	$P(D, B) = \sum_A P(A, B)P(D A) \quad P(F, B) = \sum_A \left(P(A, B) \sum_C P(C B)P(F C) \right)$ $P(D, F, B) = \sum_A \left(P(A, B)P(D A) \sum_C P(C B)P(F C) \right) \Rightarrow P(D B)P(F B) =$ $P(D, F B) \Rightarrow D \perp\!\!\!\perp F B \dots\dots\dots$	73
12.2	(using LSTM cell) \dots\dots\dots	105

List of Tables

Chapter 1

Math

1.1 Convolution

- Definition

- $f * g(z) = \int_{\mathbb{R}} f(x)g(z-x)dx$, where $f(x), g(x)$ are functions in \mathbb{R}

- Statistical Meaning

- Notation

- X, Y : independent random variables, with pdf's given by f and g
 - $Z = X + Y$, with pdf given by $h(z)$:

- $\Rightarrow h(z) = f * g(z)$

- derivation

$$\begin{aligned} H(z) &= P(Z < z) = P(X + Y < z) \\ &= \int_x P(X = x)P(X + Y < z | X = x)dx \\ &= \int_x f(x)P(Y < z - x)dx \\ &= \int_x f(x)G_Y(z - x)dx \\ \Rightarrow h(x) &= \frac{d}{dz}H(z) = \frac{d}{dz} \int_x f(x)G_Y(z - x)dx \\ &= \int_x f(x) \frac{dG_Y(z - x)}{dz} dx \\ &= \int_x f(x)g(z - x)dx \\ &= f * g(z) \end{aligned}$$

1.2 Linear Algebra

1.2.1 Essence

Vector

- Interpretation
 - Movement

- direction
 - distance
 - Numeric in High Dimensions
 - in 1- D : +/- represents direction
 - in n - D : +/- alone each dimension combined to represent an overall direction (direction of the n - D numeric - vector)
 - Abstraction
 - a vector space with operations that can be applied on different math concept (e.g. function derivatives is actually in vector space)
 - Numerics Multiplication
 - Scaling
 - the number scales the distance of vector (direction remains)
 \Rightarrow such number thus also called scalar
 - \Rightarrow scale alone each axis by that scalar
 $2\mathbf{x} = 2x_1e_1 + \dots + 2x_ne_n$,
 where e_1, \dots, e_n are vector defining coordinates
 - Linear Combination
 - Vector Adding: Generalization of Numerical Adding
 - in 1- D : joint movement along single axis
 - in n - D : joint movement along each axis \Rightarrow a joint movement in n - D space
 - Definition: $\mathbf{x} = a_1\mathbf{x}_1 + \dots + a_n\mathbf{x}_n$
 - the vectors $\mathbf{x}_1, \dots, \mathbf{x}_n$ only altered linearly (as only being scaled)
 \Rightarrow \mathbf{x} direction & size are linear combination of that in $\mathbf{x}_1, \dots, \mathbf{x}_n$
 - Span of $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$
 - the n - D space S constructed by linear combination of $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$
 - x_0 linearly dependent on $\{x_1, \dots, x_n\}$
 - x_0 can be constructed by linear combination of $\{x_1, \dots, x_n\}$
 (already in the span space S)
 - \Leftrightarrow function $a_0x_0 + \dots + a_nx_n = 0$ has other solution than $a_0 = \dots = a_n = 0$
 - x_0 linearly INdependent with $\{x_1, \dots, x_n\}$
 - x_0 can NOT be constructed by linear combination of $\{x_1, \dots, x_n\}$
 (not in the span space S , will increase the dimension of S if adopted)
 - \Leftrightarrow function $a_0x_0 + \dots + a_nx_n = 0$ has and only has solution $a_0 = \dots = a_n = 0$
 - Special Vectors
 - n - D Zero Vector \mathbf{x}
 -
 - Unit Vector
 - Basis of Vector Space S^n
 - general basis: a set of linearly independent vectors that span the space
 (i.e. a set of linearly independent n - D vectors)
 - unit basis: a general basis where every vector is unit vector
 - orthogonal basis: a general basis where all vectors are orthogonal with each other
 - unit orthogonal basis: a basis that is also a unit basis and an orthogonal basis
- \Rightarrow coordinate: the scalar to composite a vector given a specific basis

Linear Transformation and Maps

- Linear Transformation
 - Transformation
 - a function mapping: vector \rightarrow vector
 - a vector movement: scale & rotate all possible input vectors (i.e. a vector space)
 - Transformation with Linearity $L(\cdot)$
 - intuition: lines remain lines & origin remains origin
 - definition: a transformation $L(\cdot)$ is linear if
 - additivity: $L(\mathbf{x}_1 + \mathbf{x}_2) = L(\mathbf{x}_1) + L(\mathbf{x}_2)$
 - scaling: $L(a\mathbf{x}) = aL(\mathbf{x})$, where a is scalar
 - Features of $L(\cdot)$ given $\mathbf{x} = x_1e_1 + \dots + x_ne_n$
 - same scalar for coordinates
 - $\Rightarrow L(\mathbf{x}) = L(x_1e_1 + \dots + x_ne_n)$
 - $= L(x_1e_1) + \dots + L(x_ne_n)$
 - $= x_1L(e_1) + \dots + x_nL(e_n)$
 - \Rightarrow transformed vector $\mathbf{x}' = L(\mathbf{x})$ has the same coord under the transformed basis
- Linear Map
 - Definition
 - map $F : V \rightarrow X$ is a linear map if it is a linear transformation, where V, X are vector spaces
- Multilinear Maps
 - Definition
 - map $F : \underbrace{V \times \dots \times V}_{k \text{ copies}} \rightarrow X$ is multilinear/ k -linear if it is linear in each slot
 - i.e. $F(\mathbf{v}_1, \dots, a\mathbf{v}_i + b\mathbf{v}'_i, \dots, \mathbf{v}_k) = aF(\mathbf{v}_1, \dots, \mathbf{v}_i, \dots, \mathbf{v}_k) + bF(\mathbf{v}_1, \dots, \mathbf{v}'_i, \dots, \mathbf{v}_k)$
 - i.e. for fixed $\mathbf{v}_1, \dots, \mathbf{v}_{i-1}, \mathbf{v}_{i+1}, \dots, \mathbf{v}_k$, F reduced to linear map with \mathbf{v}_i as variable (where V, X are vector spaces)
 - Alternating Maps
 - map F is alternating if, its output is $\mathbf{0}$ whenever two vectors in inputs are identical
 - for Multilinear Map F : F Alternating $\Leftrightarrow F(\dots, \mathbf{v}, \dots, \mathbf{w}, \dots) = -F(\dots, \mathbf{w}, \dots, \mathbf{v}, \dots)$
 - i.e. for multilinear map F , F alternating \Leftrightarrow swapping two inputs flips sign of output
 - proof: given multilinear and alternating map F , for any \mathbf{v}, \mathbf{w}
 - $0 = F(\dots, (\mathbf{v} + \mathbf{w}), \dots, (\mathbf{v} + \mathbf{w}), \dots)$
 - $= F(\dots, \mathbf{v}, \dots, \mathbf{v}, \dots) + F(\dots, \mathbf{w}, \dots, \mathbf{w}, \dots) + F(\dots, \mathbf{v}, \dots, \mathbf{w}, \dots) + F(\dots, \mathbf{w}, \dots, \mathbf{v}, \dots)$
 - $= F(\dots, \mathbf{v}, \dots, \mathbf{w}, \dots) + F(\dots, \mathbf{w}, \dots, \mathbf{v}, \dots)$
 - $\Rightarrow F(\dots, \mathbf{v}, \dots, \mathbf{w}, \dots) = -F(\dots, \mathbf{w}, \dots, \mathbf{v}, \dots)$
 - proof: given multilinear map $F : F(\dots, \mathbf{v}, \dots, \mathbf{w}, \dots) = -F(\dots, \mathbf{w}, \dots, \mathbf{v}, \dots)$
 - $\Rightarrow F(\dots, \mathbf{v}, \dots, \mathbf{v}, \dots) = -F(\dots, \mathbf{v}, \dots, \mathbf{v}, \dots)$
 - $\Rightarrow F(\dots, \mathbf{v}, \dots, \mathbf{v}, \dots) = 0$, hence alternating

Matrix

- Matrix for Linear Transformation $L : S \rightarrow S'$
 - Representing Space Transformation
 - package the transformed basis under the original basis using matrix M
i.e. represent the $e'_1, \dots, e'_n = L(e_1), \dots, L(e_n)$ under the original basis e_1, \dots, e_n
 $\Rightarrow M = [e'_1, \dots, e'_n]$, with all transformed basis as column vectors
 $\Rightarrow M$ represent the result of linear transformation for the basis of S
 - hence, determine a linear space transformation $L : S \rightarrow S'$ using e_1, \dots, e_n
for $M_{m \times n}$ matrix: a linear transformation from n -D space to m -D space
 - $m < n$: projecting to subspace
 - $m > n$: expanding into a hyper-plane/-line/etc (constrained in hyper-space)
 - Performing Space Transformation
 - $\forall i \in \{1, \dots, n\}, x'_i = i^{\text{th}}$ component of $\mathbf{x}' = L(\mathbf{x})$, then $x'_i = (e'_{1i} + \dots + e'_{ni})x_i$
(as proved above)
 - \Rightarrow output vector $\mathbf{x}' = L(\mathbf{x}) = M\mathbf{x}$ under the original basis e_1, \dots, e_n
hence the rule for matrix multiplication

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = x \begin{bmatrix} a \\ c \end{bmatrix} + y \begin{bmatrix} b \\ d \end{bmatrix} = \begin{bmatrix} ax + by \\ cx + dy \end{bmatrix}$$

where green and red columns the $\{e'_1, e'_2\}$ under $\{e_1, e_2\}$, $[x, y]$ the input vector \mathbf{x}

- Matrices for Composition of Linear Transformation
 - Transformation L_1, L_2 as Matrix M_1, M_2
 - as easy to prove $M_2 \cdot (M_1 \cdot \mathbf{x}) = (M_1 \cdot M_2) \cdot \mathbf{x}$
 $\Rightarrow L_2(L_1(\cdot)) \Leftrightarrow$ the composition of transformation defined by $M_2 \cdot M_1$
 - $\Rightarrow \mathbf{x}$ first transformed by L_1 then $L_2 \Leftrightarrow$ transformed by $M_2 \cdot M_1$
 - Linear Transformation on Space
 - given a basis matrix $E = [e_1, \dots, e_n]$, a transformed basis $L_1(E) = L_1(e_1), \dots, L_1(e_n)$
(e_1, \dots, e_n as column vectors)
 $\Rightarrow L_2(L_1(e_k))$ performs L_2 transformation on the k^{th} vector of $L_1(E)$
 - hence to perform L_2 on the transformed space $L_1(E) \Rightarrow L_2(L_1(E))$
 - given that $L_1(E) = M_1 \Rightarrow L_2(L_1(E)) = M_2 \cdot M_1$
(due to the derivation of linear transformation as matrix)
 - hence, $M_2 \cdot M_1$ denotes
 - a further L_2 transformation on a transformed space $L_1(E)$
(all result represented under the original basis E)
 - the final transformed basis (first L_1 then L_2) under original basis E
 \Rightarrow denotes a composite transformation of $L_2(L_1(\cdot))$
 - Multiplication between Matrices
 - \Rightarrow composite linear transformations into single linear transformation
 - \Rightarrow linear transformation on vectors/space (generalized from single vector)
- Understanding Properties
 - $(AB)C = A(BC)$
 - both meaning apply transformation C then B then A ...

$$\overbrace{\begin{bmatrix} a & b \\ c & d \end{bmatrix}}^{M_2} \overbrace{\begin{bmatrix} e & f \\ g & h \end{bmatrix}}^{M_1} = \begin{bmatrix} ae + bg & af + bh \\ ce + dg & cf + dh \end{bmatrix}$$

where $M_1 = [L_1(e_1), L_1(e_2)]$, $M_2 = [L_2(e_1), L_2(e_2)]$, the result $= [L_2(L_1(e_1)), L_2(L_1(e_2))]$
(all under basis $E = [e_1, e_2]$)

Inverse of Matrix

- Inverse of Linear Transformation in Square Matrix
 - Interpretation
 - a transformation L^{-1} to inverse the effect of another transformation L
 $\Rightarrow \forall \mathbf{x}, L^{-1}L(\mathbf{x}) = \mathbf{x}$
 - $\Rightarrow M^{-1}M = I$, where M for L , M^{-1} for L^{-1} , I the identity matrix
 - Requirement for $M_{n \times n}$
 - transformed basis in M still span the space!
 (otherwise, transformed vectors projected onto subspace & can NOT be reversed)
 - \Rightarrow column vectors (transformed basis) in $M_{n \times n}$ linearly Independent
- Non-Square Matrix
 - Definition
 - left inversion: $M_{n \times m}^{-1} \cdot M_{m \times n} = I_{n \times n}$
 - right inversion: $M_{m \times n} \cdot M_{n \times m}^{-1} = I_{m \times m}$

Rank and Dimension

- Dimension
 - Definition
 - the number of linearly Independent vectors in a vector space
 \Rightarrow the least number of vector required to span the space
 - Dimension of Column Space
 - \Rightarrow the linearly independent col vectors
 i.e. the rank of M
- Measuring Linear Transformation on Dimension
 - Measuring the Transformed Space
 - rank of $M : R(M) = r$: basis in M span a r -dimension (column) space
 note: the column space of M : transformed space defined by column basis
 - full rank: the dimension of column space is as high as possible
- Null Space (Kernel) of $M_{m \times n}$
 - Vectors in Null Space
 - $\forall \mathbf{x}, M\mathbf{x} = \mathbf{0}$
 (i.e. all the vectors in null space transformed into $\mathbf{0}$ by M , hence the name)
 - $\mathbf{0}$ always in null space
 - Dimension of Null Space

- $D(\text{null space}) = n - R(M)$
- understanding: $R(M)$ vectors chosen for basis of col space of M
 $\Rightarrow \mathbf{x}$ able to freely combine the remaining col vectors, with $M\mathbf{x} = 0$ still satisfied
- Understanding Properties
 - $R(M_{m \times n}) \leq \min\{m, n\}$
 - $m < n$: projecting a n -D vector to a m -D subspace, hence at most of rank m
 - $m > n$: M consists of n vectors of m dimensions, define at most n -D space
 - $R(M_{m \times n}) = \min\{m, n\} \Leftrightarrow M$ Full Rank
 - from definition, it is as high as possible
 - m vectors with n dimensions, span at most a $\min\{m, n\}$ space
 either linearly dependent ($m < n$), or not enough independent vectors ($m > n$)
- Linear Dependency
 - $R(M_{m \times n}) = \min\{m, n\}$ (Full Rank)
 - $m = n$: M col vectors linear INdependent as n vector spans an N -D space
 - $m < n$: M col vectors linear dependent as n vector spans an M -D subspace
 - $m > n$: M col vectors linear INdependent as n vector spans at least an N -D space
 - $R(M_{m \times n}) < \min\{m, n\}$
 - $m \leq n$: M col vector linear dependent
 - $m > n$: M col vector MAY be linear independent/dependent

Dot Product

- Projecting to 1-D
 - Unit Orthogonal Basis
 - $\mathbf{x} \cdot \mathbf{y} = x_1 y_1 + \dots + x_n y_n = \mathbf{x}^T \cdot \mathbf{y} = \mathbf{y}^T \cdot \mathbf{x}$
 (\mathbf{x}, \mathbf{y} assumed to be column vector / matrix with single column)
 - General Basis $E = [e_1, \dots, e_n]$
 - $\Rightarrow \mathbf{x} = x_1 e_1 + \dots + x_n e_n = E\mathbf{x}, \mathbf{y} = y_1 e_1 + \dots + y_n e_n = E\mathbf{y}$
 $\Rightarrow \mathbf{x} \cdot \mathbf{y} = (E\mathbf{x}) \cdot (E\mathbf{y}) = \mathbf{x}^T E^T E \mathbf{y}$
 - understanding:
 1. $\mathbf{x} = E\mathbf{x}$: transfer back to a representation under unit orthogonal basis
 2. for $E = I$, back to the unit orthogonal case
 - Understanding
 - each basis $e_i \in E$ projected to 1-D space (a scalar) by transformation \mathbf{x}/\mathbf{y}
 - $\Rightarrow \mathbf{x} \cdot \mathbf{y}$: projecting \mathbf{x}/\mathbf{y} to 1-D line using transformation \mathbf{y}/\mathbf{x}
 (direction/scaling effect of projection can be alter by choice of basis E though)
- Duality
 - Dual Vector
 - the vector $\in \mathbb{R}^n$ represent a projection (linear transformation) to 1-D line
 (hence the vector equivalent to the matrix with 1 row defining the projection)
 - \Rightarrow performing transformation on a vector \Leftrightarrow taking product with the dual vector

Exterior (Wedge) Product

- Overview

- Definition

- k th exterior product $\wedge^k V$ is a vector space, with a map $\psi : \underbrace{V \times \dots \times V}_{k \text{ times}} \rightarrow \wedge^k V$,

where $V \in \mathbb{R}^n$ a vector space, map ψ the exterior multiplication

(note: $\psi(\mathbf{v}_1, \dots, \mathbf{v}_k) = \mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k$, also called k -blade or k -form)

- Properties of Space $\wedge^k V$

- ψ is alternating multilinear map
- for basis $\{e_1, \dots, e_n\}$ of $V \in \mathbb{R}^n$
 $\Rightarrow \{e_{i_1} \wedge \dots \wedge e_{i_k} | 1 \leq i_1 < \dots < i_k \leq n\}$ a basis of $\wedge^k V$
 $\Rightarrow e_{i_1} \wedge \dots \wedge e_{i_k}$ can form any permutation of the same input by swapping order (due to its alternating multilinearity)
 e.g. for $\wedge^2 V : e_1 \wedge e_1 = e_2 \wedge e_2 = \mathbf{0}, e_1 \wedge e_2 = -e_2 \wedge e_1 \Rightarrow$ not linearly independent
- sum of k -wedge is still in the vector space $\wedge^k V$ (as can be expressed by its basis)
- $F : \underbrace{V \times \dots \times V}_k \rightarrow X$ uniquely factors into $\underbrace{V \times \dots \times V}_k \xrightarrow{\psi} \wedge^k V \xrightarrow{F} X$,
 where F any alternating multilinear map, ψ exterior multiplication, \underline{F} linear map (since ψ also an alternating multilinear map & \underline{F} to preserve linearity)

- Inferred Propositions

- with $V \in \mathbb{R}^n$, $\dim \wedge^k V = \binom{n}{k} = C_n^k$, due to the form of its basis
 note: for $k > n$, $\wedge^k V = \{0\}$
 (top exterior power)
- k -blade $\mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k = 0 \Leftrightarrow \mathbf{v}_1, \dots, \mathbf{v}_k$ linearly dependent
 (k -blade $\mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k \neq 0 \Leftrightarrow \mathbf{v}_1, \dots, \mathbf{v}_k$ linearly INdependent)
 - if $\mathbf{v}_1, \dots, \mathbf{v}_k$ linearly dependent $\Rightarrow \mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k = 0$ as ψ alternating
 - if $\mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k = 0$, assume $\mathbf{v}_1, \dots, \mathbf{v}_k$ INdependent
 $\Rightarrow \mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k$ a basis of $\wedge^k V \Rightarrow$ basis $= 0$
 \Rightarrow assumption failed, $\mathbf{v}_1, \dots, \mathbf{v}_k$ linearly dependent
 (wedge dependence lemma)
- $\mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k = c(\mathbf{w}_1 \wedge \dots \wedge \mathbf{w}_k) \Leftrightarrow \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_k) = \text{span}(\mathbf{w}_1, \dots, \mathbf{w}_k)$,
 where $\mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k, \mathbf{w}_1 \wedge \dots \wedge \mathbf{w}_k \in \wedge^k V$ & non-zero, scalar $c \neq 0$
 - if $\text{span}(\mathbf{v}_1, \dots, \mathbf{v}_k) = \text{span}(\mathbf{w}_1, \dots, \mathbf{w}_k)$
 \Rightarrow each $\mathbf{v}_1, \dots, \mathbf{v}_k$ expressed by a linear combination of $\mathbf{w}_1, \dots, \mathbf{w}_k$
 $\Rightarrow \mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k = c(\mathbf{w}_1 \wedge \dots \wedge \mathbf{w}_k)$ as ψ alternating multilinear
 $(c \neq 0 \text{ as both pure wedges non-zero})$
 - if $\mathbf{v}_1 \wedge \dots \wedge \mathbf{v}_k = c(\mathbf{w}_1 \wedge \dots \wedge \mathbf{w}_k)$
 let $U_v = \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_k), U_w = \text{span}(\mathbf{w}_1, \dots, \mathbf{w}_k), U_{vw} = U_v \cap U_w$
 \Rightarrow assume $U_v \neq U_w, U_{vw}$ has dimension l
 $\Rightarrow \exists l$ vectors in $\{\mathbf{v}_1, \dots, \mathbf{v}_k\}, \{\mathbf{w}_1, \dots, \mathbf{w}_k\}$ as 2 sets of basis of U_{vw}
 \Rightarrow change l vectors to be a common basis for U_{vw} : $\{\mathbf{u}_1, \dots, \mathbf{u}_l\}$
 \Rightarrow now $\mathbf{v}_1, \dots, \mathbf{v}_{k-l}, \mathbf{u}_1, \dots, \mathbf{u}_l, \mathbf{w}_1, \dots, \mathbf{w}_{k-l}$ linearly independent
 $\Rightarrow \psi(\mathbf{v}_1, \dots, \mathbf{v}_{k-l}, \mathbf{u}_1, \dots, \mathbf{u}_l) \psi(\mathbf{w}_1, \dots, \mathbf{w}_{k-l}, \mathbf{u}_1, \dots, \mathbf{u}_l)$ 2 different basis for $\wedge^k V$
 while $\psi(\mathbf{v}_1, \dots, \mathbf{v}_{k-l}, \mathbf{u}_1, \dots, \mathbf{u}_l) = c_v \psi(\mathbf{v}_1, \dots, \mathbf{v}_k)$
 similarly $\psi(\mathbf{w}_1, \dots, \mathbf{w}_{k-l}, \mathbf{u}_1, \dots, \mathbf{u}_l) = c_w \psi(\mathbf{w}_1, \dots, \mathbf{w}_k)$
 $\Rightarrow \psi(\mathbf{v}_1, \dots, \mathbf{v}_k) \neq c \psi(\mathbf{w}_1, \dots, \mathbf{w}_k), (c \neq 0)$ (as basis linearly independent)
 \Rightarrow conflict, assumption failed, hence $\text{span}(\mathbf{v}_1, \dots, \mathbf{v}_k) = \text{span}(\mathbf{w}_1, \dots, \mathbf{w}_k)$

- Induced Maps

-
- Understanding
 - Construction from Geometry
 - e.g. $V \in \mathbb{R}^3, \wedge^3 V$ has basis for:
each x, y, z axis, each xy, xz, yz plane and finally xyz volume

Determinant

- Construction from Exterior Product
 - Definition
 - Independence of Basis
 - Multiplicativity
 - Relationship to Invertibility
- Measuring Linear Transformation on Volume
 - Unit Volume
 - unit volume $v = \|e_1 \wedge \dots \wedge e_n\|$ with basis e_1, \dots, e_n
(for orthogonal basis, $v = \|e_1\| \times \dots \times \|e_n\|$)
 - after transformation: $L(v) = \|L(e_1) \wedge \dots \wedge L(e_n)\| = \|Me_1 \wedge \dots \wedge Me_n\|$
where \wedge is the exterior product
 - Measuring Change of Unit Volume
 - $\Rightarrow \det(M) = \frac{L(v)}{v} = \|M_0 \times \dots \times M_n\|$,
assuming unit orthogonal basis $E = I$, where I is identity matrix
(note: interpretation of $\det(\cdot) \leftrightarrow$ spacial interpretation of cross product \times)
 - hence, the rule of calculating $\det(\cdot)$
- Measuring Change of Orientation
 - Orientation of Space
 - jointly defined by the direction & order of the sequence of basis vector
i.e. the positive/negative part of each axis in sequence
 - Measuring the Change
 - $\det(M) < 0$ if axes flipped over once (for an odd times)
 $\det(M) > 0$ if flipped for even times
 - interpretation: the flipped axis approaches 0 then expanded into the negative
(measured by original basis E)
- Linear Dependency
 - $\det(M) = 0$
 - volume in current space becomes 0
 \Rightarrow dimensions decreases after transformation applied
 \Rightarrow i.e. transformed basis not able to span the current space
 - \Rightarrow basis in M NOT linearly INdependent! $\Leftrightarrow \det(M) = 0$
 - $\det(M) \neq 0$
 - volume still exist
 \Rightarrow dimensions remain & transformed basis still span the space

■ \Rightarrow basis in M is linearly INdependent $\Leftrightarrow \det(M) \neq 0$

- Understanding Properties

- $\det(M_1 M_2) = \det(M_1) \det(M_2)$:

- left: final volume & orientation changed after transformation M_2 then M_1
 - right: the volume scaled by one transformation, then further by the other; (similar for orientation, as measured by sign)

Cross Product

- Calculation via Determinant

- Practice

- for $\mathbf{v}_1, \dots, \mathbf{v}_{n-1} \in \mathbb{R}^n$, $\mathbf{v}_1 \times \dots \times \mathbf{v}_{n-1} = \det \begin{pmatrix} - & \mathbf{v}_1 & - \\ & \vdots & \\ - & \mathbf{v}_{n-1} & - \\ e_1 & \dots & e_n \end{pmatrix}$
 - construct matrix with each vectors as row & basis e_1, \dots, e_n as the last row
 \Rightarrow an $N \times N$ matrix (fake e_1, \dots, e_n as number)
 \Rightarrow re-combined into a vector $\mathbf{w} = \mathbf{v}_1 \times \dots \times \mathbf{v}_{n-1} \in \mathbb{R}^n$
 - actually, carry out the hodge star operator implicitly (via determinant)

- Direction

- perpendicular to the $\text{span}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})$

- Understanding of Calculation via Determinant

- Equivalent Linear Transformation

- view the constructed $\det(\cdot)$ as transformation $F : (\mathbf{v}_1, \dots, \mathbf{v}_{n-1}, \mathbf{x}) \rightarrow y \in \mathbb{R}$
 \Rightarrow a linear transformation to 1-D number line
 $\Rightarrow \exists \mathbf{p} \in \mathbb{R}^n, y = F(\mathbf{x}) = \mathbf{p} \cdot \mathbf{x}$ (due to duality of dot product)
 - note that: transformation \mathbf{p} (or F) is generated by vector $\mathbf{v}_1, \dots, \mathbf{v}_{n-1}$

- Algebraic View

- view variable $\mathbf{x} = x_1, \dots, x_n$ as the coefficient for each basis e_1, \dots, e_n
 - $\Rightarrow \mathbf{p}$ describes the linear transformation of F on the vector space (on arbitrary vector \mathbf{x})
 - \Rightarrow dual vector $\mathbf{p} = \mathbf{v}_1 \times \dots \times \mathbf{v}_{n-1}$

- Geometric View

- as $\det(\mathbf{v}_1, \dots, \mathbf{v}_{n-1}, \mathbf{x})$ measures their high-dimension volume y
 $\Rightarrow y = \mathbf{p} \cdot \mathbf{x} = \|\mathbf{p}\| \cdot \mathbf{x}_{\parallel \mathbf{p}}$,
where $\mathbf{x}_{\parallel \mathbf{p}}$ the component of \mathbf{x} parallel to \mathbf{p}
 - while $y = \text{volume}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1}) \cdot \mathbf{x}_{\perp \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})}$ (integration perspective),
where $\mathbf{x}_{\perp \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})}$ the component of \mathbf{x} perpendicular to $\text{span}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})$
 - \Rightarrow find a \mathbf{p} , s.t. $\forall \mathbf{x}, \det(\mathbf{v}_1, \dots, \mathbf{v}_{n-1}) \cdot \mathbf{x}_{\perp \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})} = \|\mathbf{p}\| \cdot \mathbf{x}_{\parallel \mathbf{p}}$
 - \Rightarrow only solution: $\mathbf{p} \perp \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})$ and $\|\mathbf{p}\| = \det(\mathbf{v}_1, \dots, \mathbf{v}_{n-1})$
(naturally accounting for signed volume via direction of \mathbf{p})

Linear System of Equations

- Linearity

- Linear Combination of Variables

- coefficients matrix A : holding the coefficient for each equation (in row)

- variables vector \mathbf{x} : holding variables as column vector
 - constants vector \mathbf{v} : holding target constant for each equations as column vector
- $\Rightarrow A\mathbf{x} = \mathbf{v}$ for a set of linear equations
- Linear Transformation Perspective
 - \mathbf{x}/\mathbf{v} as original/transformed vectors
 - columns of A (coefficients for the same variable) as transformed basis
- \Rightarrow finding a start position \mathbf{x} which, after transformation A , lands on \mathbf{v}
- Existence of Solution(s)
 - Transformed Basis Linearly INdependent
 - single solution exists, as $A\mathbf{x} = \mathbf{v} \Leftrightarrow \mathbf{x} = A^{-1}\mathbf{v}$
(use the inverse transformation A^{-1} to find the input \mathbf{x} using output \mathbf{v})
 \Rightarrow single unique \mathbf{x} found
 - Transformed Basis Linearly Dependent ($\det(A) = 0$)
 - multiple solutions: transformed basis in A linearly dependent with \mathbf{v}
 \Rightarrow i.e. \mathbf{v} in the column space of A
(special case where $\mathbf{v} = \mathbf{0}$: solution space = null space of A)
 - no solution: basis in A linearly INdependent with \mathbf{v}
 \Rightarrow i.e. can NOT possibly be described by transformed basis in A
- Cramer's Rule
 - Basis Represented by Determinant
 - $\forall i = 1, \dots, n, \det(e_1, \dots, e_{i-1}, \mathbf{x}, e_{i+1}, \dots, e_n) = \mathbf{x}_{\parallel e_i}$,
where $\mathbf{x}_{\parallel e_i}$ the component of \mathbf{x} on basis e_i (as the i^{th} coord of \mathbf{x})
 - similarly, $\det(a_1, \dots, a_{i-1}, \mathbf{v}, a_{i+1}, \dots, a_n) = \mathbf{v}_{\parallel a_i}$,
where a_i the i^{th} col vector of A
 \Rightarrow represent the i^{th} coord of transformed vector \mathbf{v} under transformed basis A
 - Transformation Described by Determinant
 - $\det(A)$ measures the change of unit volume from I to A
 - $\det(e_1, \dots, e_{i-1}, \mathbf{x}, e_{i+1}, \dots, e_n)$ the volume before transformation
 - $\det(a_1, \dots, a_{i-1}, \mathbf{v}, a_{i+1}, \dots, a_n)$ the volume after transformation
 - $\Rightarrow \det(a_1, \dots, a_{i-1}, \mathbf{v}, a_{i+1}, \dots, a_n) = \det(e_1, \dots, e_{i-1}, \mathbf{x}, e_{i+1}, \dots, e_n) \cdot \det(A)$
 - Solving \mathbf{x}
 - hence, $\mathbf{x}_{\parallel e_i} = \det(e_1, \dots, e_{i-1}, \mathbf{x}, e_{i+1}, \dots, e_n) = \frac{\det(a_1, \dots, a_{i-1}, \mathbf{v}, a_{i+1}, \dots, a_n)}{\det(A)}$

Eigenvector and Eigenvalue

- Invariability in Linear Transformation
 - Invariable Line
 - vector before and after transformation A are on the same **line**
 $\Rightarrow A\mathbf{x} = \lambda\mathbf{x}$ (under the same basis)
 \Rightarrow transformation A NOT rotate those \mathbf{x} , but only stretch them, by a factor λ
 - \mathbf{x} the eigenvector for those invariable line;
 λ the eigenvalue for corresponding stretching factor
 - Representing Transformation
 - for rotation, the eigenvector&value can be more expressive than a matrix

- Solving Eigenvector&value
 - solving $A\mathbf{x} = \lambda\mathbf{x}$ for \mathbf{x} and λ
 \Rightarrow equivalent with solving $(A - \lambda I)\mathbf{x} = \mathbf{0}$
 - a necessary condition for non-zero solution \mathbf{x} : $\det(A - \lambda I) = 0$
 (understanding 0: as solution existence condition for linear system of equations)
 (understanding 1: non-zero \mathbf{x} transformed by $A - \lambda I$ into $\mathbf{0}$)
 (understanding 2: non-zero \mathbf{x} perpendicular to all row vectors in $A - \lambda I$)
 (understanding 3: otherwise, $\mathbf{x} = (A - \lambda I)^{-1}\mathbf{0} = \mathbf{0}$)
 - \Rightarrow with no/single/multiple $\lambda \in \mathbb{R}$ solved, bring back to relation $(A - \lambda I)\mathbf{x} = \mathbf{0}$
 \Rightarrow to solve for a line/span, consisting of eigenvector(s)
 (note: no $\lambda \in \mathbb{R}$ solved, no eigenvector then)

Translation of Basis

- Translation between Coordinate
 - Notation
 - $\mathbf{x} = (x_1, \dots, x_n)$: a coordinate under basis $A_{n \times n}$
 - $\mathbf{b} = (b_1, \dots, b_n)$: a coordinate under basis $E = I$
 - \vec{v} : the same vector described by \mathbf{x} & \mathbf{b}
 - $A_{n \times n}$: the transformed basis described by basis before transformation (i.e. E)
 (with col vectors of A linearly INdependent)
 - $A_{n \times n}^{-1}$: the transformed basis s.t. $A^{-1}AE = E$
 \Rightarrow describe transformation $A \rightarrow E$ by basis before transformation (i.e. $AE = A$)
 - $M_{n \times n}$: a matrix for arbitrary transformation $L : E \rightarrow ?$, described by basis E
 - Translating Coordinate
 - $A\mathbf{x}$: describe \vec{v} by basis E
 \Rightarrow translate the vector coord from basis A to basis E
 - $A^{-1}\mathbf{b}$: describe \vec{v} by basis A
 \Rightarrow translate the vector coord from basis E to basis A
 - $\Rightarrow A\mathbf{x} = \mathbf{b}$
 - Translating Transformation
 - translate coord \mathbf{x} (under A) to be under $E = I$: $A\mathbf{x}$
 (translate to the basis where M known)
 - transform the vector by M : $MA\mathbf{x}$ (under basis E)
 - translate the vector back to be under basis A : $A^{-1}MA\mathbf{x}$
 - $\Rightarrow \mathbf{x}' = A^{-1}MA\mathbf{x}$, where $A^{-1}MA$ the transformation L described under basis A

Diagonal Matrix

- Overview
 - Definition
 - must be a square matrix
 - non-zero values allowed only on the diagonal of a matrix
- Transformation in Diagonal Matrix
 - Interpretation
 - as the transformation only stretch original basis I by diagonal value
 \Rightarrow all transformed basis is eigenvectors, with eigenvalue the diagonal value!

- Eigenbasis
 - the basis formed by eigenvectors of a matrix $A_{n \times n}$
- Diagonalizing Matrix
 - Practice
 - for transformation $A_{n \times n}$ under basis I , solve all eigenvalues $\lambda = [\lambda_1, \dots, \lambda_m]$
 - if $m = n$, check all eigenvectors $Q = [\mathbf{v}_1, \dots, \mathbf{v}_m]$ to be linearly INdependent (i.e. check if eigenvectors able to form a valid basis for n -D space)
 - \Leftrightarrow check there are n distinct eigenvalues u
 - if so, A said to be **diagonalizable**
 - \Rightarrow reconstruct transformation A under eigenbasis (view V as transformed basis)
 - $\Rightarrow Q^{-1}AQ = \Lambda$, or $A = Q\Lambda Q^{-1}$, where $\Lambda = \text{diag}[\lambda_1, \dots, \lambda_n]$
 - Understanding
 - numerical proof: $Q^{-1}AQ = Q^{-1}[\lambda_1 \mathbf{v}_1, \dots, \lambda_n \mathbf{v}_n] = \Lambda$
 - translating transformation A into a transformation A' under eigenbasis
 - \Rightarrow for \mathbf{x} , transformed into eigenbasis $Q^{-1}\mathbf{x}$
 - \Rightarrow perform k transformation in eigenbasis $\Lambda^k Q^{-1}\mathbf{x}$
 - \Rightarrow then transformed back $Q\Lambda^k Q^{-1}\mathbf{x}$
 - \Rightarrow easy to perform any diagonalizable transformation

1.2.2 Interchanging Coordinates

N-dimensional Spherical Coordinates

- Notation
 - N -dim Euclidean Space E_N
 - e_1, e_2, \dots, e_N : a group of orthonormal basis of E_N
 - $\mathbf{x} = (x_1, x_2, \dots, x_N)$: vector in E_n
 - \mathbf{x}_{i-N} : projection of \mathbf{x} onto subspace spanned by e_i, \dots, e_N
 - $\Rightarrow \mathbf{x}_{i-N} = \sum_{n=i}^N x_n e_n$
 - Spherical Coordinates
 - $r = \|\mathbf{x}\|$: the norm of \mathbf{x}
 - $\phi_i \in [0, \pi]$: angle between \mathbf{x}_{i-N} and e_i
 - $r_i = \|\mathbf{x}_{i-N}\|$: norm of projection \mathbf{x}_{i-N} , with $r_1 = r$
- Observation
 - Space e_1, \dots, e_N :
 - $\cos \phi_1 = \frac{\mathbf{x}e_1}{\|\mathbf{x}\|\|e_1\|} = \frac{x_1}{r}$
 - $\Rightarrow x_1 = r \cos \phi_1$
 - $\Rightarrow \mathbf{x} = r \cos \phi_1 e_1 + \sum_{n=2}^N x_n e_n$
 - Space e_2, \dots, e_N :
 - from above: $\mathbf{x}^2 = r^2 \cos^2 \phi_1 + \sum_{n=2}^N x_n^2 = r^2$
 - $\Rightarrow \sum_{n=2}^N x_n^2 = r^2 \sin^2 \phi_1$

$$\begin{aligned}
& \blacksquare \mathbf{x}_{2-N} = \sum_{n=2}^N x_n e_n \\
& \Rightarrow \begin{cases} \|\mathbf{x}_{2-N}\|^2 = \sum_{n=2}^N x_n^2 = r^2 \sin^2 \phi_1 = r_2^2 \\ \cos \phi_2 = \frac{\mathbf{x}_{2-N}^{n=2} \cdot e_2}{\|\mathbf{x}_{2-N}\| \|e_2\|} = \frac{x_2}{r_2} \end{cases} \\
& \Rightarrow \begin{cases} r_2 = r \sin \phi_1 \\ x_2 = r_2 \cos \phi_2 = r \sin \phi_1 \cos \phi_2 \end{cases} \quad (\text{as } \phi_1 \in [0, \pi]) \\
& \Rightarrow \mathbf{x}_{2-N} = r \sin \phi_1 \cos \phi_2 e_2 + \sum_{n=3}^N x_n e_n
\end{aligned}$$

◦ Space e_3, \dots, e_N :

$$\begin{aligned}
& \blacksquare \text{from above: } \mathbf{x}_{2-N}^2 = r^2 \sin^2 \phi_1 \cos^2 \phi_2 + \sum_{n=3}^N x_n^2 = r^2 \sin^2 \phi_1 \\
& \Rightarrow \sum_{n=3}^N x_n^2 = r^2 \sin^2 \phi_1 \sin^2 \phi_2 \\
& \blacksquare \mathbf{x}_{3-N} = \sum_{n=3}^N x_n e_n \\
& \Rightarrow \begin{cases} \|\mathbf{x}_{3-N}\|^2 = \sum_{n=3}^N x_n^2 = r^2 \sin^2 \phi_1 \sin^2 \phi_2 = r_3^2 \\ \cos \phi_3 = \frac{\mathbf{x}_{3-N}^{n=3} \cdot e_3}{\|\mathbf{x}_{3-N}\| \|e_3\|} = \frac{x_3}{r_3} \end{cases} \\
& \Rightarrow \begin{cases} r_3 = r \sin \phi_1 \sin \phi_2 \quad (\text{as } \phi_1, \phi_2 \in [0, \pi]) \\ x_3 = r_3 \cos \phi_3 \end{cases} \\
& \Rightarrow \mathbf{x}_{3-N} = r \sin \phi_1 \sin \phi_2 e_3 + \sum_{n=4}^N x_n e_n
\end{aligned}$$

• Proof for x_i

◦ Procedure

$$\begin{aligned}
& \blacksquare \mathbf{x}_{i-N} = \sum_{n=i}^N x_n e_n \\
& \Rightarrow \cos \phi_i = \frac{\mathbf{x}_{i-N} \cdot e_i}{\|\mathbf{x}_{i-N}\| \|e_i\|} = \frac{x_i}{r_i} \\
& \Rightarrow x_i = r_i \cos \phi_i
\end{aligned}$$

• Induction

◦ Goal

$$\blacksquare \forall i \geq 2, r_i = r \prod_{j=1}^{i-1} \sin \phi_j$$

◦ Base Case ($i = 2$)

$$\blacksquare \text{as in observation, } r_2 = r \sin \phi_1 = r \prod_{j=1}^{2-1} \sin \phi_j$$

◦ Step Case

- assumption $r_i = r \prod_{j=1}^{i-1} \sin \phi_j$
- procedure: $\mathbf{x}_{i-N} = \sum_{n=i}^N x_n e_n = r_i \cos \phi_i + \sum_{n=i+1}^N x_n e_n$
- $\Rightarrow \|x_{i-N}\|^2 = r_i^2 \cos^2 \phi_i + \sum_{n=i+1}^N x_n^2 = r_i^2$
- $\Rightarrow \|x_{i+1-N}\|^2 = \sum_{n=i+1}^N x_n^2 = r_i^2 \sin^2 \phi = r_{i+1}^2$
- $\Rightarrow r_{i+1} = r_i \sin \phi_i = r \prod_{j=1}^i \sin \phi_j$

- Derivation

- x_i from Combined Proofs

- $x_i = \begin{cases} r \cos \phi_1 & i = 1 \\ r \cos \phi_i \prod_{j=1}^{i-1} \sin \phi_j & 2 \leq i \leq N \end{cases}$

- Last 2 Dimensions

- $\mathbf{x}_{(N-1)-N} = x_{N-1} \cdot e_{N-1} + x_N \cdot e_N$, with $r_{N-1} = r \prod_{j=1}^{N-2} \sin \phi_j$
 - $\Rightarrow \|\mathbf{x}_{(N-1)-N}\| = f(\phi_{N-1}, \phi_N) = r_{N-1}$
 - $\Rightarrow \phi_{N-1}, \phi_N$ not independent!
 - (actually, if $e_N = e_{N-1} + \frac{\pi}{2}$, then $\phi_N = \phi_{N-1} - \frac{\pi}{2}$)
 - \Rightarrow define $\theta \in [0, 2\pi)$ instead of $\phi_{N-1}, \phi_N \in [0, \pi]$
 - $\Rightarrow x_{N-1} = r_{N-1} \sin \theta, x_N = r_{N-1} \cos \theta$ (interchangeable)

- Final Spherical Coordinates

- $x_i = \begin{cases} r \cos \phi_1 & i = 1 \\ r \cos \phi_i \prod_{j=1}^{i-1} \sin \phi_j & 2 \leq i \leq N-1 \\ r \sin \theta \prod_{j=1}^{N-2} \sin \phi_j & i = N-1 \\ r \cos \theta \prod_{j=1}^{N-2} \sin \phi_j & i = N \end{cases}$

1.3 Calculus

1.3.1 Integral

Interchanging Coordinates in Integral

- General Theory

- Notation

- (x, y) : coordinate under Field D
 - (u, v) : coordinate under Field D'

$$\blacksquare T : \begin{cases} x = x(u, v), \\ y = y(u, v) \end{cases} : \text{transformation from } D \text{ to } D'$$

○ Assumption

- $f(x, y)$ continuous in D
- transformation T 's partial 1st order derivatives continuous on D'
- transformation T 's Jacobian $J(u, v) = \frac{\partial(x, y)}{\partial(u, v)} \neq 0$
- transformation $T : D \rightarrow D'$ is 1-1 mapping

○ Derivation

- take infinitely small square in D' :

$$\begin{array}{ll} M_4'(u, v + \delta v), & M_3'(u + \delta u, v + \delta v), \\ M_1'(u, v), & M_2'(u + \delta u, v) \end{array}$$

\Rightarrow after transformation to D :

$$\begin{array}{ll} M_4(x(u, v + \delta v), y(u, v + \delta v)), & M_3(x(u + \delta u, v + \delta v), y(u + \delta u, v + \delta v)), \\ M_1(x(u, v), y(u, v)), & M_2(x(u + \delta u, v), y(u + \delta u, v)) \end{array}$$

$$\Rightarrow x_2 - x_1 = x(u + \delta u, v) - x(u, v) = \frac{\partial x}{\partial u}|_{(u, v)} \delta u$$

$$x_4 - x_1 = x(u, v + \delta v) - x(u, v) = \frac{\partial x}{\partial v}|_{(u, v)} \delta v$$

$$y_2 - y_1 = y(u + \delta u, v) - y(u, v) = \frac{\partial y}{\partial u}|_{(u, v)} \delta u$$

$$y_4 - y_1 = y(u, v + \delta v) - y(u, v) = \frac{\partial y}{\partial v}|_{(u, v)} \delta v$$

as $\delta u, \delta v \rightarrow 0$, curvilinear boundary quadrilateral $M_1 M_2 M_3 M_4 \rightarrow$ parallelogram

$$\begin{aligned} \Rightarrow S_{M_1 M_2 M_3 M_4} &= |\overrightarrow{M_1 M_2} \times \overrightarrow{M_1 M_4}| = \left| \begin{vmatrix} x_2 - x_1 & y_2 - y_1 \\ x_4 - x_1 & y_4 - y_1 \end{vmatrix} \right| \\ &= \left| \begin{vmatrix} \frac{\partial x}{\partial u} \delta u & \frac{\partial y}{\partial u} \delta u \\ \frac{\partial x}{\partial v} \delta v & \frac{\partial y}{\partial v} \delta v \end{vmatrix} \right| = \left| \begin{vmatrix} \frac{\partial x}{\partial u} & \frac{\partial y}{\partial u} \\ \frac{\partial x}{\partial v} & \frac{\partial y}{\partial v} \end{vmatrix} \right| \delta u \delta v \\ &= |J(u, v)| \delta u \delta v \end{aligned}$$

- \Rightarrow infinitely small area $\delta\sigma = dxdy = |J(u, v)| \delta u \delta v$

$$\Rightarrow \int \int_D f(x, y) dxdy = \int \int_{D'} f(x(u, v), y(u, v)) |J(u, v)| dudv$$

• Integral in Cartesian \rightarrow Polar

○ Result

$$\blacksquare dxdy = r dr d\theta$$

○ Derivation

- from general transformation: $x = r \cos(\theta), y = r \sin(\theta)$
 $\Rightarrow dxdy = |J(r, \theta)| dr d\theta = r dr d\theta$
- from direct calculation of infinite small area in polar coordinate
 $\Rightarrow d\sigma = \frac{1}{2}(r + dr)^2 d\theta - \frac{1}{2}r^2 d\theta = r dr d\theta + \frac{1}{2}(dr)^2 d\theta$
 $\Rightarrow d\sigma = r dr d\theta$, when $dr, d\theta \rightarrow 0$

Gaussian Integral

• Gaussian Function

$$\circ f(x) = e^{-a(x+b)^2}$$

- special form: $f(x) = e^{-(x)^2}$
- alternative form: $f(x) = e^{ax^2+bx+c}$
- no indefinite integral $\int_a^b e^{-x^2}$
- only definite integral $\int_{-\infty}^{+\infty} e^{-x^2}$

- Direct Integral

$$\begin{aligned}
 \circ \left(\int_{-\infty}^{+\infty} e^{-a(x+b)^2} dx \right)^2 &= \int_{-\infty}^{+\infty} e^{-a(x+b)^2} dx \int_{-\infty}^{+\infty} e^{-a(y+b)^2} dy \\
 &= \int \int_{-\infty}^{+\infty} e^{-a[(x+b)^2+(y+b)^2]} d(x+b)d(y+b) = \int \int_{-\infty}^{+\infty} e^{-a(x^2+y^2)} dx dy \\
 &= \int_0^{2\pi} \int_0^{+\infty} e^{-ar^2} r dr d\theta \\
 &= \frac{\pi}{a} \\
 \Rightarrow \int_{-\infty}^{+\infty} e^{-a(x+b)^2} dx &= \sqrt{\frac{\pi}{a}}, \text{ alternatively } \int_{-\infty}^{+\infty} e^{ax^2+bx+c} dx = \sqrt{\frac{\pi}{-a}} \cdot e^{\frac{b^2}{4a}+c}
 \end{aligned}$$

- Even Moment of Gaussian Function

$$\begin{aligned}
 \circ \int_{-\infty}^{+\infty} x^{2n} e^{-ax^2} dx &= (-1)^n \int_{-\infty}^{+\infty} \frac{\partial^n}{\partial a^n} e^{-ax^2} dx \\
 &= (-1)^n \frac{\partial^n}{\partial a^n} \int_{-\infty}^{+\infty} e^{-ax^2} dx && \text{by parameter differeation} \\
 &= (-1)^n \sqrt{\pi} \frac{\partial^n}{\partial a^n} a^{-\frac{1}{2}} \\
 &= \sqrt{\frac{\pi}{a}} \frac{(2n-1)!!}{(2a)^n}, && \text{where !! is double factorial}
 \end{aligned}$$

1.4 Probability Theory

1.4.1 Introduction

Background

- Measuring Uncertainty
 - Source of Uncertainty
 - noise in reality & observation
 - finite size of data (limited information)
- Derivation
 - Quantifying Belief
 - by Cox (1946): if numerical values used to represent degrees of belief, a simple set of axioms encoding common sense properties of such beliefs will lead uniquely to a set of rules for manipulating degrees of belief that are equivalent to the sum and product rules of probability
 - Measuring Uncertainty
 - by Jaynes (2003): probability theory can be regarded as an extension of Boolean logic to situations involving uncertainty

- Common Destination
 - numerical quantities to measure uncertainty, derived from different sets of properties/axioms, behave precisely according to the rules of probability

The Basic

- Notation
 - X, Y : random variable
- Discrete
 - $P(X, Y)$: joint probability of X, Y taking their values
 - $P(X)$: marginal probability of X taking its value
 - $P(X|Y)$: conditional probability of X taking its value given Y observed / determined
- Continuous
 - $P(x) = P_X(x)$: cumulative probability of value for variable $X < x$
 - $p(x)$: probability density,
 - where $\lim_{\delta x \rightarrow 0} P(X \in (x, x + \delta x)) = \lim_{\delta x \rightarrow 0} p(x)\delta x \Rightarrow P(X \in (a, b)) = \int_a^b p(x)dx$
 - $\Rightarrow p(x) \geq 0$ and $\int_{-\infty}^{+\infty} p(x) = 1$
 - $\Rightarrow P(z) = \int_{-\infty}^z p(x)dx$
- Basic Rules
 - Sum Rule
 - $P(X) = \sum_Y P(X, Y)$, where X, Y are discrete
 - $P(X) = \int_Y P(X, Y)$, where X, Y are continuous
(formal justification requires measure theory)
 - Product Rule
 - $P(X, Y) = P(Y|X)P(X)$
 - $P(X, Y) = P(Y)P(X)$, where X, Y are independent
 - \Rightarrow Bayes' Rule
 - $P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} = \frac{P(X|Y)P(Y)}{\sum_Y P(X|Y)P(Y)}$, where Y are discrete
 - $P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)} = \frac{P(X|Y)P(Y)}{\int_Y P(X|Y)P(Y)}$, where Y are continuous
- Interpretation of Bayes
 - Normalization
 - the \sum, \int can be interpreted as a **normalization constant**
 $\Rightarrow \text{posterior} \propto \text{likelihood} \times \text{prior}$
 - Prior
 - $P(Y)$: available probability of desired variable **before** anything observed
 $\Rightarrow Y$ usually model parameters
 - Posterior

- $P(Y|X)$: obtained probability of desired variable **after** observation
 \Rightarrow if X, Y independent, observation has no effect \Rightarrow prior = posterior
- Likelihood
 - $P(X|Y)$: how probable/likely of X being observed under different setting of Y
- Prior \rightarrow Posterior
 - a process of incorporating the evidence provided by observation

1.4.2 Expectations and Covariances

Expectation

- Definition
 - Expectation of $f(x)$ under $p(x)$
 - discrete x : $\mathbb{E}_p[f] = \sum_x p(x)f(x)$
 - continuous x : $\mathbb{E}_p[f] = \int p(x)f(x)dx$
 - approximation with N points drawn from $p(x)$: $\mathbb{E}_p[f] \simeq \frac{1}{N} \sum_{n=1}^N f(x_n)$
 (when $N \rightarrow \infty$, \simeq becomes $=$)
 - Multivariate Expectation
 - Marginal Expectation of $f(x, y)$ on x : $\mathbb{E}_x[f(x, y)] = \sum_x p(x)f(x, y)$
 (hence a function of y)
 - Conditional Expectation $f(x)$ on $p(x|y)$: $\mathbb{E}[f|y] = \sum_x p(x|y)f(x)$
- Independence
 - Independent x, y
 - $\mathbb{E}_{xy}[x, y] = \sum_{x,y} p(x, y)xy = \sum_{x,y} p(x)p(y)xy = \mathbb{E}[x]\mathbb{E}[y]$

Variance

- Definition
 - Variance of $f(x)$:
 - $\text{var}[f] = \mathbb{E}[(f(x) - \mathbb{E}[f(x)])^2]$
 $= \mathbb{E}[f(x)^2] - \mathbb{E}[f(x)]^2$
 - Covariance

Covariance

- Definition
 - between Variables x, y
 - $\text{cov}[x, y] = \mathbb{E}_{x,y}[(x - \mathbb{E}[x])(y - \mathbb{E}[y])]$
 $= \mathbb{E}_{x,y}[xy] - \mathbb{E}[x]\mathbb{E}[y]$
 - between Vectors \mathbf{x}, \mathbf{y} (column vectors)

- $\text{cov}[\mathbf{x}, \mathbf{y}] = \mathbb{E}_{\mathbf{x}, \mathbf{y}}[(\mathbf{x} - \mathbb{E}[\mathbf{x}]) \cdot (\mathbf{y}^T - \mathbb{E}[\mathbf{y}^T])]$
 $= \mathbb{E}_{\mathbf{x}, \mathbf{y}}[\mathbf{x}\mathbf{y}^T] - \mathbb{E}[\mathbf{x}]\mathbb{E}[\mathbf{y}^T]$
(pairwise covariance between components of \mathbf{x}, \mathbf{y})
 - within Vector \mathbf{x}
 - $\text{cov}[\mathbf{x}] \equiv \text{cov}[\mathbf{x}, \mathbf{x}]$
(pairwise covariance between its components)
- Independence Variable
 - Independent x, y
 - $\text{cov}[x, y] = \mathbb{E}_{x, y}[xy] - \mathbb{E}[x]\mathbb{E}[y] = 0$

1.4.3 Transformations of Random Variables

Inverse Image

- Definition
 - Notation
 - function $g : \mathbb{R} \rightarrow \mathbb{R}$
 - set A in \mathbb{R}
 - Inverse Image on Set A
 - $g^{-1}(A) = \{x \in \mathbb{R} | g(x) \in A\}$
 $\Leftrightarrow x \in g^{-1}(A)$ if and only if $g(x) \in A$
interpretation: for each element in A , get its original value before g applied
- Properties
 - $g^{-1}(\mathbb{R}) = \mathbb{R}$, as g is defined on \mathbb{R}
 - \forall set $A, g^{-1}(A^c) = g^{-1}(A)^c$, where A^c is the complement of set A
 - \forall collection of sets $\{A_\lambda | \lambda \in \Lambda\}, g^{-1}\left(\bigcup_\lambda A_\lambda\right) = \bigcup_\lambda g^{-1}(A_\lambda)$
 - General Transformation $Y = g(X)$
 - $P(Y \in A) = P(g(X) \in A) = P(X \in g^{-1}(A))$

Discrete Variable

- Variable
 - X : random variable with probability mass function $P_X(x)$
 - $Y = g(X)$, with probability mass function $P_Y(y)$
- Probability Mass Function
 - $P_Y(y) = \sum_{x \in g^{-1}(y)} P_X(x)$, as $X = x$ is independent and mutually exclusive
note: $g^{-1}(y)$ denotes $g^{-1}(\{y\})$
 - Example
 - uniform random variable X on $\{-n, \dots, n\}$ with $Y = |X|$
 $\Rightarrow P_X(x) = \frac{1}{2n+1}$
 $\Rightarrow P_Y(y) = \begin{cases} \frac{1}{2n+1}, & x = 0, \\ \frac{2}{2n+1}, & x \neq 0. \end{cases}$

Continuous

- Variable
 - X : random variable with cumulative distribution $P_X(x)$, density $p_X(x)$
 - $Y = g(X)$, with cumulative distribution $P_Y(y)$, density $p_Y(y)$
- Cumulative Distribution
 - Strictly Monotone Increasing g
 - $P_Y(y) = P(Y \leq y) = P(g(X) \leq y) = P(X \leq g^{-1}(y)) = P_X(g^{-1}(y))$
 - Strictly Monotone Decreasing g
 - $P_Y(y) = P(Y \leq y) = P(g(X) \leq y) = P(X \geq g^{-1}(y)) = 1 - P_X(g^{-1}(y))$
- Probability Density
 - Strictly Monotone g (an one-to-one function)
 - $p_Y(y) = \frac{d}{dy}P_Y(y) = \frac{dP_Y(y)}{dg^{-1}(y)} \frac{dg^{-1}(y)}{dy} = p_X(g^{-1}(y)) \left| \frac{d}{dy}g^{-1}(y) \right|$,
as g^{-1} has the same monotony as g , combined with the sign in P_Y to give the $|\cdot|$


1.4.4 Gaussian Distribution**Definition**

- Univariate Gaussian
 - Variable
 - mean: μ
 - variance: $\sigma^2 \Rightarrow$ reciprocal of the variance $\beta = \frac{1}{\sigma^2}$ (also called precision)
 - Probability Dense Function (PDF)
 - $\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2}(x - \mu)^2 \right\}$
 \Rightarrow satisfying probability axioms: $\mathcal{N}(x|\mu, \sigma^2) > 0$ and $\int_{-\infty}^{+\infty} \mathcal{N}(x|\mu, \sigma^2) dx = 1$
 - Expectation
 - $\mathbb{E}[x] = \int_{-\infty}^{+\infty} \mathcal{N}(x|\mu, \sigma^2) x dx = \mu$
 $\Rightarrow \mathbb{E}[x^2] = \int_{-\infty}^{+\infty} \mathcal{N}(x|\mu, \sigma^2) x^2 dx$
 $= \frac{1}{(2\pi\sigma^2)^{1/2}} \int_{-\infty}^{+\infty} x^2 \exp \left\{ -\frac{1}{2\sigma^2}(x - \mu)^2 \right\} dx$
 $= \pi^{-\frac{1}{2}} \int_{-\infty}^{+\infty} (\sqrt{2\sigma^2}x + \mu)^2 \exp(-x^2) dx$, substituting $a = \frac{x - \mu}{\sqrt{2\sigma^2}}$
 $= \pi^{-\frac{1}{2}} (2\sigma^2 \int_{\mathbb{R}} x^2 e^{-x^2} dx + 2\mu\sqrt{2\sigma^2} \int_{\mathbb{R}} x e^{-x^2} dx + \mu^2 \int_{\mathbb{R}} e^{-x^2} dx)$
 $= \pi^{-\frac{1}{2}} (2\sigma^2 \int_{\mathbb{R}} x^2 e^{-x^2} dx + 2\mu\sqrt{2\sigma^2} \cdot 0 + \mu^2 \sqrt{\pi})$
 $= 2\sigma^2 \pi^{-\frac{1}{2}} \int_{\mathbb{R}} x^2 e^{-x^2} dx + \mu^2$
 $= \sigma^2 + \mu^2$, by 2nd moment of Gaussian or $(xe^{-x^2})' = e^{-x^2} - 2x^2e^{-x^2}$

- Variance
 - $\text{var}[x] = \mathbb{E}[x^2] - \mathbb{E}[x]^2 = \sigma^2$
- Multivariate (d -dimensional) Gaussian
 - Variable
 - mean: $\mu \in \mathbb{R}^d$
 - covariance matrix: $\Sigma_{d \times d}$
 - Probability Dense Function (PDF)
 - $\mathcal{N}_d(\mathbf{x}|\mu, \Sigma) = \frac{1}{(2\pi)^{d/2}|\Sigma|^{1/2}} \exp\left\{-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right\}$,
noted as $X \sim \mathcal{N}_d(x|\mu, \Sigma)$

Multivariate Gaussian

- Dimensionality
 - Volume of High Dimensional Sphere
 - for $n = 2k, k \in \mathbb{N}^+, V_{2k}(R) = \frac{\pi^k}{k!} R^{2k}$
 - for $n = 2k + 1, k \in \mathbb{N}, V_{2k+1}(R) = \frac{2^{k+1} \pi^k}{(2k+1)!!} R^{2k+1}$
 - $\Rightarrow \lim_{D \rightarrow +\infty} \frac{V_D(1) - V_D(1 - \epsilon)}{V_D(1)} = \lim_{D \rightarrow +\infty} 1 - (1 - \epsilon)^D = 1 \Rightarrow$ the volume of a D -sphere concentrate in a thin shell near the surface!
(actually, in the corner of a high dimensional cube as shown below)
 - Volume of High Dimensional Cube
 - \Rightarrow volume ratio of hyper sphere and hyper cube: \Rightarrow the volume of a D -cube concentrates in its corner! \Rightarrow distance function in high dimensional space CAN be useless
 - High Dimensional Distribution
 - High Dimensional Gaussian
 - probability density with respect to radius r for various dimension D
 \Rightarrow most density are in a thin shell at a specific r



- Facing High Dimensionality

- Convolution of Gaussian

- Integral of Gaussians $\int G_1 G_2 dx$

- $G_1 \sim \mathcal{N}_d(x|a, A), G_2 \sim \mathcal{N}_d(x|b, B)$

$$\begin{aligned}
&\Rightarrow \int \mathcal{N}_d(x|a, A) \mathcal{N}_d(x|b, B) dx \\
&= \int \frac{1}{(2\pi)^{d/2} |A|^{1/2}} e^{-\frac{1}{2}(x-a)^T A^{-1}(x-a)} \frac{1}{(2\pi)^{d/2} |B|^{1/2}} e^{-\frac{1}{2}(x-b)^T B^{-1}(x-b)} dx \\
&= \int \frac{1}{(2\pi)^{d/2} |A|^{1/2}} \frac{1}{(2\pi)^{d/2} |B|^{1/2}} e^{-\frac{1}{2}[(x-a)^T A^{-1}(x-a) + (x-b)^T B^{-1}(x-b)]} \\
&= \int \frac{1}{(2\pi)^{d/2} |A|^{1/2}} \frac{1}{(2\pi)^{d/2} |B|^{1/2}} e^{-\frac{1}{2}[(x-c)^T (A^{-1} + B^{-1})(x-c) + (a-b)^T C(a-b)]}, \\
&\quad \text{where } c = (A^{-1} + B^{-1})^{-1}(A^{-1}a + B^{-1}b), C = A^{-1}(A^{-1} + B^{-1})^{-1}B^{-1} = (A + B)^{-1} \\
&= \frac{|(A^{-1} + B^{-1})^{-1}|^{1/2}}{(2\pi)^{d/2} |A|^{1/2} |B|^{1/2}} e^{-\frac{1}{2}(a-b)^T C(a-b)} \int \frac{1}{(2\pi)^{d/2} |(A^{-1} + B^{-1})^{-1}|^{1/2}} e^{-\frac{1}{2}(x-c)^T (A^{-1} + B^{-1})(x-c)} dx \\
&= \frac{|(A^{-1} + B^{-1})^{-1}|^{1/2}}{(2\pi)^{d/2} |A|^{1/2} |B|^{1/2}} e^{-\frac{1}{2}(a-b)^T C(a-b)} \\
&= \frac{1}{(2\pi)^{d/2} (|A||B||A^{-1} + B^{-1}|)^{1/2}} e^{-\frac{1}{2}(a-b)^T (A+B)^{-1}(a-b)} \\
&= \frac{1}{(2\pi)^{d/2} |ABA^{-1} + ABB^{-1}|^{1/2}} e^{-\frac{1}{2}(a-b)^T (A+B)^{-1}(a-b)} \\
&= \frac{1}{(2\pi)^{d/2} |ABA^{-1} + A|} e^{-\frac{1}{2}(a-b)^T (A+B)^{-1}(a-b)} \\
&= \frac{1}{(2\pi)^{d/2} |A(B+A)A^{-1}|} e^{-\frac{1}{2}(a-b)^T (A+B)^{-1}(a-b)} \\
&= \frac{1}{(2\pi)^{d/2} |A+B|^{1/2}} e^{-\frac{1}{2}(a-b)^T (A+B)^{-1}(a-b)}
\end{aligned}$$

- \Rightarrow Convolution of Gaussians $G_1 * G_2$

- $G_1 \sim \mathcal{N}_d(a, A), G_2 \sim \mathcal{N}_d(b, B)$

$$\begin{aligned}
G_1 * G_2(z) &= \int G_1(x) G_2(z-x) dx \\
&= \int \mathcal{N}_d(x|a, A) \mathcal{N}_d(z-x|b, B) dx \\
&= \int \mathcal{N}_d(x|a, A) \cdot \frac{1}{(2\pi)^{d/2} |B|^{1/2}} e^{-\frac{1}{2}(z-x-b)^T B^{-1}(z-x-b)} dx \\
&= \int \mathcal{N}_d(x|a, A) \mathcal{N}_d(x|z-b, B) dx \\
&= \frac{1}{(2\pi)^{d/2} |A+B|^{1/2}} e^{-\frac{1}{2}(z-(a+b))^T (A+B)^{-1}(z-(a+b))} \\
&= \mathcal{N}_d(z|a+b, A+B)
\end{aligned}$$

1.4.5 Bayesian Interpretation of Probability

Contrasting Frequentist Estimator

- Posterior $p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w})p(\mathbf{w})}{p(\mathcal{D})}$

- Notation
 - \mathcal{D} the observed dataset
 - \mathbf{w} the vector for model parameters
- Bayesian
 - only one single dataset \mathcal{D} (the observed one)
 - uncertainty expressed as distribution over \mathbf{w}
 - model's error: use likelihood / posterior directly (or after taking log)
 - pros
 1. naturally incorporating prior knowledge as prior distribution (of \mathbf{w})
 - cons
 1. prior usually selected for mathematic convenience
- Frequentist Estimator
 - parameters \mathbf{w} already determined / fixed by 'estimator' (model)
 - error bars of the model obtained by considering the distribution over \mathcal{D}
 - model's error: bootstrap procedure
 1. generate dataset(s) by drawing from the observed \mathcal{D} with replacement
 2. sampling L datasets with the same size as \mathcal{D}
 3. error = variability of predictions between the sampled datasets
 - pros
 1. protect the conclusion from false prior knowledge
 - cons
 1. sensitive to observation (extreme cases), especially under small dataset

Parameter Estimation

- Bias vs. Variance
- Taking Logarithm
 - Reason
 - monotonically increasing function $\Rightarrow \arg \max_{\theta} f(x; \theta) = \arg \max_{\theta} \log f(x; \theta)$
 - simplify mathematical analysis $\Rightarrow \prod \rightarrow \sum$
 - numerical stability \Rightarrow avoid \prod (small probabilities)
(may otherwise underflow the numerical precision)
- Maximum Likelihood Estimation for Gaussian
 - Notation
 - $X = \{x_1, \dots, x_N\}$: observed N data points
 - Assumption
 - data points are i.i.d. (identically and independently distributed)
 - underlying distribution is Gaussian $\mathcal{N}(\mu, \sigma)$
 - Likelihood
 - $p(X|\mu, \sigma^2) = \prod_{n=1}^N p(x_n|\mu, \sigma^2)$
 - $\Rightarrow \ln p(X|\mu, \sigma) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (x_n - \mu)^2 - \frac{N}{2} \ln \sigma^2 - \frac{N}{2} \ln(2\pi)$

◦ Solution

$$\blacksquare \text{ let } \frac{\partial}{\partial \mu} \ln p(X|\mu, \sigma^2) = 0 \Rightarrow \mu_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N x_n$$

$$\blacksquare \text{ let } \frac{\partial}{\partial \sigma^2} \ln p(X|\mu, \sigma^2) = 0 \Rightarrow \sigma_{\text{ML}}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2$$

◦ Analysis

$$\blacksquare \mathbb{E}[\mu_{\text{ML}}] = \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N x_n\right] = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n] = \mu$$

(as x_1, \dots, x_N i.i.d, drawn from $\mathcal{N}(\mu, \sigma^2)$, thus $\sim \mathcal{N}(\mu, \sigma^2)$)

\Rightarrow unbiased estimation of mean

$$\blacksquare \mathbb{E}[\sigma_{\text{ML}}^2] = \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2\right]$$

$$= \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N (x_n^2 - 2\mu_{\text{ML}}x_n + \mu_{\text{ML}}^2)\right]$$

$$= \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n^2] - \mathbb{E}\left[\frac{1}{N} \sum_{n=1}^N 2x_n\mu_{\text{ML}}\right] + \mathbb{E}[\mu_{\text{ML}}^2]$$

$$= \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n^2] - 2\mathbb{E}[\mu_{\text{ML}}^2] + \mathbb{E}[\mu_{\text{ML}}^2]$$

$$= \frac{1}{N} \sum_{n=1}^N \mathbb{E}[x_n^2] - \frac{1}{N^2} \sum_{i,j=1}^N \mathbb{E}[x_i x_j]$$

$$= \frac{1}{N} \sum_{n=1}^N (\sigma^2 + \mu^2) - \frac{1}{N^2} [N(N-1)\mu^2 + N(\sigma^2 + \mu^2)]$$

(by 2nd moment of Gaussian $\mathbb{E}[x^2]$ and i.i.d assumption)

$$= \left(\frac{N-1}{N}\right) \sigma^2$$

\Rightarrow biased variance !

$$\Rightarrow \text{unbiased variance } \hat{\sigma}^2 = \frac{N}{N-1} \sigma_{\text{ML}}^2 = \frac{1}{N-1} \sum_{n=1}^N (x_n - \mu_{\text{ML}})^2$$

interpretation: $N-1$ degree of freedom,
(as calculating σ^2 needs μ , which help pin down x_N given x_1, \dots, x_{N-1})

Predictive Distribution

• Probabilistic Prediction

◦ Notation

- \mathbf{x}, \mathbf{t} : vector of data examples and corresponding ground truth
- \mathbf{w} : model parameters
- x, t : new data example for prediction and its ground truth

◦ Prediction by Model

- $p(t|x, \mathbf{w}')$, where \mathbf{w}' is the best fit parameters founded

◦ Prediction by Data

- $p(t|x, \mathbf{x}, \mathbf{t}) = \int p(t|x, \mathbf{w})p(\mathbf{w}|\mathbf{x}, \mathbf{t})d\mathbf{w}$, where \mathbf{w} marginalized over its posterior
(consider all possible \mathbf{w})

Chapter 2

Introduction

2.1 General Concern

2.1.1 Types of Learning

Supervised Learning

- Overview
 - training data comprises examples of input vectors with corresponding target vectors
- Regression
 - output one or more continuous variable
- Classification
 - assign input to one of a finite number of discrete categories

Unsupervised Learning

- Overview
 - training data consists of a set of input vectors without target vectors
- Clustering
 - Goal: discover groups of similar examples
- Density Estimation
 - Goal: determine the distribution of data within the input space
- Dimension Reduction
 - Goal: project data into low dimension, for the purpose of such as visualization

Reinforcement Learning

- Overview
 - input with time series & discover optimal output by a process of trial and error
- Goal
 - find actions to take under given circumstance to maximize a reward

2.2 Decision Theory

2.2.1 Overview

Formulation

- Analysis Target
 - Human Behavior
 - goal-directed behaviors in the presence of options (alternative actions)
- Goal
 - Normative Decision Theory
 - how the decision should be made, in order to be rational, etc.
 \Rightarrow
 - an normative theory is *weakly falsified*, if \exists a decision problem in which an agent can perform in contradiction with the theory without being irrational
 - an normative theory is *strictly falsified*, if \exists a decision problem in which an agent performing in accordance with the theory cannot be a rational agent
 - Descriptive Decision Theory
 - how decisions are actually made in practice
 - Specific Concern
 - interaction of agents: collective decision-making
 - behavior of individuals in decisions
 - rationality in decisions
 - coordinating decision over time, in a changing environment
- Preference in Decision Making
 - Preference Logic
 - defined relation \geq for two options a, b , according to agent's preference
 \Rightarrow form a partially ordered set (S, \geq) , where S the set of options
 - other relation $<, \equiv$ can be derived from \geq
 - Completeness
 - complete: $\forall i, j \in X$, a relation is defined
 (otherwise incomplete preference)
 - yet, usually incomplete preference, as determining preference takes effort
 - Transitivity
 - $a \geq b \wedge b \geq c \rightarrow a \geq c$ holds true for all a, b, c
 (hence all other relation $<, \equiv$)
 - \Rightarrow problem of accumulated negligible indifference (\equiv)
 e.g. a_0 as no sugar, a_{100} as full sugar, given $a_i \equiv a_{i+1}$ due to negligible difference
 \Rightarrow finally $a_0 \equiv a_{100}$ due to transitivity
 - Guideline
 - given the constructed partial order set (S, \geq) , choose one of its upper bound
 (with prevalent assumption of transitivity and completeness)
- Utility in Decision Making
 - Numerical Representation

- using pre-defined relation in number, with a preference of maximal utility
 \Rightarrow naturally transitive & complete
- relative number \Rightarrow meaningful in comparison within current option set
 (i.e. not comparable across different decision making process)
- **Comparability**
 - only comparable between options in the same decision process
- Guideline
 - maximization: choose one of the options with maximal utility
 (employed by default)
 - satisfying: choose one of the options with sufficient utility
- Decision Matrices
 - Outcome $O_{n \times m} = \text{actions } \{a_1, \dots, a_n\} \times \text{states } \{s_1, \dots, s_m\}$
 - with $\{a_1, a_2\} = \{\text{umbrella, no umbrella}\}; \{s_1, s_2\} = \{\text{rain, no rain}\}$

	rain	no rain
umbrella	dry&heavy	dry&heavy
no umbrella	wet&light	dry&light
 - Utility Assignment
 - requirement: shared perspective
 \Rightarrow all participants share a common concern
 (otherwise, different incomparable utility assigned to the same outcome)
 - assign each outcomes with utility \Rightarrow utility matrices U

	rain	no rain
umbrella	15	15
no umbrella	0	18
- Probability Estimation of Environment
 - Modeling State
 - certainty: deterministic knowledge of the environment after each action
 \Rightarrow thus deterministic outcomes for each action
 - risk: complete probabilistic knowledge of environment
 (may conditioned on action)
 - uncertainty: partial probabilistic knowledge
 - ignorance: no probabilistic knowledge, or not meaningful
 - Objective Probability
 - based on empirical known frequencies
 - indirect estimation (e.g. similar experiment), with calibration
 may use subjective probability, which is unreliable due to psycho-effect
 \Rightarrow better not use subjective estimates of prob
 - Bayesian Description
 - a coherent and complete set of probabilistic beliefs
 (in compliance with laws of probability)
 - probability subjected to the observation
- Action Discovering
 - Category
 - closure: settle down with current available actions

- active postponement: searching for other possible actions (postpone the decision)
 - semi-closure: consider only reversible actions & searching for other actions
- Guideline
 - all current actions have severe drawbacks?
 - is searching possible actions costly?
 - is problem aggregating over time?
 - etc...
- Understanding
 - a meta-decision problem before the scenario-specific decision
- \Rightarrow Assumption
 - Closed Set of Actions
 - no new alternative action allowed to be added
(v.s. open: new actions can be discovered and taken into consideration)
 - Mutually Exclusive Actions
 - no two actions can be both realized
 - Defined States
 - all possible states of environment are recognized
 - Outcome & Utility
 - the joint result of environment (state) and chosen action
 - utility assigned to each outcome (similar to rewards in RL)
(though the utility for each result is usually subject to agent)
 \Rightarrow completeness & transitivity assumed

Decision under Risk

- Maximizing Expected Utility
 - Procedure
 - for action a , expected utility $u = \mathbb{E}_{p(s)} U(a, s) = \sum_{s \in S} p(s) U(a, s)$,
where $U(a, s)$ the utility assigned to outcome $O(a, s)$
 - Objective Utility
 - utility assigned according to the objective rewards (e.g. money)
 - \Rightarrow St, Petersburg paradox: a fair coin tossed, if its head comes up at n^{th} toss, 2^n gold coins are rewarded, stop tossing if its back comes up
 \Rightarrow expected wealth $= \lim_{n \rightarrow \infty} \frac{1}{2} \cdot 1 + \frac{1}{4} \cdot 2 + \dots + \frac{1}{2^n} \cdot 2^{n-1} = +\infty$
 \Rightarrow should play the game for any finite entry fee
 - Subjective Utility
 - in economic: utility of next increment of wealth $\propto \frac{1}{w}$, where w the current wealth
 \Rightarrow utility of wealth $= \log(w)$
 - Understanding
 - objective utility to maintain inter-subjective validity
(hence expert advice remains effective)
 - maximizing utility to maximize long-term outcome
(suitable for large-scale / long-term decision)
 \Rightarrow law of large number to level out randomness

- \Rightarrow suitability depends on the scale (the leveling-out effect)
(e.g. isolated cases may not be suitable to maximize expected utility)
- may not be suitable in extreme effect, to maintain fairness
(e.g. avoid imposing high-probability risk on individuals)
- Variations
 - Conditionalized Expected Utility
 - probability of state $p(s)$ described conditional on action
 $\Rightarrow u = \mathbb{E}_{p(s|a)} U(a, s) = \sum_{s \in S} p(s|a) U(a, s)$
 - model the influence of actions on states
 - Generalized Expected Utility
 - process utility: accounts for attitude towards risk and certainty
 - maximize expected (utility of outcome + process utility)
 - Regret Theory
 - regret: the gap between reward received and highest possible reward from other actions
 - maximize expected utility of outcome & minimize regret
 - Prospect Theory
 - function $f(0, 1) \rightarrow (0, 1)$ to adjust probability
(description to be merely weights, not satisfying probability law) \Rightarrow overweight the probability close to 0&1, as tenuous to certainty
 - maximize weighted utility
- Causal Decision

Decision under Uncertainty

- Estimation of Environment
 - Binary Measure
 - segment out a range for probability describing the state
e.g. $p(s) \in (0.05, 0.2)$
 - Second-Order Probability
 - a Bayesian probability for each state s
 - a further measurement of the reliability of probability estimates $p(s)$
(may consider as a meta subjective probability)
 - Fuzzy Set Membership
 - vagueness, instead of randomness, to describe uncertainty
 \Rightarrow assign a degree of membership for each $p(s)$
 \Rightarrow based on laws of fuzzy membership (non-statistical concept)
 - measure the degree of " $p(s)$ is the probability of state s happening" being true
 - Epistemic Reliability
 - a weight $\in (0, 1)$ as the measure of reliability of a probability, with NO mathematical properties pre-defined
- Decision Criteria
 - Maximin Expected Utility

- maximize the minimal expected utility
⇒ an extremely prudent/pessimistic criterion
- Reliability-weighted Expected Utility
 - use the weighted average probability as true probability
(then reduced to maximizing expected utility)
⇒ an optimistic criterion
- Index
 - p' the best prob estimate and ρ as its degree of confidence
⇒ calculate u_{best} using p' ; u_{min} the minimal expected utility
 - summarize as $u = \rho u_{\text{best}} + (1 - \rho)u_{\text{min}}$
⇒ maximize the u (choose the action with maximal u)
 - $\rho \in (0, 1)$ as a balance factor between pessimism & optimism
- Clipped Maxmin Expected Utility
 - discard any prob with a reliability lower than a threshold
(then max-min expected utility with remaining prob)
- Filtered Maxmin Expected Utility
 - 3 filters applied in sequential order on actions (rather than prob)
 - E -test: action a passes, iff \exists reliable prob $p(s) > 0$ & $\forall a' \in A, U(a, s) \geq U(a', s)$
(action a is possible to be the best)
 - P -test: passes, iff it can be the best in all candidates
 - S -test: passes, iff it is the best under maxmin expected utility in all candidates
⇒ directly produce action to choose

Decision under Ignorance

- Category
 - Known States
 - the prob for each state $p(s) > 0 \Rightarrow$ unknown non-zero prob
 - Unknown States
 - the prob $p(s) \geq 0 \Rightarrow$ unknown prob
- Unknown Non-zero Prob
 - Max-Min Utility
 - maximize the minimal utility (extremely pessimistic)
 - lexicographic maximin: consider their 2nd worst outcome
 - Max-Max Utility
 - maximize the maximal utility (extremely optimistic)
 - Index
 - calculate $u_{\text{max}}, u_{\text{min}}$ by max-max, max-min
 - $u = \alpha u + (1 - \alpha)u_{\text{min}}$
⇒ α as tge degree of optimism
 - Min-Max Regret
 - produce a regret matrix R from utility matrix U ,
where $R(a, s) = \text{maximal possible utility in state } s - U(a, s)$
 - then, minimize the maximal regret, according to R
 - Common Prior

- assign uniform distribution over states, thus reduced to decision under risk
- yet, depends on how the states are partitioned
i.e. the num of states affect the prob
- Unknown Prob
 - Challenge
 - action may lead to catastrophic outcomes (unforeseen outcomes)
⇒ hard to decide if the chance of a severe consequence is negligible or not
⇒ uncertainty towards severe consequence
 - Empirical Guideline
 - choosing / not choosing an action may have different uncertainty (asymmetry of uncertainty)
 - novelty: empirically, novelty brings in more uncertainty
 - spatial and temporal limitations: more limitation, less uncertainty
 - more interference with complex system, the more uncertainty

Social Decision

- Assumption
 - Conflicting Concern
 - participants usually have different/conflicting goals&concerns
- Cyclic Preference
 - No Stable Actions
 - the directed set (actions, preference) forms a directed graph
⇒ \forall action a , \exists action $a', a' > a$
 - Arrow's Theorem
 - 4 rationality criteria are satisfied by decision rule
⇒ then cyclic preference unavoidable
 - understanding: some rationality demands are NOT compatible

Decision for Statistical Model


- Formulation
 - Probability Description
 - known distribution $p(\mathbf{t}|\mathbf{x})$ from model inference, where \mathbf{x} the input, \mathbf{t} the label
⇒ all states (label) estimated with probability
- Decision Criteria
 - Minimizing Misclassification Rate (Maximizing Correct-classification Rate)
 - action: for each class k , assign a region \mathcal{R}_k s.t. if $\mathbf{x} \in \mathcal{R}_k$, predict $\hat{\mathbf{t}} := \mathcal{C}_k$

$$\Rightarrow P(\text{mistake}) = \sum_{k=1}^K P(\mathbf{x} \notin \mathcal{R}_k, \mathbf{t} = \mathcal{C}_k) = 1 - \underbrace{\sum_{k=1}^K P(\mathbf{x} \in \mathcal{R}_k, \mathbf{t} = \mathcal{C}_k)}_{P(\text{correct})}$$

$$= 1 - \sum_k \int_{\mathcal{R}_k} p(\mathcal{C}_k|\mathbf{x})p(\mathbf{x})d\mathbf{x}$$

 ⇒ segment \mathbf{x} to be in region \mathcal{R}_k governed by the maximal $p(\mathcal{C}_k|\mathbf{x})$
 i.e. predict $\hat{\mathbf{t}}$ to be the \mathcal{C}_k with maximal $p(\mathcal{C}_k|\mathbf{x})$

- in binary classification ($K = 2$) with $\mathbf{x} = x$ being a scalar:



where \hat{x} a sub-optimal decision region segment, x_0 the optimal one

- Minimizing Expected Loss (Maximizing Expected Utility)

- loss for a kind of misclassification can vary from other kinds (e.g. recall v.s. precision)
 \Rightarrow a loss function (utility function) $f : (\text{label } \mathbf{t} = \mathcal{C}_k, \text{pred } \mathbf{x} \in \mathcal{R}_j) \rightarrow \text{loss } L_{kj}$
 $\Rightarrow L_{kj}$ as the loss of mis-classify \mathcal{C}_k into \mathcal{C}_j
 $(\forall k, L_{kk} = 0 \text{ \& utility as negative loss})$
- expected loss $\mathbb{E}[L] = \sum_{k,j} L_{kj} P(\mathbf{x} \in \mathcal{R}_j, \mathbf{t} = \mathcal{C}_k) = \sum_{k,j} \int_{\mathcal{R}_j} L_{kj} p(\mathcal{C}_k|\mathbf{x}) p(\mathbf{x}) d\mathbf{x}$
 \Rightarrow expected loss from $\mathcal{R}_j : \sum_k \int_{\mathcal{R}_j} L_{kj} p(\mathcal{C}_k|\mathbf{x}) p(\mathbf{x}) d\mathbf{x}$
 \Rightarrow segment \mathbf{x} into region \mathcal{R}_j governed by the minimal $\sum_k L_{kj} p(\mathcal{C}_k|\mathbf{x})$
i.e. predict $\hat{\mathbf{t}}$ to be the \mathcal{C}_j with minimal $\sum_k L_{kj} p(\mathcal{C}_k|\mathbf{x})$

- Regret Region

- Decision Error Source
 - the largest posterior $p(\mathcal{C}_k|\mathbf{x}) \ll 1$
 - all posteriors are comparable

2.3 Information Theory

2.3.1

2.4 Recommended Practice

2.4.1 Data

Data Augmentation

- Artificial Data Synthesis
 - Practice
 - easily prepare a large amount of similar (yet, different) data
 - add canonical noise to the similar data
 - Understanding
 - convert similar data to the desired distribution
 - Caution

- collected canonical noise may only represent a subset of all possible noise
 - ⇒ may overfit to those collected noise
 - (e.g. distortion on image from game for car detection: too less unique cars)
- Examples
 - in the task of classifying image uploaded by users
 - ⇒ collect web image & blur the image
 - in the task of in-car NLP interaction
 - ⇒ collect well-recorded sentence audio & add in-car noise
- Distortion
 - Practice in Computer Vision
 - mirroring horizontally/vertically, rotation, shirring, etc.
 - random crop a reasonably large subset of image
 - color shifting: e.g. PCA color augmentation
 -

Data Preprocessing


- Mean Centering
 - Practice
 - for all training examples, compute mean (on each features) $\mu = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$,
 where $\{\mathbf{x}_1, \dots, \mathbf{x}_N\} = X_{\text{train}}$ the training set
 - preprocess each $\mathbf{x} \in X_{\text{train}}, X_{\text{val}}, X_{\text{test}}$ to be $\mathbf{x}' = \mathbf{x} - \mu$ (all data go through the same process)
 - Pros
 - (training) data has a zero mean (statistically, most data close to 0)
 - Cons
 - different features may reside in various scales
- Standardizing
 - Practice
 - compute mean μ , standard deviation $\sigma = \left(\frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \mu)^2 \right)^{1/2}$,
 where $\{\mathbf{x}_1, \dots, \mathbf{x}_N\} = X_{\text{train}}$ the training set
 - preprocess each $\mathbf{x} \in X_{\text{train}}, X_{\text{val}}, X_{\text{test}}$ to be $\mathbf{x}' = \frac{\mathbf{x} - \mu}{\sigma}$
 (all data go through the same process)
 - note: with big data, usually computed iteratively due to limited memory
 - Pros
 - (training) data has zero mean & unit variance
 - ⇒ approximated to normal distribution
 - for deep learning: different features in same small range close to 0
 - ⇒ weights for different features are in roughly the same scale
 - ⇒ easier to train
- Cleaning Incorrect Label

- Practice
 - before cleaning: measure its contribution to the error rate & its cause
 - random error (e.g. occasional mistake, etc.) with a big dataset: okay to ignore
 - systematic error: should be corrected, at least for val&test set
⇒ to evaluate the model on the target data distribution
 - if mislabeled data cause inability to evaluate&compare model: must be cleaned
 - val&test set should be cleaned together
⇒ to have the same distribution
- Understanding
 - in train set: statistic model (deep net etc.) quite robust to random errors while model can learn the systematic error ⇒ not able to generalize
 - in val set: random error can cause inability
⇒

2.4.2 Dataset

Train-Val-Test

- Reason
 - Iterative Process



- intuition usually do NOT transfer across domains (NLP, CV, Search, etc.)
 - do NOT hope to have the correct hyperparameters at the first try
- ⇒ need feedback from experiment result
- ⇒ make sure the feedback is CORRECT and FAST
- Recommended Usage
 - Splitting
 - classic split for small dataset ⇒ train:val:test = 60 : 20 : 20, or K-fold
 - in big data (e.g. 100 million) ⇒ train:val:test = 98 : 1 : 1
(as long as val-test sets cover enough data variance)
 - Training Set
 - to find the model parameter estimation (used for learning process of model)
⇒ over-fit by complex model
 - can incorporate methods to train the model to have desired property
(where augment data goes)
 - Validation Set (Val)
 - to indicate generalization ability of a range of trained models on target data
(correct if enough various input covered)
⇒ for model comparison, selection & hyperparameters tuning

- should have consistent distribution with test set
(as val set is also evaluating the generalization ability)
- Test Set
 - to evaluate the **generalization ability** of final model on target data
(correct if enough various input covered)
 - should represent the distribution of target data
i.e. data that the deployed model will need to handle
- Training-Validation Set (Train-Val)
 - another val set split from original training set
 - used if training set are from different distribution then the val/test set
(e.g. due to augmented data etc.)
 - performance gap between train&val set: variance + distribution mismatch
⇒ separate each measurement
 - performance gap between train&train-val set: measuring variance
⇒ performance gap between train-val&val set: measuring distribution mismatch
- Potential Problem
 - Mismatched Distribution across Sets
 - classic supervised learning assumption: all sets drawn from SAME distribution
(yet transfer/adaptive learning focus on violation of such assumption)
 - measured by train-val set
 - should ensure at least that val&test set have the SAME distribution
 ⇒ yet, make sure val&test set from the SAME distribution as the desired one
 - **Overfitting Val Set**
 - iteratively tuning model is a processing of learning (fitting to the val set)
⇒ with enough iteration, val set can be overfit
 - may consider test set as 2nd val set, and further have 3rd, 4th... val sets
 - Limited Data
 - better model ⇒ more training data
 - ⇒ less validation ⇒ noisy estimation of generalization ability

Train-Test

- No Val Set
 - Practice
 - may use the "test" set as val set ⇒ generalization ability NOT reported
 - should be confident in that dataset cover/represent true distribution of data
(yet, not recommended)
 - Understanding
 - to utilize as many data as possible for ultimate performance
- K-fold Cross Validation
 - Procedure
 - split all data into K folds, $K - 1$ folds for train, 1 for validation
 - ⇒ average over all C_K^1 combination to indicate the generalization ability
 - extreme case: leave-out-one ⇒ $K = N$, where N is number of all data
 - Cons:
 - $\mathcal{O}(K)$ ⇒ slow, especially if training process already slow
⇒ trade off between time vs. constraint on validation
 - hence, **not** often used in big data era

2.4.3 Orthogonalization Practice

Definition

- Decoupling Goals and Models
 - Designing Metrics
 - to evaluate the models
 - ⇒ capture how well the problem solved as desired
 - decouple different aspect of concern into different metrics
 - Designing Models
 - to do well on the previously chosen metrics
 - (including training & tuning hyperparameters)
- Decoupling Hyperparameters
 - disjoint set of hyperparameters to optimize for train-val-test set
 - hyperparameter taking effect on single goal
 - at least, NOT to impose negatively related effect on multiple goals
 - (e.g. early stopping on performances on train and val sets ⇒ NOT preferred)
 - ⇒ clearer control on model behaviors

Practice

- Designing Metrics (Goals)
 - Single Metric Reporting Overall Performance
 - a metric accounting for multiple metrics
 - e.g. F1 score instead of precision and recall
 - weighted average over metrics
 - (capturing tendency by different weights)
 - Satisficing Metrics
 - optimizing single metric with some minimum requirement must being satisfied
 - ⇒ single optimizing metric + several satisficing metrics
 - ⇒ optimizing under constraints
 - e.g. optimizing accuracy with false positive rate < 0.2 satisfied
- Tuning Hyperparameters (Designing Model)
 - Inherent Separation for Set-level Goals
 - fit model on train set for good fitting
 - ⇒ tune model/network structure, optimization, preprocessing, etc.
 - evaluate on val set for good generalization ability
 - ⇒ tune regularization, etc.
 - evaluate on test set for hopefully good generalization ability reported
 - ⇒ consider bigger val set (if an overfit val set indicated)
 - apply in real world hopping model to generalize well indeed
 - ⇒ consider mismatched data distribution, redesign cost function / metrics etc.
 - (if failed to generalize)
 - note: size of dataset can be hyperparameter sometimes
 - Separating Tuning for Performance Metrics
 -

Orthogonalization

- Definition
 - Decoupling Tuning for Different Goals
 - disjoint set of hyperparameters to optimize on train-val-test set
 - hyperparameter taking effect on single goal
(at least, NOT to impose negatively related effect on multiple goals)
 - Single Metric Reporting Performance
 - a metric accounting for multiple metrics
e.g. F1 score instead of precision and recall
 - optimizing under constraints (must-satisfied metrics)
e.g. optimizing accuracy with false positive rate < 0.2 satisfied
- Practice
 - Separating Tuning for Set-level Goals
 - for train set: model/network structure, optimization, preprocessing, etc.
 - for val set: regularization, etc.
 - for test set: bigger val set (as indicating an overfit val set)
 - for real world: mismatched cost function / data distribution, etc.
 - Separating Tuning for Performance Metrics
 -
- Understanding
 - Inherently Separated Goals
 - fit model on train set: adjust model for good fitting
 - evaluate on val set: adjust model for good result on metrics
(indicating generalization ability)
 - evaluate on test set: hope to report good generalization ability
 - apply in real world: hope to generalize well indeed
 - Clear Control on Behaviors
 - form iterative process among various goals
 - prevent tuning practice with unaware negatively coupled effect on different goals
(e.g. early stopping on performances on train and val sets \Rightarrow NOT preferred)

2.4.4 Tuning Hyperparameters

especially for deep learning

Hyperparameters

- Overview
 - Structures and Architectures
 - type of layers and size of layers
 - type of activation
 - depth of networks
 - Learning
 - learning rate

- optimizer (learning process)
- Robustness and Generalizability
 - regularization(s)
 - data preprocessing/augmentation
- Challenges
 - NO Consistent Prescience
 - popular choices from one domain usually NOT carry over to other domains
 - NOT Predictable Effect
 - hyperparameter does NOT have predictable effect on specific model behavior
⇒ need a search for the best one

Systematic Searching

- Random Sampling
 - Reason
 - NOT able to know the importance of different hyperparameters
⇒ not wasting grid search step on the unimportant
 - NOT able to know the effective range of a hyperparameters
⇒ may be skipped by grid search step
 - decouple the search for different hyperparameters ⇒ more richly explore
(whereas grid search fix one while searching on others)
 - Coarse to Fine Scheme
 - explore whole space uniformly (equally random)
 - exploit region where good results show up (with more densely sampled)
 - Sampling on Scale
 - instead of sampling the value of hyperparameter, sample the scale of it
e.g. sampling learning rate $\alpha = 10^r, r \sim U(-4, 0)$
 - ⇒ distribute the density across desired scales
(by using transfered scale, e.g. applying \log, e^x e.t.c)
 - reason: depends on the
 - use of hyperparameter e.g. in an exponential/linear/log way
 - whether intend to sample on scale e.g. across one or more scales
- Swarm Optimization
 - Intuition
 - searching over a space with continuous×discrete across various scale
⇒ encoded into a list
 - search using permutation / group behavior
⇒ inherently imposing explore-exploit strategy
 - Popular Framework
 - genetic algorithm (GA)
 - particle swarm optimization (PSO)

Tunning Practice

- Single Model
 - Practice
 - supervising one model at a time
 - interactively justify the hyperparameter in training process
 - ⇒ gain knowledge through interaction & ensure a good performance
 - ⇒ early feedback
 - Reason
 - too many data (online advertisement, computer vision etc.)
 - few computing resource
- Parallel Training
 - Practice
 - shoot out multiple model with various settings
 - compare at the end (after trained & evaluated)
 - Reason
 - small data/model, enough computability / fast training process

2.5 Model Analysis

2.5.1 Measurements of Problem

Performance Metrics

- Intersection over Union (IoU)
 - Input
 - two regions, from prediction and label
e.g. bounding box, segmentation mask
 - Definition
 - I = intersection of the regions
 - U = union of the regions
 - $\Rightarrow \text{IoU} = \frac{I}{U}$
 - Use Case
 - evaluation of object detection/segmentation
 - post-processing in object detection/segmentation
(e.g. non-max suppress)
 - Understanding
 - measure how well two regions overlap
⇒ usually $\text{IoU} > 0.5$ considered a decent match
 - average IoU over each prediction class & scale of threshold for an overview report
(e.g. average over all obstacle types, threshold range [0.05 0.95] with step 0.05)
 - Extension
 - intersection over label/prediction region \Rightarrow in case of unstable label/prediction region
(while matching relation considered important in the task)

- Bilingual Evaluation Understudy (BLEU)
 - Input
 - one sequence as prediction; other sequence(s) as references
(can be more than one) i.e. label
 - Definition
 - n -gram g_n : n continuous tokens
 - count of g_n : the number of appearance of g_n in a sequence
 \Rightarrow num of $g_n = \sum_{g_n} (\text{count of } g_n) = l - n + 1$, in sequence of len l
 - ref count: $\sum_{g_n \in \text{pred}} (1 \text{ if it appears in any reference; else } 0)$
 \Rightarrow count of g_n in prediction that also appears in a reference
 $\Rightarrow g_n \in \text{pred}$ matched as long as $g_n \in$ a reference
 - \Rightarrow precision = $\frac{\text{true positive}}{\text{positive pred}} = \frac{\text{ref count}}{\text{num of } g_n \in \text{pred}}$
 (yet, "the the the the" has high score when $n = 1$, as "the" almost always appear)
 - clipped ref count: $\min\{\text{count of } g_n \text{ in pred}, \max(\text{num of } g_n \text{ in a reference})\}$
 \Rightarrow count only unique n -gram in pred
 $\Rightarrow g_n \in \text{pred}$ matched up to the max num of $g_n \in$ reference
 - \Rightarrow clipped precision $p_n = \frac{\text{clipped true positive}}{\text{positive pred}} = \frac{\text{clipped ref count}}{\text{num of } g_n \in \text{pred}}$
 - \Rightarrow BLUE score = $BP \exp\left(\frac{1}{N} \sum_{n=1}^N p_n\right)$,
 where $BP = 1$ if $\text{len}_{\text{pred}} > \text{len}_{\text{ref}}$; else $\exp(1 - \text{len}_{\text{pred}}/\text{len}_{\text{ref}})$
 \Rightarrow the brevity penalty to penalize too short pred
 (as short pred has larger chance to have all its gram contained in ref)
 - Use Case
 - machine translation
 - image caption
 - Understanding
 - clipped ref count: has a maximal number for g_n appearance, given the reference
 \Rightarrow same g_n exceeding the number becomes false positive, as can NOT be matched
 (analogy: metrics in obj detection with bounding box)
 - evaluate the appearance of generated tokens (prediction) in references (label)
 - highly correlated with human evaluation
 - bias towards statistical model (when compared against rule-based model)

Expected Generalization Ability Measurement

- Bayes Optimal Performance
 - Measurement
 - the theoretically best performance on all data
 \Rightarrow modeling only the indent mapping without the noise (a perfect model)
 (note: in practice, only approximated bayes performance available)
 - Understanding
 - measure the inherent noise in data as the best possible performance on all data
 - \Rightarrow measure **avoidable** bias: indicate the upper-bound performance

- \Rightarrow measure degree of **overfitting**: indicate how much the model fit to the noise
- Human Performance: Approximation to Bayes Performance
 - Definition
 - for best approximation: best achievable performance by human
(e.g. group of experts, as bayes optimal performance is even better)
 - for specific focus: depends on use case
e.g. for self-diagnose model, may define as the performance of a normal doctor
 - Practice
 - usually done in supervised learning
note: the label (i.e. 0-error) is NOT bayes performance
 - on unstructured data, human almost achieves bayes performance
(as human good at natural perception task, like cv, nlp)
 - Cons
 - hard to distinguish surpassing human performance from overfitting training set
- (Avoidable) Bias
 - Measurement
 - gap between train set performance and bayes performance
(note: in practice, only approximated bayes performance available)
 - Understanding
 - measure the model capacity of handling given (train) data
as (approximately) measuring the gap between the theoretically best model
- Variance
 - Measurement
 - performance gap between val set & train set (if under same distribution)
 - if different distribution for train&val set \Rightarrow train-val set instead of val set
 - Understanding
 - measure the generalization ability:
as measuring how much model can cope with unseen data
 \Rightarrow model the indent mapping, instead of the noise
- Mismatch Distribution
 - Measurement
 - performance gap between train-val set & val set
 - Understanding
 - train-val set contains the unseen train data; val set the unseen target data
 \Rightarrow gap only caused by different distribution between sets
- 2. Interaction with regularization:
 - Improper λ : - large $\lambda =$ high bias - small $\lambda =$ high variance - Choosing λ : - try $\lambda = 0, 0.01, 0.02, 0.04, \dots, 10$ - select the model with lowest $J_{cv}(\theta)$ without regularization term
- 3. Interaction with training set size:
 - Normal Learning curve:
 - ! [Normal learning curve] (../Machine)
 - Learning curve with high bias:
 - where getting more training data ****doesn't**** help

! [Learning curve with high bias] (./../Machine
 - Learning curve with high variance:
 - where getting more training data ****helps****
 ! [Learning curve with high variance] (./../Machine
 4. Ways to fix:
 - High bias: - more features / more polynomial terms of features - decreasing λ
 - High variance:
 - larger data set / fewer features - increasing λ
 - ****In neural network:****
 - High bias = \downarrow larger neural networks (more hidden layers / more units in one layer)
 - High variance = \downarrow smaller neural networks
****Larger network with regularization (λ) is more powerful****

2.5.2 Improving Model

Bias-Variance Guideline

- Solving High (Avoidable) Bias
 - Increasing Model Capability
 - increase complexity: more weights, latent variable / hidden layer etc.
 - use more suitable model specifically designed for the data (e.g. CNN for image)
- \Rightarrow until fitting training set well
- Solving High Variance
 - Data Augment
 - get/simulate more training data (via crowd sourcing, distortion, GANs, etc.)
 - Model Regularization
 - control the complexity of model (e.g. L0/1/2 normalization)
- Solving Trade-off
 - Iterative Process
 - solve bias, then solve variance, iteratively
 - Complexity + Data/Regularization
 - increase complexity to solve bias
without hurting variance (via more data/regularization)
 - more data/regularization to solve variance
without hurting bias (with enough complexity)

Behavior Detail

- Low Bias, High Variance (Over-fitting)
 - Symptom
 - good performance on train set & poor generalization (bad on val)
 - \Rightarrow good at fitting train set; bad at representing/modeling underlying data source
 - Cause
 - too much representation ability (to fit even the noise)
 - directly model the likelihood instead of posterior

- Remedy
 - larger dataset
 - regularization (model the posterior by accounting prior)
- High Bias, Low Variance (Under-fitting)
 - Symptom
 - bad at fitting training examples & modeling underlying data source (bad at train & val)
 - \Rightarrow poor performance on train set & good generalization (though meaningless)
 - Cause
 - lack of representation ability (not enough flexibility)
 - Remedy
 - try model with better representative ability (more complexity, flexibility)
- High Bias, High Variance (Over&Under-fitting)
 - Symptom
 - bad at fitting some general cases; while good at some rare and special cases (especially in high dimensional space)
 - \Rightarrow fitting largely noise
 - Cause
 - model probably not suitable for the dataset
 - Remedy
 - switch to other types of model
 - dataset preprocessing
- Low Variance, High Mismatch
 - Symptom
 - good at generalizing (on train-val); bad at target data (on val)
 - Cause
 - model not able to generalize across mismatch distribution (yet generalized well in the same distribution: as good at train-val)
 - Remedy
 - transform train data towards (more like) target data
e.g. data synthesis: adding noise that is special in target data, etc.
 - ensure train set contains enough / assign larger weight to, the desired target data
 - transfer learning, adaptive learning, etc.
- Low Bias, Low Variance, Low Mismatch, High val-test Variance
 - Symptom
 - model with specific hyperparameter overfitting the val set
 \Rightarrow val set NO longer reveal model generalizability
 - Cause
 - val set overfit by iteration of hyperparameter tuning (as a practice of fitting)
 - Remedy
 - more data for val&test set

- re-design/choose the model after val set overfitting fixed
(after true generalizability reported)
- Low Bias, Low Variance
 - Behavior
 - good at fitting training examples & modeling underlying data source
(good at train & val)
 - \Rightarrow good performance on training set & good generalization
- High Bias, Low Variance, Lower/Negative Mismatch
 - Behavior
 - perform better on val& test set then on train set
 \Rightarrow target data distribution easier than train set distribution
 \Rightarrow able to do well on desired data, even if not good on train set
 (better convinced by measuring human performance on both distribution)

Error Analysis

- Categorizing Error Source
 - Practice
 - create histogram on val set reflecting data categories
(e.g. for image: blurry, rotated, incorrect label, etc...)
 \Rightarrow categorize data first
 - \Rightarrow data leading to error scattered into different categories
 \Rightarrow evaluate the contribution to error from different categories
 - note:
 - Understanding
 - find out the most important error source
 \Rightarrow prioritize the direction of tuning model
- Ceiling Analysis
 - Definition & Practice
 -
 - Understanding

Approaches Analysis

- End-to-End Approach
 - Definition
 - use single network to learn the mapping from input directly to desired output
(no intermediate result)
 - Pros
 - reveal the data statics: avoid any specific prior
 - large & auto feature extraction
 - Cons
 - need enough data for effective end-to-end model
 - hard to inject effective prior into model
 \Rightarrow exclude potentially hand-designed component/knowledge

- Regularization: add penalty for θ being large into cost function
- $J(\theta) = \dots + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$, **bias θ_0 shouldn't be penalized**

2.7 Linear Regression

- Notation
 - t : observed data
 - $y(\mathbf{x}, \mathbf{w}) = \sum_{i=0}^M \phi_i(\mathbf{x}) w_i = \mathbf{w}^T \phi(\mathbf{x})$: model generating ground truth, with
 - \mathbf{w} : weight vector
 - $\phi(\mathbf{x})$: basis function for feature vector \mathbf{x} , with usually $\phi_0(\mathbf{x}) = 1$ as bias
- Assumption
 - Deterministic Model with Observation Noise
 - $t = y(\mathbf{x}, \mathbf{w}) + \epsilon$, where $\epsilon \sim \mathcal{N}(0, \beta^{-1})$ is Gaussian noise where precision (inverse variance) β
 - \Rightarrow consequence
 1. likelihood $p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1})$
 2. $\mathbb{E}[t|\mathbf{x}] = \int t \cdot p(t|\mathbf{x}) dt = y(\mathbf{x}, \mathbf{w})$
 3. unimodal distribution $p(t|\mathbf{x}) \Rightarrow$ extended by conditional mixture model
- Joint Likelihood
 - $P(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n|\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$, where
 - $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}, \mathbf{t} = \{t_1, \dots, t_N\}$
 - Log Likelihood
 - $\ln P(\mathbf{y}|\mathbf{X}, \theta, \beta) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta \frac{1}{2} \sum_{i=1}^m (h_\theta(x^i) - y^i)^2$
- Log Posterior leads to regularization
 - Maximizing the likelihood function \Rightarrow (often) excessively complex models & overfitting
 - Regularization term comes from the Prior:
 - assume Prior $p(\theta) = \mathcal{N}(\theta|0, \alpha^{-1}I)$, so that Posterior & Prior are of the same distribution to maximize log Posterior :

$$\Rightarrow \ln p(\theta|X) \propto -\frac{\beta}{2} \sum_{i=1}^n (y^i - h_\theta(x))^2 - \frac{\alpha}{2} \theta^T \theta + \text{const}$$
 - If $\alpha \rightarrow 0$ (Prior is most useless), maximise Posterior is equivalent to maximizing likelihood
 - Maximize Posterior \Leftrightarrow Minimize cost function with regularization, where $\lambda = \alpha/\beta$
- Predictive Distribution: $p(y|x, X, Y)$

- $p(y|x, X, Y) = \int p(y, \theta|x, X, Y)d\theta = \int p(y|\theta, x, X, Y)p(\theta|x, X, Y)d\theta$
- $p(y|\theta, x, X, Y) = p(y|\theta, x) = \mathcal{N}(y|h(x, \theta), \beta^{-1})$
 based on assumption: $y = y(x, \theta) + \epsilon$, where ϵ is Gaussian noise
 $p(\theta|x, X, Y) = p(\theta|X, Y) = \text{posterior}$
- $\Rightarrow p(y|x, X, Y) = \int p(y|\theta, x)p(\theta|X, Y)d\theta$
- Expected Lost = $(\text{bias})^2 + \text{variance} + \text{noise}$
- Notation:
 - $t = y(x, w) + \epsilon$, where ϵ is Gaussian noise
 - \hat{y} is prediction function to approximate $y = y(x, w)$
- Procedure:
 - $\mathbb{E}[(t - \hat{y})^2] = \mathbb{E}[t^2 - 2t\hat{y} + \hat{y}^2]$
 $= \mathbb{E}[t^2] + \mathbb{E}[\hat{y}^2] - \mathbb{E}[2t\hat{y}]$
 $= \text{Var}[t] + \mathbb{E}[t]^2 + \text{Var}[\hat{y}] + \mathbb{E}[\hat{y}]^2 - 2y\mathbb{E}[\hat{y}]$
 $= \text{Var}[t] + \text{Var}[\hat{y}] + (y^2 - 2y\mathbb{E}[\hat{y}] + \mathbb{E}[\hat{y}]^2)$
 $= \text{Var}[t] + \text{Var}[\hat{y}] + (y - \mathbb{E}[\hat{y}])^2$
 $= \text{Var}[t] + \text{Var}[\hat{y}] + \mathbb{E}[t - \hat{y}]^2$
 $= \sigma^2 + \text{Var}[\hat{y}] + \text{Bias}[\hat{y}]^2$
 where $\sigma^2 = \text{Var}[\epsilon]$ is the noise
 (formula used: $\text{Var}[x] = \mathbb{E}[x^2] - \mathbb{E}[x]^2 \Leftrightarrow \mathbb{E}[x^2] = \text{Var}[x] + \mathbb{E}[x]^2$)
 - Matrix inverse can be evil & avoid inverse operation:
 $A = U\Lambda U^T$, where Λ is diagonal matrix
 $\Rightarrow A^{-1} = U\Lambda^{-1}U^T$
 but number on the diagonal line of Λ can be small = maybe 0 depending on accuracy of computer

2.8 Bayesian Regression

- Assumption:
 - $t = y(x, w) + \epsilon$, where ϵ is Gaussian noise; $y(x, w)$ approximated by $\phi(x)w$
- Bayesian view:
- Gaussian Prior : $p(w) = \mathcal{N}(w|m_0, S_0)$
 Reason: to be conjugate
- Likelihood : $p(\mathbf{t}|w) = \prod_{n=1}^N \mathcal{N}(t_n|w^T \phi(x_n), \beta^{-1}) = \mathcal{N}(\mathbf{t}|\Phi w, \beta^{-1}I)$
- \Rightarrow Posterior : $p(w|\mathbf{t}) = \mathcal{N}(w|m_N, S_N)$
 where $m_N = S_N^{-1}(S_0^{-1}m_0 + \beta\Phi^T\mathbf{t})$, $S_N^{-1} = S_0^{-1} + \beta\Phi^T\Phi$
- Maximum Likelihood:
 - Likelihood : $p(\mathbf{t}|w) = \prod_{n=1}^N \mathcal{N}(t_n|\phi(x_n)w, \beta^{-1})$
 - meaning: how probable the observed dataset is w.r.t the model setting (under parameter w)

$$\begin{aligned}
\circ \ln \text{Likelihood} &= \sum_{n=1}^N \left[-\ln \frac{\beta}{\sqrt{2\pi}} - \frac{\beta}{2} (t_n - \phi(x)w)^2 \right] \\
\circ \frac{\partial}{\partial w} \ln \text{Likelihood} &= \beta \Phi^T (\mathbf{t} - \Phi w) \\
\text{let } \frac{\partial}{\partial w} \ln \text{Likelihood} &= 0 \\
\Rightarrow w_{ML} &= (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \\
\circ \frac{\partial}{\partial \beta} \ln \text{Likelihood} &= -N\beta^{-\frac{1}{2}} + \beta^{-\frac{3}{2}} (\mathbf{t} - \Phi w)^T (\mathbf{t} - \Phi w) \\
\text{let } \frac{\partial}{\partial \beta} \ln \text{Likelihood} &= 0 \\
\Rightarrow \beta^{-1} &= \frac{1}{N} (\mathbf{t} - \Phi w)^T (\mathbf{t} - \Phi w) \\
\text{Note: solve } w &= w_{ML} \text{ first}
\end{aligned}$$

- Maximum Posterior:

$$\begin{aligned}
\circ \text{Posterior} &= p(w|\mathbf{t}), \text{Prior} = p(w), \text{Likelihood} = p(\mathbf{t}|w), \text{Normalization} = p(\mathbf{t}) \\
\Rightarrow \text{Posterior} &= \frac{\text{Likelihood} * \text{Prior}}{\text{Normalization}} \\
\Rightarrow \text{Posterior} &\propto \text{Likelihood} * \text{Prior} \\
\circ \text{assume Prior } p(w) &= \mathcal{N}(w|m_0, S_0), \\
\text{so that Prior \& Likelihood are conjugate} &\Rightarrow \text{Gaussian Posterior} \\
\circ \text{Likelihood } p(\mathbf{t}|w) &= \prod_{n=1}^N \mathcal{N}(t_n|\phi(x_n)w, \beta^{-1}) = \mathcal{N}(\mathbf{t}|\Phi w, \beta^{-1}I) \\
\circ \Rightarrow \text{Posterior } p(w|\mathbf{t}) &= \mathcal{N}(w|m_N, S_N), \\
\text{where } m_N &= S_N(S_0^{-1}m_0 + \beta\Phi^T\mathbf{t}), S_N^{-1} = S_0^{-1} + \beta\Phi^T\Phi \\
\Rightarrow w_{MAP} &= \text{mean of the Gaussian} = m_N \\
\text{Note: can also get } w_{MAP} &\text{ from taking gradient}
\end{aligned}$$

- Simple Prior:

$$\begin{aligned}
\text{Prior } p(w) &= \mathcal{N}(w|0, \alpha^{-1}I) \\
\Rightarrow \text{Posterior } p(w|\mathbf{t}) &= \mathcal{N}(w|m_N, S_N), \\
\text{where } m_N &= \beta(\alpha I + \beta\Phi^T\Phi)^{-1}\Phi^T\mathbf{t}, S_N^{-1} = \alpha I + \beta\Phi^T\Phi \\
w_{MAP} &\rightarrow w_{ML}, \text{ when } \alpha \rightarrow 0 \text{ (most useless Prior)}
\end{aligned}$$

- Maximize Posterior \Leftrightarrow Minimize cost function with regularization:

$$\text{Simple Prior} \Rightarrow \ln p(w|\mathbf{t}) = -\frac{\beta}{2} (\mathbf{t} - \Phi w)^T (\mathbf{t} - \Phi w) - \frac{\alpha}{2} w^T w + \text{const}$$

- If $\alpha \rightarrow 0$ (Prior is most useless), maximize Posterior is equivalent to maximizing likelihood

- Maximize Posterior equal to minimize sum-of-squares error function with the addition of a quadratic regularization term with $\lambda = \alpha/\beta$

- Regularization term comes from the Prior

- Predictive Distribution:

- Assume: Prior : $p(x|\alpha) = \mathcal{N}(x|0, \alpha^{-1}I)$, ($m_0 = 0, S_0 = \alpha^{-1}I$)

- $p(t|x, X, \mathbf{t}) = \int p(t|w, x)p(w|X, \mathbf{t})dw$

- $\Rightarrow p(t|x, X, \mathbf{t}) = \mathcal{N}(t|m_N^T \phi(x), \sigma_N^2(x))$
where $\sigma_N^2(x) = \frac{1}{\beta} + \phi(x)^T S_N \phi(x)$; m_N, S_N from Posterior($m_N = w_{MAP}$)
- Sequential data:
 - Posterior from previous data \Leftrightarrow the Prior for the arriving data
 - Sequential data and data in one go is equivalent when finding the Posterior
- Gradient descent
 - Hypothesis function:
 - $h_\theta(x) = x\theta, \theta = [\theta_0, \theta_1, \dots, \theta_n]^T, x = [x_0, x_1, \dots, x_n], x_0 = 1$
 - x is one instance
 - Cost function: $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^i) - y^i)^2 + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$
 - Update rule: $\forall \theta_j \in \theta, \theta_j := \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}, \frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m [(h_\theta(x^i) - y^i)x_j^i] + \frac{\lambda}{m} \theta_j -$
 $\frac{d}{d\theta} J(\theta) = \frac{1}{m} ((X\theta - y)^T X)^T + \frac{\lambda}{m} [0, \theta_1, \dots, \theta_m]^T$ (θ_0 shouldn't be penalized)
 - simultaneously for all $\theta_j \in \theta$
 - Normal equation (mathematical solution)
 - $\theta = (X^T X)^{-1} X^T y$

2.9 Logistic Regression (Classification)

- Decision Theory:

- classes C_1, \dots, C_K , decision regions $\mathcal{R}_1, \dots, \mathcal{R}_K$ - Minimize $p(\text{mistake}) = \sum_{k=1}^K (\int_{\mathcal{R}_k} \sum_{i \neq k} p(x, C_i) dx)$

(can have weight on each misclassification though) - Maximize $p(\text{correct}) = \sum_{k=1}^K \int_{\mathcal{R}_k} p(x, C_k) dx$

- Models for Decision Problems:

- Find a discriminant function - Discriminative Models: less powerful, yet less parameter = easier to learn

- Infer **posterior** $p(C_k|x)$, $C_k: x \in C_k, x$ is examples in training set - Use decision theory to

assign a new x - Generative Models: more powerful, but computationally expensive - Infer con-

ditional probabilities $p(x|C_k)$ - Infer prior $p(C_k)$ - Find **either** **posterior** $p(C_k|x)$, **or**

joint distribution $p(x, C_k)$ (using Bayes' theorem) - Use decision theory to assign a new x

- **= Able to create synthetic data using $p(x)$ **

- Naive Bayes on Discrete Features:

- Assumption:

- Discrete Features: data point $x \in \{0, 1\}^D$

- Naive Bayes: all features conditioned on class C_k are independent with each other

$$\Rightarrow p(x|C_k) = \prod_{i=1}^D \mu_{ki}^{x_i} (1 - \mu_{ki})^{1-x_i}$$

1. Linear Discriminant (Least Squares Approach)

- Prediction:

- $y(x) = xw + w_0$, with bias $= w_0$, where $w = [w_1, \dots, w_n]^T, x = [x_1, \dots, x_n]$ - $y(x) > 0$:

positive confidence to assign x to current class - $-w_0$ called threshold sometimes

- Decision Boundary $y(x) = w^T x + w_0 = 0$:

- w is orthogonal to vectors on the boundary:

assume x_1, x_2 on the boundary

$$\Rightarrow 0 = y(x_1) - y(x_2) = (x_1 - x_2)w$$

- Distance from origin to boundary is $-\frac{w_0}{\|w\|}$:

assume distance is k

$$\Rightarrow k \frac{w}{\|w\|} \text{ on boundary, thus } k \frac{w}{\|w\|} w + w_0 = 0$$

$$\Rightarrow k = -\frac{w_0}{\|w\|}$$

- $y(x)$ is a signed measure of distance from point x to boundary:

assume distance is r

$$\Rightarrow y(x) = \overbrace{(x_\perp + r \frac{w}{\|w\|})}^x w + w_0 = \overbrace{x_\perp w + w_0}^0 + r\|w\| = r\|w\|$$

$$\Rightarrow r = \frac{y(x)}{\|w\|}$$

- Multi-class (k-classes):

- prediction: x is of class C_k if $\forall j \neq k, y_k(x) > y_j(x)$

$\Rightarrow y(x) = xW$, where $W = [w_1, \dots, w_k], \forall x_i \in X, x_{i0} = 1$ (bias), $y(x)$ is 1-of-k coding

- sum-of-squares error: $E_D(W) = \frac{1}{2} \text{tr}\{(XW - T)(XW - T)^T\}$

\Rightarrow optimal $W = (X^T X)^{-1} X^T T$

note that $\text{tr}\{AB\} = A^T B^T$

2. Fisher's Linear Discriminant

- Basic idea:

- Take linear classification $y = w^T x$ as dimensionality reduction (projection onto 1-D) - $=_i$ find a projection (denoted by vector w) which maximally preserves the class separation - $=_i$ if $y > -w_0$ then class C_1 , otherwise C_2

- Goal:

- Distance within one class is small - Distance between classes is large

- Mean & Variance of Projected Data:

- Mean: $\tilde{m}_k = w^T m_k$, where $m_k = \frac{1}{N_k} \sum_{x \in C_k} x$ - Variance: $\tilde{s}_k = \sum_{x \in C_k} (w^T x - w^T m_k)^2 =$

$$w^T \left[\sum_{x \in C_k} (x - m_k)(x - m_k)^T \right] w$$

- 2-Classes to 1-D line:

- Maximize Fisher criterion: $J(w) = \frac{|\tilde{m}_1 - \tilde{m}_2|^2}{\tilde{s}_1^2 + \tilde{s}_2^2}$

- Between-class covariance: $S_B = (m_1 - m_2)(m_1 - m_2)^T$

- Within-class covariance: $S_k = \sum_{n \in C_k} (x_n - m_k)(x_n - m_k)^T$

$$\Rightarrow J(w) = \frac{w^T S_B w}{w^T S_W w}$$

- Lagrangian: $L(w, \lambda) = w^T S_B w + \lambda(1 - w^T S_W w)$

fix $w^T S_W w$ to 1 to avoid infinite solution (any multiple of a solution is a solution)

$$\Rightarrow \frac{\partial}{\partial w} L = 2S_B w - 2\lambda S_W w = 0$$

$$\Rightarrow S_B w = \lambda S_W w$$

$$\Rightarrow (S_W^{-1} S_B) w = \lambda w$$

To maximize $J(w)$, w is the largest eigenvector of $S_W^{-1} S_B$ if S_W invertible

- K-classes to a d-D subspace: N_k is num in class k, N is the total example num

- Between-class covariance: $S_B = \sum_{k=1}^K N_k (m_k - m)(m_k - m)^T$, where $m = \frac{1}{N} \sum_{n=1}^N x_n$

reduce to $(m_1 - m_2)(m_1 - m_2)^T$ when K=2 (constant ignored)

- Within-class covariance: $S_W = \sum_{k=1}^K S_k$, where $S_k = \sum_{n \in C_k} (x_n - m_k)(x_n - m_k)^T$, $m_k =$

$$\frac{1}{N_k} \sum_{n \in C_k} x_n$$

- Maximize Fisher criterion: $J(w) = \frac{\text{tr}\{W^T S_B W\}}{\text{tr}\{W^T S_W W\}}$

- Lagrangian:

Solve for each $w_i \in W \Rightarrow (S_W^{-1} S_B) w_i = \lambda_i w_i$

$\Rightarrow W$ consists of the largest d eigenvectors

$S_W^{-1} S_B$ is not guaranteed to be symmetric $\Rightarrow W$ might not be orthogonal

Need to minimize the whole covariance matrix ($J(w)$ as a matrix) \Rightarrow not choosing same eigenvectors twice

- Maximum Possible Projection Directions = $K - 1$:

$$r(S_W^{-1} S_B) \leq \min(r(S_W^{-1}), r(S_B)) \leq r(S_B)$$

$$r(S_B) \leq \sum_{k=1}^K r((m_k - m)(m_k - m)^T) = K, \text{ as } r(m_k - m) = 1$$

$$\sum_{k=1}^K m_k = m \Rightarrow r(m_1 - m, \dots, m_K - m) = K - 1$$

$$\Rightarrow r(S_B) \leq K - 1$$

$$\Rightarrow r(S_W^{-1} S_B) \leq K - 1$$

3. Perceptron Algorithm

- Generalised linear model $y = f(w^T \phi(x))$, where $\phi(x)$ is basis function; $\phi_0(x) = 1$

- Nonlinear activation function: $f(a) = \begin{cases} 1, & a \geq 0 \\ -1, & a < 0 \end{cases}$

- Target coding: $t = \begin{cases} 1, & \text{if } C_1 \\ -1, & \text{if } C_2 \end{cases}$

- Cost function:

- All correctly classified patterns: $w^T \phi(x_n) t_n > 0$

- Add the errors for all misclassified patterns (denoted as set \mathcal{M}):

$$\Rightarrow E_P(w) = - \sum_{n \in \mathcal{M}} w^T \phi(x_n) t_n$$

- Algorithm: (Aim: minimize total num of misclassified patterns)

- loop

choose a training pair (x_n, t_n)

update the weight vector w : $w = w - \eta \nabla E_P(w) = w + \phi_n t_n$

where $\eta=1$ because $y = f(\cdot)$ does not depend on $\|w\|$

- Perceptron Convergence Theorem:

- If the training set is linearly separable, the perceptron algorithm is guaranteed to find a solution in a finite number of steps

(Also is the algorithm to find whether the set is linearly separable = Halting Problem)

4. Maximum Likelihood

- Assumption: - $p(x|C_k) \sim \mathcal{N}(\mu_k, \Sigma)$, and all $p(x|C_k)$ share the same Σ - $p(C_1) = \pi, p(C_2) = 1 - \pi$, π unknown - Likelihood of whole data set \mathbf{X}, \mathbf{t} , N is the num of data - $p(\mathbf{X}, \mathbf{t} | \pi, \mu_1, \mu_2, \Sigma) =$

$$\prod_{n=1}^N [\pi \mathcal{N}(x_n | \mu_1, \Sigma)]^{t_n} [(1 - \pi) \mathcal{N}(x_n | \mu_2, \Sigma)]^{1-t_n} \rightarrow \text{when info of label } t \text{ lost: mixture of Gaussian}$$

$$\ln(\text{Likelihood}) = \sum_{n=1}^N [t_n (\ln \pi + \ln \mathcal{N}(x_n | \mu_1, \Sigma)) + (1 - t_n) (\ln(1 - \pi) + \ln \mathcal{N}(x_n | \mu_2, \Sigma))] - \text{Parameters}$$

$$\text{when maximum reached: } - \pi = \frac{N_1}{N_1 + N_2}, N_1 \text{ is the num of class } C_1 - \mu_1 = \frac{1}{N_1} \sum_{n=1}^N t_n x_n, \mu_2 =$$

$$\frac{1}{N_2} \sum_{n=1}^N (1-t_n)x_n, \text{ (mean of each class) } - \Sigma = \frac{N_1}{N} S_1 + \frac{N_2}{N} S_2, \text{ where } S_k = \frac{1}{N_k} \sum_{n \in C_k} (x_n - \mu_k)(x_n - \mu_k)^T$$

5. Logistic Regression

- Sigmoid function: $\sigma(a) = \frac{1}{1 + e^{-a}}$

- $p(x|C_k) \sim \mathcal{N} \implies p(C_k|x) = \sigma(w^T x + w_0)$ (2-classes) (Generative model)

- Assumption:

$p(x|C_k) = \mathcal{N}(\mu_k, \Sigma)$ (can also be a number of other distributions)

$\forall k, p(x|C_k)$ shares the same Σ

-

$$p(C_1|x) = \frac{p(x|C_1)p(C_1)}{p(x|C_1)p(C_1) + p(x|C_2)p(C_2)} = \sigma(a),$$

$$\text{where } a = \ln \frac{p(x, C_1)}{p(x, C_2)}$$

$$= \ln \frac{p(x|C_1)p(C_1)}{p(x|C_2)p(C_2)}$$

$$= \dots \text{ (assumption applied)}$$

$$= \ln \frac{\exp(\mu_1^T \Sigma^{-1} x - \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1)}{\exp(\mu_2^T \Sigma^{-1} x - \frac{1}{2} \mu_2^T \Sigma^{-1} \mu_2)} + \ln \frac{p(C_1)}{p(C_2)}$$

$$\implies a = w^T x + w_0 \text{ where,}$$

$$w = \Sigma^{-1}(\mu_1 - \mu_2)$$

$$w_0 = -\frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2} \mu_2^T \Sigma^{-1} \mu_2 + \ln \frac{p(C_1)}{p(C_2)}$$

- $\implies p(C_1|x) = \sigma(w^T x + w_0)$

\Rightarrow Find parameters in Gaussian model using Maximal Likelihood Solution

as: $p(C_1|x) \propto p(x|C_1)p(C_1) = p(x, C_1)$

- Generalize to K-classes:

$$a_k(x) = \ln[p(x|C_k)p(C_k)], p(C_k|x) = \frac{\exp(a_k)}{\sum_i \exp(a_i)}$$

$$\Rightarrow a_k(x) = w_k^T x + w_{k0}, \text{ where } w_k = \Sigma^{-1} \mu_k; w_{k0} = -\frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + p(C_k)$$

- Assume directly $p(C_k|x) = \sigma(w^T x + w_0)$ (2-classes) (Discriminative model)

- Assume directly: $p(C_1|w, x) = \sigma(w^T x), x_0 = 1$

\Rightarrow less parameters to fit (compared to Gaussian)

- Likelihood function:

$$p(\mathbf{t}|w, X) = \prod_{n=1}^N p_n^{t_n} (1 - p_n)^{1-t_n}, \text{ where, } p_n = p(C_1|x_n), t_n \text{ is the class of } x_n$$

Define error function :

$$E(w) = -\ln(\text{Likelihood}) = -\sum_{n=1}^N [t_n \ln p_n + (1 - t_n) \ln(1 - p_n)]$$

$$\Rightarrow \nabla E(w) = \sum_{n=1}^N (p_n - t_n) x_n$$

- Find Posterior $p(w|\mathbf{t})$:

Likelihood is product of sigmoid

Conjugate Prior for "sigmoid distribution" is unknown

\Rightarrow Assume Prior $p(w) = \mathcal{N}(w|m_0, S_0)$

$$\Rightarrow \ln p(w|\mathbf{t}) \propto -\frac{1}{2} (w - m_0)^T S_0^{-1} (w - m_0) + \sum_{n=1}^N [t_n \ln p_n + (1 - t_n) \ln(1 - p_n)]$$

find w_{MAP} , calculate $S_N = -\nabla \nabla \ln p(w|\mathbf{t}) = S_o^{-1} + \sum_{n=1}^N p_n(1 - p_n)\phi_n\phi_n^T$

$\Rightarrow p(w|\mathbf{t}) \simeq \mathcal{N}(w|w_{MAP}, S_N)$, via Laplace Approximation

- Laplace Approximation:

- Fit a gaussian to $p(z)$ at its ****mode**** (mode of $p(z)$: point where $p'(z) = 0$)

- Assume $p(z) = \frac{1}{Z}f(z)$, with normalization $Z = \int f(z)dz$

Taylor expansion of $\ln f(z)$ at z_0 : $\ln f(z) \simeq \ln f(z_0) - \frac{1}{2}A(z - z_0)^2$,

where $f'(z_0) = 0$, $A = -\frac{d^2}{dz^2} \ln f(z)|_{z=z_0}$

Take its exponentiating: $f(z) \simeq f(z_0)\exp -\frac{A}{2}(z - z_0)^2$

\Rightarrow Laplace Approximation $= (\frac{A}{2\pi})^{1/2} \exp -\frac{A}{2}(z - z_0)^2$, where $A = \frac{1}{\sigma^2}$

- Requirement:

$f(z)$ differentiable to find a critical point

$f''(z_0) < 0$ to have a maximum & so that $\nabla \nabla \ln f(z_0) = A > 0$ as $A = \frac{1}{\sigma^2}$

- In Vector Space: approximate $p(z)$ for $z \in \mathcal{R}^M$

Assume $p(z) = \frac{1}{Z}f(z)$

Taylor expansion: $\ln f(z) \simeq \ln f(z_0) - \frac{1}{2}(z - z_0)^T A (z - z_0)$,

Hessian $A = -\nabla \nabla \ln f(z)|_{z=z_0}$

$$\Rightarrow \text{Laplace approximation} = \frac{|A|^{1/2}}{(2\pi^{M/2})} \exp -\frac{1}{2}(z - z_0)^T A (z - z_0) \quad (2.1)$$

$$= \mathcal{N}(z|z_0, A^{-1}) \quad (2.2)$$

- Gradient descent:

- Hypothesis function: $h_\theta(x) = \sigma(x\theta) = \frac{1}{1+e^{-x\theta}}$

- Cost function:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^i \ln(h_\theta(x^i)) - (1 - y^i) \ln(1 - h_\theta(x^i))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- Update rule: $\forall \theta_j \in \theta, \theta_j := \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}, \frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m [(h_\theta(x^i) - y^i)x_j^i] + \frac{\lambda}{m} \theta_j$

2.10 Latent Variable Analysis

2.10.1 Principal Component Analysis (PCA)

1. Motivation:

- Data compression (reduce highly related features) - Data visualization

2. Assumption:

- Gaussian distributions for both the latent and observed variables

3. Two Equivalent Definition of PCA:

- Linear projection of data onto lower dimensional linear space (principal subspace) such

that:

\Rightarrow variance of projected data is maximized

\Rightarrow distortion error from projection is minimized

4. Maximum Variance Formulation

- Goal:

- project data from D dimension to M while maximizing the variance of projected data

- Eigenvalues λ of covariance matrix S express the variance of data set X in direction of

corresponding eigenvectors

- Projection Vectors:

- $U = (u_1, \dots, u_M)$, where $\forall i \in \{1, \dots, M\}, u_i \in \mathbb{R}^D$ s.t. $u_i^T u_i = 1$ (only consider direction)

- Projected Data:

- Mean = $\bar{x}^T U$, where $\bar{x} = \frac{1}{N} \sum_{i=1}^N x^i$ - Variance = $\text{tr}\{U^T S U\}$, where $S = \sum_{i=1}^N (x^i - \bar{x})(x^i - \bar{x})^T$ (outer product)
- Lagrangian to maximize Variance:
- $L(U, \lambda) = \text{tr}\{U^T S U\} + \text{tr}\{(I - U^T U)\lambda\}$
- constraint $u_i^T u_i = 1$ to prevent $u_i \rightarrow +\infty$

$$\text{For each } u_i \in U, \frac{\partial}{\partial u_i} L = 2S u_i - 2\lambda_i u_i = 0 \quad (2.3)$$

$$\Rightarrow S u_i = \lambda_i u_i \quad (2.4)$$

$$\Rightarrow U \text{ consists of eigenvectors corresponding to the first } M \text{ large eigenvalue of } S \quad (2.5)$$

(S symmetric $\Rightarrow U$ orthogonal)

5. Minimum Error Formulation:

- Introduce Orthogonal Basis Vector for D dimension:

- $U = (u_1, \dots, u_D)$

- Data representation:

- Original: $x^n = \sum_{i=1}^D \alpha_i^n u_i$ - Projected: $\tilde{x}^n = \sum_{i=1}^M z_i^n u_i + \sum_{i=M+1}^D b_i u_i$

(z_1^n, \dots, z_M^n) is different for different x^n , (b_{M+1}, \dots, b_D) is the same for all x^n

- Cost function: $J = \frac{1}{N} \sum_{n=1}^N \|x^n - \tilde{x}^n\|^2$, where $\tilde{x}^n = \sum_{i=1}^M z_i^n u_i + \sum_{i=M+1}^D b_i u_i$

$$\text{- Let } \begin{cases} \frac{\partial}{\partial z_j^n} J = 0 \\ \frac{\partial}{\partial b_j} J = 0 \end{cases} \Rightarrow \begin{cases} \frac{1}{N} 2(x^n - \tilde{x}^n)^T (-u_j) = \frac{2}{N} (z_j - (x^n)^T u_j) = 0 \\ \frac{1}{N} \sum_{n=1}^N 2(x^n - \tilde{x}^n)^T (-u_j) = \frac{2}{N} \sum_{n=1}^N (b_j - (x^n)^T u_j) = 0 \end{cases}$$

$$\Rightarrow \begin{cases} z_j = (x^n)^T u_j & j \in \{1, \dots, M\} \\ b_j = \bar{x}^T u_j & j \in \{M+1, \dots, D\} \end{cases}$$

$$\text{Noticing } (x^n)^T u_j = \left(\sum_{i=1}^D \alpha_i^n u_i^T \right) u_j = a_j \Rightarrow a_j = (x^n)^T u_j$$

$$\Rightarrow x^n - \tilde{x}^n = \sum_{i=M+1}^D [(x^n - \bar{x})^T u_i] u_i$$

-

$$\Rightarrow J = \frac{1}{N} \sum_{n=1}^N \left(\sum_{i=M+1}^D [(x^n - \bar{x})^T u_i] u_i \right)^T \left(\sum_{i=M+1}^D [(x^n - \bar{x})^T u_i] u_i \right) \quad (2.6)$$

$$= \frac{1}{N} \sum_{n=1}^N \left(\sum_{i=M+1}^D u_i^T ((x^n - \bar{x})^T u_i) \right) \left(\sum_{i=M+1}^D ((x^n - \bar{x})^T u_i) u_i \right) \quad (2.7)$$

$$= \frac{1}{N} \sum_{n=1}^N \sum_{i=M+1}^D u_i^T (x^n - \bar{x})^T u_i u_i^T (x^n - \bar{x}) u_i \quad u_i \text{ orthogonal to each other} \quad (2.8)$$

$$= \sum_{i=M+1}^D u_i^T \left(\frac{1}{N} \sum_{n=1}^N (x^n - \bar{x})^T (x^n - \bar{x}) \right) u_i \quad \|u_i\| = 1 \quad (2.9)$$

$$(2.10)$$

$$\Rightarrow J = \sum_{i=M+1}^D u_i^T S u_i, \text{ where } S = \frac{1}{N} \sum_{n=1}^N (x^n - \bar{x})^T (x^n - \bar{x})$$

- Lagrangian to Minimize J :

$$- L(u_{M+1}, \dots, u_D, \lambda_{M+1}, \dots, \lambda_D) = \sum_{i=M+1}^D u_i^T S u_i + \sum_{i=M+1}^D \lambda_i (1 - u_i^T u_i)$$

constraint $\|u_i\| = 1$ to prevent $u_i = 0$

$$\text{For each } u_i, \frac{\partial}{\partial u_i} L = 2S u_i - 2\lambda_i u_i = 0$$

$$\Rightarrow S u_i = \lambda_i u_i$$

\Rightarrow To minimize J , take eigenvectors with the first $(D - M)$ small eigenvalue orthogonal to (out of) subspace

\Leftrightarrow define subspace with eigenvectors with the first M large eigenvalue

-

$$\text{Intuition: } \widetilde{x}_n = \sum_{i=1}^M ((x^n)^T u_i) u_i + \sum_{i=M+1}^D (\bar{x}^T u_i) u_i \quad (2.11)$$

$$= \bar{x} + \sum_{i=1}^M [(x^n - \bar{x})^T u_i] u_i \quad (2.12)$$

1. Singular Value Decomposition - SVD:

- Introduce matrix $A_{m \times n}$

- $(A^T A)_{n \times n}$ symmetric matrix (actually, Gram matrix \rightarrow semi-definite) -

eigenvalue decomposition:

(2.13)

$$A^T A = V D V^T, V \text{ is normalized } (v_i^T v_i = 1) \text{ with column as eigenvector} \quad (2.14)$$

- $AV = (Av_1, \dots, Av_n)_{m \times n}$

- Let $r(A) = r$

$$\Rightarrow r(A^T A) = r(A) = r \quad (2.15)$$

$$r(AV) = \min\{r(A), r(V)\} = \min\{r, n\} = r \quad (2.16)$$

- Reduce AV to basis (Av_1, \dots, Av_r)

- Let $U = (u_1, \dots, u_r) = (\frac{Av_1}{\sqrt{\lambda_1}}, \dots, \frac{Av_r}{\sqrt{\lambda_r}})$, λ_i is i -thh eigenvalue of $A^T A$

- Orthogonal: $\forall i \neq j, u_i^T u_j = \frac{1}{\sqrt{\lambda_i \lambda_j}} v_i^T A^T A v_j = \frac{\lambda_j}{\sqrt{\lambda_i \lambda_j}} v_i^T v_j = 0$

- Unit: $\|u_i\| = \frac{\|Av_i\|}{\sqrt{\lambda_i}} = \frac{\sqrt{\langle Av_i, Av_i \rangle}}{\sqrt{\lambda_i}} = 1$

$\Rightarrow U$ is standard orthogonal (orthonormal) basis

- $AV = U \Sigma$, where $\Sigma = D^{\frac{1}{2}}$

- Expand U to orthonormal in $\mathbb{R}^m : (u_i, \dots, u_m)$

- Expand corresponding part in Σ with 0

- $A = U \Sigma V^T$, with singular value in Σ in decreasing order

2. SVD with PCA:

- X is data matrix in row (centered - zero mean)

- Eigenvectors of covariance matrix $S = X^T X$ are in V , where $X = U \Sigma V^T$

- When using $S = U \Sigma V^T \Rightarrow U = V \wedge S = V \Sigma V^T$

reduced to eigenvalue decomposition

- $S = V D V^T$ with V orthonormal:

Eigenvalues λ of covariance matrix S express the variance of data set X in direction of corresponding eigenvectors

- Projection:

- $\tilde{X} = X V_M$, where V_M contains first M -large eigenvectors - Projection direction is **not** unique

3. Reconstruction (approximate):

- Data is projected onto k dimension using SVD with $S = U\Sigma V^T$ - $x_{approx} = U_{reduce} \cdot z$, U_{reduce} is $n \times k$ matrix, z is $k \times 1$ vector - [Reconstruction from data Compression] (./../Machine

4. Choosing k (num of principal components):

- choose the **smallest** k making $\frac{J}{V} \leq 0.01 \Rightarrow 99$

- $[U, S, V] = \text{svd}(\text{Sigma}) \Rightarrow \frac{J}{V} = 1 - \frac{\sum_{i=1}^k S_{ii}}{\sum_{i=1}^n S_{ii}}$, S is diagonal matrix

\Rightarrow check $\frac{J}{V}$ before compress data

5. Data Preprocessing:

- PCA vs. Normalization: - Normalization: Individually normalized but still correlated - PCA: create decorrelated data - whitening - Whitening: projection with normalization - $S = VDV^T$, where S is Gram matrix over X^T - $\forall n, y_n = D^{-\frac{1}{2}} V^T (x^n - \bar{x})$, where \bar{x} is the mean of X

$\Rightarrow y^n$ has zero mean (2.17)

$$\text{cov}(\{y^n\}) = \frac{1}{N} \sum_{n=1}^N y_n y_n^T = D^{-\frac{1}{2}} V^T S V D^{-\frac{1}{2}} = I \quad (2.18)$$

6. Tips for PCA:

- Do NOT use PCA to prevent overfitting, use regularization instead - Try original data before implement PCA - Train PCA only on training set

2.10.2 Independent Component Analysis (ICA)

1. Goal: - Recover original signals from a mixed observed data - Source signal $S \in \mathbb{R}^{N \times K}$; mixing matrix A ; Observed data $X = SA$ - Maximizes statistical independence - Find A^{-1} to maximize independence of columns of S 2. Assumption: - At most one signal is Gaussian distributed - Ignore amplitude and order of recovered signals - Have at least as many observed mixtures as signals - A invertible 3. Independence vs. Uncorrelatedness - Independence \Rightarrow Uncorrelatedness - $p(x_1, x_2) = p(x_1)p(x_2) \Rightarrow \mathbb{E}(x_1 x_2) - \mathbb{E}(x_1)\mathbb{E}(x_2) = 0$ 4. Central Limit Theorem 5. FastICA algorithm

2.10.3 t-SNE

1. Problem & Focus 2. Compared to PCA: - No whitening function to use for new data - PCA can only capture linear structure inside the data - t-SNE preserves the local distances in the original data

2.10.4 Anomaly Detection

1. Problem to solve:

- Given dataset x^1, x^2, \dots, x^m , build density estimation model $p(x)$ - $p(x^{test} < \epsilon) \Rightarrow x^{test}$ anomaly

2. Hypothesis function:

$$p(x) = \prod_{i=1}^n p(x_i), x \in \mathbb{R}^n, \forall i \in [1, n], x_i \sim N(\mu_i, \sigma_i^2) - \mu = \frac{1}{m} \sum_{i=1}^m x^i, \sigma^2 = \frac{1}{m} \sum_{i=1}^m (x^i - \mu)^2 -$$

assume x_1, \dots, x_n independent from each other

3. Multivariate Gaussian:

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu)\right),$$

$x \in \mathbb{R}^n, \mu \in \mathbb{R}^n, \Sigma \in \mathbb{R}^{n \times n}$, where Σ is covariance matrix

$$- \mu = \frac{1}{m} \sum_{i=1}^m x^i, \Sigma = \frac{1}{m} \sum_{i=1}^m (x^i - \mu)(x^i - \mu)^T - x_1, \dots, x_n \text{ can be correlated but **not** linearly}$$

dependent - need $m > n$ ($m \geq 10n$ suggested) or else Σ non-invertible

4. Algorithm:
 - choose features - compute μ, σ - compute $p(x)$ for new example, anomaly if $p(x) < \epsilon$
5. Evaluation (real-number):
 - Labeled data into normal/anomalous set (okay if some anomalies slip into normal set)
 - training set: unlabeled data from normal set (60- CV set: labeled data from normal (20- test set: labeled data from normal (20
 - Use evaluation metrics (skewed data)
6. When to use:
 - Anomaly detection: - Very small num of positive data (0-20 commonly); Large num of negative data - Difficult to learn from positive data (not enough data, too many features...)
 - Future anomalies may look nothing like given data - Supervised Learning: - Larger num of positive & negative data - Enough positive data for algorithm to learn - Future positive example is likely to be similar to given data
7. Example:
 - Anomaly detection: - Fraud detection, Manufacturing, Monitoring machines in data center... - Supervised learning: - Email spam classification (enough data), Weather prediction (sunny/rainy/etc), Cancer classification...
8. Tips:
 - Non-gaussian feature: transformation / using other distribution - Choosing features: compare anomaly data with normal data

2.10.5 Recommender System

1. Problem Formulation:
 - $r_{i,j} = 1$ if item i is rated by user j
 - $y_{i,j}$ = rating of item i given by user j
 - θ^j = parameter vector for user j
 - x^i = feature vector for movie i
 - =, for user j , movie i , ($r_{i,j} = 0$), predict rating $x^i \theta^j$
2. Content Based Recommendations:
 - Treat each user as a separate linear regression problem with the feature vectors of its rated items as training set
 - **Assume features for each items (x^i) are available and known**
 - =, given X estimate Θ
 - Cost Function for θ_j :

$$J(\theta^j) = \frac{1}{2} \sum_{i:r_{i,j}=1} (x^i \theta^j - y_{i,j})^2 + \frac{\lambda}{2} \sum_{k=1}^n (\theta_k^j)^2, \theta^j \in R^{n+1} (\theta_0 \text{ not regularized})$$
 - Cost Function for Θ :

$$J(\Theta) = \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r_{i,j}=1} (x^i \theta^j - y_{i,j})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^j)^2,$$

$$\theta^j \in R^{n+1} (\theta_0 \text{ not regularized}), n_u \text{ is num of users}$$
 - Update Rule: $\forall \theta_k^j \in \theta^j, \theta_k^j := \theta_k^j - \alpha \frac{\partial J(\Theta)}{\partial \theta_k^j}, \frac{\partial J(\Theta)}{\partial \theta_k^j} = \sum_{i:r_{i,j}=1} (x^i \theta^j - y_{i,j}) x_k^i + \lambda \theta_k^j, \text{ for } k \neq 0 (\theta^j \in R^{n+1})$
3. Collaborative Filtering
 - **Assume preference of each users (θ^j) are available and known**
 - =, given Θ estimate X
 - Cost Function for x^i : $J(x^i) = \frac{1}{2} \sum_{j:r_{i,j}=1} (x^i \theta^j - y_{i,j})^2 + \frac{\lambda}{2} \sum_{k=1}^n (x_k^i)^2$ - Cost Function for X :

$$J(X) = \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r_{i,j}=1} (x^i \theta^j - y_{i,j})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^i)^2$$

$x^j \in R^{n+1}$ (x_0 not regularized), n_m is num of items - Update Rule: $\forall x_k^i \in x^i, x_k^i := x_k^i - \alpha \frac{\partial J(X)}{\partial x_k^i}$,

$$\frac{\partial J(X)}{\partial x_k^i} = \sum_{j:r_{i,j}=1} (\theta^j x^i - y_{i,j}) \theta_k^j + \lambda x_k^i, \text{ for } k \neq 0 \ (x^i \in R^{n+1})$$

- Basic Idea:

- Randomly initialize Θ

- loop:

Estimate X

Estimate Θ

- Cost Function:

$$J(X, \Theta) = \frac{1}{2} \sum_{(i,j):r_{i,j}=1} (x^i \theta^j - y_{i,j})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^i)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^j)^2, x \in R^n, \theta \in R^n$$

(the sum term in $J(\Theta)$, $J(X)$, and $J(X, \Theta)$ is the same)

- Update Rule:

$$\forall x_k^i \in x^i, x_k^i := x_k^i - \alpha \frac{\partial J(X, \Theta)}{\partial x_k^i}, \frac{\partial J(X, \Theta)}{\partial x_k^i} = \frac{\partial J(X)}{\partial x_k^i} = \sum_{j:r_{i,j}=1} (\theta^j x^i - y_{i,j}) \theta_k^j + \lambda x_k^i, x^i \in R^n$$

$$\forall \theta_k^j \in \theta^j, \theta_k^j := \theta_k^j - \alpha \frac{\partial J(X, \Theta)}{\partial \theta_k^j}, \frac{\partial J(X, \Theta)}{\partial \theta_k^j} = \frac{\partial J(\Theta)}{\partial \theta_k^j} = \sum_{i:r_{i,j}=1} (\theta^j x^i - y_{i,j}) x_k^i + \lambda \theta_k^j, \theta^j \in R^n$$

- **Algorithm**

- Initialize X, Θ to **small random values**

= \hookrightarrow for symmetry breaking (similar to random initialization in neural network)

= \hookrightarrow so that algorithm learns features x^1, \dots, x^{n_m} that are different from each other

- Minimize $J(X, \Theta)$

- Predict $y_{i,j} = x^i \theta^j$ ($Y = X\Theta$)

- Finding Related Item to Recommend

- $\|x^i - x^j\|$ is small = \hookrightarrow item i and j is similar

- Mean Normalization:

- Problem: if user j hasn't rated any movie, $\theta^j = [0, \dots, 0]$

= \hookrightarrow predicted rating of user j on all item = 0

= \hookrightarrow useless prediction

- Algorithm (row version):

compute vector $\mu, \forall \mu_i \in \mu, \mu_i = \text{mean of } Y_i$, where Y_i is the i^{th} row in Y

manipulate $Y: \forall y_{i,j} \in Y \wedge r_{i,j} = 1, y_{i,j} - \mu_i = \tilde{y}_{i,j}$ the mean of each row in Y is 0

predict rating for user j on item $i = x^i \tilde{\theta}^j + \mu_i$

- For item i with no rating

= \hookrightarrow apply column version of mean normalization

(but user with no rating is generally more important)

2.11 Large Scale Machine Learning

- Compute $cost(\theta, (x^i, y^i))$ before updating

For every k update iterations, plot average $cost(\theta, (x^i, y^i))$ over the last k examples

- Checking curves:

Increasing k result in smoother line and less noise, but the result is more delayed

2.11.1 Online Learning

1. Situation: - Has too many data (can be considered as infinite) - When data comes in as a continuous stream - Can adapt to changing user preference
2. Procedure: - Use one example only once (Similar to stochastic gradient decent in this sense)

2.11.2 Map-reduce

1. In Batch Gradient Descent:

- Update rule $\theta_j = \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)x_j^i$ - Parallelize the computation of $\sum_{i=1}^m (h_{\theta}(x^i) - y^i)x_j^i$ by dividing the data set into multiple sections

2. Ability to reduce:

- Contain operation over the whole data set

(Neural Network can be map-reduced)

Chapter 3

Linear Regression

Chapter 4

Linear Classification

Chapter 5

Kernel Methods

Chapter 6

Graphical Models

Chapter 7

Mixture Models and EM

Chapter 8

Approximate Inference

Chapter 9

Sampling Methods

Chapter 10

Continuous Latent Variable

Chapter 11


Sequential Data

- Challenge
 - Violated Identically Independent Draw Assumption
 - current data depends on previous data \Rightarrow i.i.d. assumption NOT hold
 - \Rightarrow distribution changing while drawing data
- Assumption
 - yet Gaussian still usually assumed s.t. model complexity remains in iterations (as exponential family's prior-posterior in the same family)
- Overview
 - Stationary
 - data evolves in time; generative distribution stays the same
 - Non-stationary
 - generative distribution changes along the time

11.1 Markov Model

11.1.1 Markov Chain


First-order Markov Chain

- Assumption
 - each data only depends on the most recent node (direct predecessor)
- Description
 - Bayesian Networks
 - 
 - \Rightarrow Distribution
 - joint distribution: $p(x_1, \dots, x_N) = p(x_1) \prod_{n=2}^N p(x_n | x_{n-1})$
 - conditional dependency: $p(x_n | x_1, \dots, x_{n-1}) = \dots = p(x_n | x_{n-1})$

Second-order Markov Chain

- Motivation
 - Modeling Trend
 - at least 2 observations to model a trend

- Description
 - Bayesian Networks



- Distribution

- joint distribution: $p(x_1, \dots, x_N) = p(x_1)p(x_2|x_1) \prod_{n=3}^N p(x_n|x_{n-1}, x_{n-2})$

- Generalization: M^{th} -order Markov Chain

- Distribution

- joint distribution: $p(x_1, \dots, x_N) = p(x_1)p(x_2|x_1) \cdots p(x_M|x_1, \dots, x_{M-1}) \times \prod_{n=M+1}^N p(x_n|x_{n-1}, \dots, x_{n-M})$


- Understanding
 - Analysis of M (assume K states for an observation)
 - $M = 0$: no markov parameter, data drawn i.i.d.
 - $M = 1$: first-order chain: $K - 1$ for previous observation $\Rightarrow K(K - 1)$ parameters
 - M : M^{th} -order chain: $K - 1$ for each of M observation $\Rightarrow K^M(K - 1)$ parameters
 \Rightarrow num of parameters grows exponentially with M

11.1.2 Hidden Markov Model

Overview

- Motivation
 - Improving High-Order Markov Chain
 - avoid exponentially growing parameters in high-order Markov chain
 - avoid fixed & restricted length of dependence
 - Sequential Correlation
 - model sequential correlations in data as an extension of mixture models
- Assumption
 - Latent Variable
 - latent variables are discrete & form a Markov chain
 \Rightarrow assumption on conditional independence
 - one latent variable z_n for each observation x_n
- Description

- Bayesian Networks



- Conditional Independence

- $z_{n+1} \perp\!\!\!\perp z_{n-1} \mid z_n$
- no blocked path between any two observed x_i, x_j
 \Rightarrow prediction depends on all previous observation

- Distribution


- joint distribution

$$\begin{aligned} p(x_1, \dots, x_N, z_1, \dots, z_N) &= p(z_1) \left[\prod_{n=2}^N p(z_n | z_{n-1}) \right] \prod_{n=1}^N p(x_n | z_n) \\ &= p(z_1) p(x_1 | z_1) \prod_{n=2}^N p(z_n | z_{n-1}) p(x_n | z_n) \end{aligned}$$

- \Rightarrow as an extension of mixture distribution:
 component densities = $p(x|z)$
 choice of component depends on previous state = $p(z_n | z_{n-1})$

- Transition Probabilities

- encoding: latent variable z with K state, in one-hot encoding
- $p(z_n | z_{n-1}) = \mathbf{A}_{K \times K} \in [0, 1]^{K \times K}$
 where $A_{ik} = p(z_{n,k} = 1 \mid z_{n-1,i} = 1), 0 \leq A_{ik} \leq 1$ and $\sum_{k=1}^K A_{ik} = 1$
 \Rightarrow num of independent parameters = $K(K-1)$
 $\Rightarrow p(z_n | z_{n-1}, \mathbf{A}) = \prod_{k=1}^K \prod_{i=1}^K A_{ik}^{z_{n-1,i} \times z_{n,k}}$
- $p(z_1) = \pi$
 where $\pi_k = p(z_{1k} = 1), 0 \leq \pi_k \leq 1$ and $\sum_k \pi_k = 1$
 $\Rightarrow p(z_1 | \pi) = \prod_{k=1}^K \pi_k^{z_{1k}}$
- transition diagram & unfolded transition diagram ($K=3$)



- Emission Probabilities

- $p(x_n | z_n, \phi) = \prod_{k=1}^K p(x_n | \phi_k)^{z_{n,k}}$
 where ϕ is a set of parameters governing the conditional distribution

- Variants of Hidden Markov Model

- Homogeneous HMM

- homogeneous assumption: $\forall n = 3, \dots, N, p(z_n|z_{n-1}) = p(z_{n-1}|z_{n-2})$
(transition probabilities are the same for all)


- \Rightarrow joint distribution

$$\begin{aligned} p(X, Z|\theta) &= p(z_1|\pi) \left[\prod_{n=2}^N p(z_n|z_{n-1}, A) \right] \prod_{n=1}^N p(x_n|z_n, \phi) \\ &= \prod_{k=1}^K \pi_k^{z_{1k}} \left[\prod_{n=2}^N \prod_{k=1}^K \prod_{i=1}^K A_{ik}^{z_{n-1,i} \times z_{n,k}} \right] \prod_{n=1}^N \prod_{k=1}^K p(x_n|\phi_k)^{z_{nk}} \end{aligned}$$


where $X = (x_1, \dots, x_N)$, $Z = (z_1, \dots, z_N)$ and $\theta = \{\pi, A, \phi\}$

- Left-to-Right HMM

- $\forall i > j, A_{ij} = 0 \Rightarrow$ can only go to larger num state
 \Rightarrow transition diagram:



- constraint over the maximal change of state Δ , in one step
 \Rightarrow unfolded transition diagram, with $\Delta = 1$:



- Maximum Likelihood Solution (via EM)

- Goal

- $\arg \max_{\theta} p(X|\theta)$, where likelihood $p(X|\theta) = \sum_Z p(X, Z|\theta)$
 \Rightarrow determine model parameter θ given observations X

- Notation

- $\gamma(z_n) = p(z_n|X, \theta)$,

where $\gamma(z_n) \in \mathbb{R}^K$, $\sum_{k=1}^K \gamma(z_{nk}) = 1, \forall k, \gamma(z_{nk}) \geq 0$

$$\Rightarrow \gamma(z_{nk}) = p(z_{nk} = 1|X, \theta) = \sum_{z_n} \gamma(z_n) z_{nk}$$

- $\xi(z_{n-1}, z_n) = p(z_{n-1}, z_n|X, \theta)$,

where $\xi(z_{n-1}, z_n) \in \mathbb{R}^{K \times K}$, $\sum_{i,j=1}^K \xi(z_{n-1,i}, z_{nj}) = 1, \forall i, j, \xi(z_{n-1,i}, z_{nj}) \geq 0$

$$\Rightarrow \xi(z_{n-1,i}, z_{nk}) = p(z_{n-1,i} = 1, z_{nk} = 1|X, \theta) = \sum_{z_{n-1}, z_n} \xi(z_{n-1}, z_n) z_{n-1,i} z_{nk}$$

- E-step

- evaluate posterior $P(Z|X, \theta) \Rightarrow$ evaluate $\gamma(z_n), \xi(z_{n-1}, z_n)$

- M-step

- $\theta = \arg \max_{\theta} \ln P(X|\theta) = \arg \max_{\theta} \sum_Z p(Z|X, \theta^{\text{old}}) \ln p(X, Z|\theta)$, where $\theta = \{\pi, A, \phi\}$


$$\begin{aligned} \blacksquare \arg \max_{\theta} Q &= \sum_{k=1}^K \gamma(z_{1k}) \ln \pi_k + \sum_{n=2}^N \sum_{i=1}^K \sum_{k=1}^K \xi(z_{n-1,i}, z_{nk}) \ln A_{ik} + \sum_{n=1}^N \sum_{k=1}^K \gamma(z_{nk}) \ln p(x_n | \phi_k) \\ \blacksquare \text{ where full combination } \sum_Z \prod_{\text{each combination of } z_{1-n}} &= \prod_X \sum_{\text{all possibility of each } x} \end{aligned}$$

○ Critical Points

$$\begin{aligned} \blacksquare \pi_k &= \frac{\gamma(z_{1k})}{\sum_{i=1}^K \gamma(z_{1i})} \\ \blacksquare A_{ik} &= \frac{\sum_{n=2}^N \xi(z_{n-1,i}, z_{nk})}{\sum_{n=2}^N \sum_{j=1}^K \xi(z_{n-1,i}, z_{nj})} \\ \blacksquare \text{ assume independent } \phi_k &\Rightarrow \text{ if Gaussian } \begin{cases} \mu_k = \frac{\sum_{n=1}^N \gamma(z_{nk}) x_n}{\sum_{n=1}^N \gamma(z_{nk})} \\ \Sigma_k = \frac{\sum_{n=1}^N \gamma(z_{nk}) (x_n - \mu_k)(x_n - \mu_k)^T}{\sum_{n=1}^N \gamma(z_{nk})} \end{cases} \end{aligned}$$

○ Solving $\gamma(z_{nk}), \xi(z_{n-1,i}, z_{nk})$


- sum-product: based on message passing tree structure in HMM
- alpha-beta algorithm: known as forward-backward algorithm (two approaches are equivalent i.e. derive the same recursion formula)
- with Bayesian networks as



\Rightarrow all paths through z_n is blocked conditioned on z_n

$\Rightarrow p(X|z_n) = p(x_1, \dots, x_n|z_n)p(x_{n+1}, \dots, x_N|z_n)$

illustration:



$$P(D, B) = \sum_A P(A, B)P(D|A)$$

$$P(F, B) = \sum_A \left(P(A, B) \sum_C P(C|B)P(F|C) \right)$$

$$P(D, F, B) = \sum_A \left(P(A, B)P(D|A) \sum_C P(C|B)P(F|C) \right)$$

$$\Rightarrow P(D|B)P(F|B) = P(D, F|B) \Rightarrow D \perp\!\!\!\perp F|B$$

6. Maximum Likelihood – EM algorithm:

- alpha-beta algorithm:

- Let $\alpha(z_n) = p(x_1, \dots, x_n, z_n)$

$$\Rightarrow \alpha(z_n) = p(x_n|z_n)p(x_1, \dots, x_{n-1}|z_n)p(z_n) \quad (11.1)$$

$$= p(x_n|z_n) \sum_{z_{n-1}} p(x_1, \dots, x_{n-1}, z_{n-1}, z_n) \quad (11.2)$$

$$= p(x_n|z_n) \sum_{z_{n-1}} p(x_1, \dots, x_{n-1}|z_{n-1})p(z_n|z_{n-1})p(z_{n-1}) \quad (11.3)$$

$$= p(x_n|z_n) \sum_{z_{n-1}} \alpha(z_{n-1})p(z_n|z_{n-1}) \quad (11.4)$$

$$\Rightarrow \alpha(z_1) = \prod_{k=1}^K [\pi_k p(x_1|\phi_k)]^{z_{1k}}$$

$$\text{- Let } \beta(z_n) = p(x_{n+1}, \dots, x_N|z_n)$$

$$\Rightarrow \beta(z_n) = \frac{1}{p(z_n)} \sum_{z_{n+1}} p(x_{n+1}, \dots, x_N, z_n|z_{n+1})p(z_{n+1}) \quad (11.5)$$

$$= \frac{1}{p(z_n)} \sum_{z_{n+1}} p(x_{n+1}, \dots, x_N|z_{n+1})p(z_n|z_{n+1})p(z_{n+1}) \quad (11.6)$$

$$= \sum_{z_{n+1}} p(x_{n+2}, \dots, x_N|z_{n+1})p(x_{n+1}|z_{n+1}) \frac{p(z_n|z_{n+1})p(z_{n+1})}{p(z_n)} \quad (11.7)$$

$$= \sum_{z_{n+1}} \beta(z_{n+1})p(x_{n+1}|z_{n+1})p(z_{n+1}|z_n) \quad (11.8)$$

$$\Rightarrow \beta(z_N) = 1, \text{ solved from } \gamma(z_N) = \frac{\alpha(z_N)\beta(z_N)}{p(X)}$$

$$\text{- } \gamma(z_n) = p(z_n|X) = \frac{P(X|z_n)p(z_n)}{p(X)} = \frac{\alpha(z_n)\beta(z_n)}{p(X)}$$

$$\text{- } 1 = \sum_{z_n} \gamma(z_n) = \frac{\sum_{z_n} \alpha(z_n)\beta(z_n)}{p(X)}$$

$$\Rightarrow p(X) = \sum_{z_n} \alpha(z_n)\beta(z_n)$$

$$\text{More conveniently, let } z_n = z_N \Rightarrow p(X) = \sum_{z_N} \alpha(z_N)$$

$$\Rightarrow \text{complexity: } \mathcal{O}(n), \text{ instead of } \mathcal{O}(2^n), \text{ where } n \text{ is the length of the chain}$$

$$\text{- } \xi(z_{n-1}, z_n) = \frac{\alpha(z_{n-1})p(x_n|z_n)p(z_n|z_{n-1})\beta(z_n)}{P(X)} :$$

$$\xi(z_{n-1}, z_n) = p(z_{n-1}, z_n|X) = \frac{p(X|z_{n-1}, z_n)p(z_{n-1}, z_n)}{p(X)} \quad (11.9)$$

$$= \frac{p(x_1, \dots, x_{n-1}|z_{n-1}, z_n)p(x_n, \dots, x_N|z_{n-1}, z_n) \times p(z_n|z_{n-1})p(z_{n-1})}{p(X)} \quad (11.10)$$

$$= \frac{p(x_1, \dots, x_{n-1}|z_{n-1})p(x_n|z_n)p(x_{n+1}, \dots, x_N|z_n) \times p(z_n|z_{n-1})p(z_{n-1})}{p(X)} \quad (11.11)$$

$$= \frac{\alpha(z_{n-1})p(x_n|z_n)p(z_n|z_{n-1})\beta(z_n)}{p(X)} \quad (11.12)$$

$$\text{where factorizing using conditional independence : } \begin{cases} [x_1, \dots, x_{n-1}], [z_n] & \text{on } z_{n-1} \\ [x_n, \dots, x_N], [z_{n-1}] & \text{on } z_n \\ [x_n], [x_{n+1}, \dots, x_N] & \text{on } z_n \end{cases}$$

$$(p(A|B) = p(A) \text{ when } A \perp\!\!\!\perp B)$$

- Algorithm description:

- Initialize $\theta = \{\pi, A, \phi\}$

- E step:

Forward recursion for $\alpha(z_n)$

Backward recursion for $\beta(z_n)$

Calculate $\gamma(z_n), \xi(z_{n-1}, z_n)$

- M step:

Maximize $Q(\theta, \theta^{\text{old}})$ using critical points

- Regularized EM

- add the prior of π, A, ϕ , in the form of $\log p(\theta)$, into $Q(\theta, \theta^{\text{old}})$ before maximization

7. Prediction

- Goal

- predict x_{N+1} given $X = \{x_1, \dots, x_N\}$

- Algorithm

- α recursion + summing over z_N

$$p(x_{N+1}) = \sum_{z_{N+1}} p(x_{N+1}, z_{N+1} | X) \quad (11.13)$$

$$= \sum_{z_{N+1}} p(x_{N+1} | z_{N+1}, X) p(z_{N+1} | X) \quad (11.14)$$

$$= \sum_{z_{N+1}} p(x_{N+1} | z_{N+1}) \sum_{z_N} p(z_{N+1} | z_N, X) p(z_N | X) \quad (11.15)$$

$$= \sum_{z_{N+1}} p(x_{N+1} | z_{N+1}) \sum_{z_N} p(z_{N+1} | z_N) p(z_N | X) \quad (11.16)$$

$$= \frac{1}{P(X)} \sum_{z_{N+1}} p(x_{N+1} | z_{N+1}) \sum_{z_N} p(z_{N+1} | z_N) \alpha(z_N) \quad (11.17)$$

\Rightarrow store the $\alpha(z_t)$ for predicting x_{t+1} **and** computing $\alpha(z_{t+1})$ once x_{t+1} observed

- Intuition: information in x_1, \dots, x_N stored in $\alpha(z_N)$

\Rightarrow enable real-time application

8. Scaling Factors

- Original $\alpha - \beta$ Recursion

$$- \alpha(z_n) = p(x_n | z_n) \sum_{z_{n-1}} \alpha(z_{n-1}) p(z_n | z_{n-1}) - \beta(z_n) = \sum_{z_{n+1}} \beta(z_{n+1}) p(x_{n+1} | z_{n+1}) p(z_{n+1} | z_n)$$

$\Rightarrow \alpha \rightarrow 0$ exponentially quickly to length of chain

\Rightarrow for long chain (100), α can exceed the dynamic range of computer

- Normalized and Rescaled α

$$- \text{normalize: } \hat{\alpha}(z_n) = \frac{\alpha(z_n)}{p(x_1, \dots, x_n)} = p(z_n | x_1, \dots, x_n)$$

\Rightarrow stay in the dynamic range

- let rescale factor $c_n = p(x_n | x_1, \dots, x_{n-1})$

$$\Rightarrow p(x_1, \dots, x_n) = \prod_{m=1}^n c_m \quad (11.18)$$

$$\Rightarrow \alpha(z_n) = \left(\prod_{m=1}^n c_m \right) \cdot \hat{\alpha}(z_n) \quad (11.19)$$

$$\Rightarrow c_n \hat{\alpha}(z_n) = p(x_n | z_n) \sum_{z_{n-1}} \hat{\alpha}(z_{n-1}) p(z_n | z_{n-1}) \quad (11.20)$$

- Rescaled β

$$- \text{let normalization } \hat{\beta}(z_n) = \frac{\beta(z_n)}{\prod_{m=n+1}^N c_m} = \frac{p(x_{n+1}, \dots, x_N | z_n)}{p(x_{n+1}, \dots, x_N | x_1, \dots, x_n)}$$

$$\Rightarrow c_{n+1}\hat{\beta}(z_n) = \sum_{z_{n+1}} \hat{\beta}(z_{n+1})p(x_{n+1}|z_{n+1})p(z_{n+1}|z_n) \text{ note: } c_{n+1} \text{ can be re-used from } \alpha \text{ recursion}$$

cursion

- EM under Rescaled $\alpha - \beta$

- monitoring likelihood: $p(X) = \prod_{n=1}^N c_n$

- E step:

$$\Rightarrow \gamma(z_n) = \hat{\alpha}(z_n)\hat{\beta}(z_n) \quad (11.21)$$

$$\xi(z_{n-1}, z_n) = c_n \hat{\alpha}(z_n) p(x_n|z_n) p(z_n|z_{n+1}) \hat{\beta}(z_n) \quad (11.22)$$

9. Viterbi Algorithm (max-sum algorithm)

- Goal:

- find the most probable sequence of latent variable

\Rightarrow find $\max_Z p(Z|X, \theta)$, where Z is sequence of latent states

- compare with: set of states being individually most probable $\Rightarrow \forall n$, find $\max_{z_n} p(z_n|X, \theta)$

\Rightarrow maximize $\gamma(z_n)$ for all n

- efficiency: searches space of paths efficiently ($\mathcal{O}(n)$ to the length of chain)

- Notation:

- $w(z_n) = \max_{z_1, \dots, z_{n-1}} \ln p(x_1, \dots, x_n, z_1, \dots, z_n)$

note: $w(z_n)$ is a function of z_n , with log probability maximized over z_1, \dots, z_{n-1}

- Recursion from Joint Distribution of HMM:

- $w(z_{n+1}) = \ln p(x_{n+1}|z_{n+1}) + \max_{z_n} \{\ln p(z_{n+1}|z_n) + w(z_n)\}$

$w(z_1) = \ln p(x_1, z_1) = \ln p(z_1) + \ln p(x_1|z_1)$

- Backtrack

- maximization over $z_n \Rightarrow$ individually done for each of the K states of z_{n+1}

- maintain a matrix record for each maximization:

let $\phi(k_n)$ be the state of z_n when $w(z_{n+1})$ getting maximum given $z_{n+1} = k$ (in state k)

$\Rightarrow k_n^{\max} = \phi(k_{n+1}^{\max})$, where k_n^{\max} is the desired state of z_n

$\Rightarrow k_{N-1}^{\max} = \phi(k_N^{\max}) = \phi(\arg \max_{z_N} w(z_N))$

11.2 Linear Dynamic System

1. Goal

- Continuous Latent Variable

- sum becomes integral

- practical sense \Rightarrow multivariate Gaussian distribution assumed

(so that complexity of posterior dose NOT increase)

- Sequential Correlation in Continuous Data

- an extension to continuous latent variable model (such as probabilistic PCA)

2. Notation

- Underlain Procedure - transition: $z_n = Az_{n-1} + w_n$, where noise $w \sim \mathcal{N}(w|0, \Gamma)$ - emission:

$x_n = Cz_n + v_n$, where noise $v \sim \mathcal{N}(v|0, \Sigma)$ - initialization: $z_1 = \mu_0 + \mu$, where noise $\mu \sim$

$\mathcal{N}(\mu|0, V_0)$ - Probabilities - transition: $p(z_n|z_{n-1}) = \mathcal{N}(z_n|Az_{n-1}, \Gamma)$ - emission: $p(x_n|z_n) =$

$\mathcal{N}(x_n|Cz_n, \Sigma)$ - initialization: $p(z_1) = \mathcal{N}(z_1|\mu_0, V_0)$ - Model Parameters - $\theta = \{A, \Gamma, C, \Sigma, \mu_0, V_0\}$

3. Maximum Likelihood - EM

- E step

- Inference problem - determine the local posterior marginals for latent variables

(sum-product algorithm)

- M step

4. Linear-Gaussian Model Features

- sequence of individually most probable latent variable \Leftrightarrow the most probable latent sequence
 \Rightarrow no need for Viterbi algorithm

- Joint Distribution

$$- P(X, Z) = p(z_1) \left[\sum_{n=2}^N p(z_n | z_{n-1}) \right] \sum_{n=1}^N p(x_n | z_n) \text{ - same form as HMM}$$

\Rightarrow an Gaussian (product of Gaussians)

\Rightarrow standard result available for its marginals and conditionals

\Rightarrow sum-product algorithm for faster computation

5. Inference

- Goal

- determine marginal distribution $P(Z|X)$ - prediction: $P(z_n, x_n | x_1, \dots, x_{n-1}, \theta)$ - used in real-time application

- Sum-Product Algorithm (Kalman Filter + Kalman Smoother)

- analogous to alpha-beta algorithm in HMM

$\Rightarrow \hat{\alpha}(z_n) = p(z_n | x_1, \dots, x_n)$, factor $c_n = p(x_n | x_1, \dots, x_{n-1})$

$$\Rightarrow c_n \hat{\alpha}(z_n) = p(x_n | z_n) \int_{z_{n-1}} \hat{\alpha}(z_{n-1}) p(z_n | z_{n-1}) dz_{n-1}$$

$$\Rightarrow c_n \hat{\alpha}(z_n) = c_n \mathcal{N}(z_n | \mu_n, V_n) \quad (11.23)$$

$$= \mathcal{N}(x_n | C z_n, \Sigma) \int \mathcal{N}(z_{n-1} | \mu_{n-1}, V_{n-1}) \mathcal{N}(z_n | A z_{n-1}, \Gamma) dz_{n-1} \quad (11.24)$$

$$= \mathcal{N}(x_n | C z_n, \Sigma) \mathcal{N}(z_n | A \mu_{n-1}, P_{n-1}), \quad (11.25)$$

$$\text{where } P_{n-1} = A V_{n-1} A^T + \Gamma \text{ (by integral of } \mathcal{N} \cdot \mathcal{N}) \quad (11.26)$$

$$\Rightarrow \mu_n = A \mu_{n-1} + K_n (x_n - C A \mu_{n-1}) \quad (11.27)$$

$$V_n = (I - K_n C) P_{n-1} \quad (11.28)$$

$$c_n = \mathcal{N}(x_n | C A \mu_{n-1}, C P_{n-1} C^T + \Sigma), \quad (11.29)$$

$$\text{where } K_n = P_{n-1} C^T (C P_{n-1} C^T + \Sigma)^{-1}, \text{ known as Kalman gain matrix} \quad (11.30)$$

- Initial Condition

$$\Rightarrow c_1 \hat{\alpha}(z_1) = p(z_1) p(x_1 | z_1), \text{ where } z_1 = p(x_1)$$

$$\Rightarrow \mu_1 = \mu_0 + K_1 (x_1 - C \mu_0) \quad (11.31)$$

$$V_1 = (I - K_1 C) V_0 \quad (11.32)$$

$$c_1 = \mathcal{N}(x_1 | C \mu_0, C V_0 C^T + \Sigma), \quad (11.33)$$

$$\text{where } K_1 = V_0 C^T (C V_0 C^T + \Sigma)^{-1} \quad (11.34)$$

- Interpretation

- $A \mu_{n-1}$: predicted mean of z_n , by projecting mean of z_{n-1} one step forward - $C A \mu_{n-1}$: predicted observation x_n , by emitting from predicted mean of z_n - $x_n - C A \mu_{n-1}$: error between predicted observation and actual observation - K_n : coefficient of error, giving a correction to the predicted mean of z_n

\Rightarrow making successive predictions & correcting them in the light of new observation

6. EM - Learning Model Parameters

Chapter 12

Deep Learning

12.1 Interview of Fame

12.1.1 Geoffrey Hinton

Knowledge Embedding

- BP
 - psychology view: knowledge in vectors
 - semantic AI: knowledge graph
 - BP algorithm can interpret & convert between feature vector and graph representation (with some embedding)
- Boltzmann Machine
 - Learning Algorithm on Density Net
 - same information in forward & backward propagation to learn feature embedding
 - Restricted Boltzmann Machine (RBM)
 - ways of learning in deep dense net with fast inference
 - iterative learning (adding layer after the above trained)
 - $\text{ReLU} \Leftrightarrow$ a stack of sigmoid functions (approximately) in RBM
 - ReLU units initialized to identity for efficient learning
- EM
 - EM with Approximate E Step
- vs. Symbolic AI
 - Symbolic AI: symbolic logic-like expression to do reasoning
 - yet, maybe state vector to represent knowledge

Brain Science

- Brain: Nets Implemented by Evolution
 - trying to train without BP
 - doing BP (get derivatives) with re-construction error (auto-encoder)

Memory in Nets

- Fast Weights for Short-term Memory
- Capsule Net
 - structured knowledge representation in each unit (feature with sets of property)
 - \Rightarrow enable nets to vote rather than filtering - thus better generalization
 - now working: published in **2017 NIPS**

Unsupervised Learning

- Importance
 - better than human eventually (as supervised learning has limited maximum)
 - GAN as a breakthrough

”Slow” Feature

- Non-linear Transform to Find Linear Transform
 - find a latent representation containing linear transform to do the work
 - e.g. change viewpoints: pixels \rightarrow coordinates \rightarrow linear transform \rightarrow back to pixels

Relations between Computers

- showing computer data to work
 - instead of programming it to work

12.1.2 Pieter Abbeel

Deep Reinforcement Learning

- Overall Challenge
 - Representation
 - Exploration Problem
 - Credit Assignment
 - Worst Case Performance
- Advantage (Deep Nets in RL)
 - network capturing the representation (state vector)
- Question in DRL
 - how to learn safely
 - how to keep learning (under small negative samples) e.g. better than human
 - can we learn the reinforcement learning program (RL in the RL)
 - long time horizon
 - use experience across tasks
- Success of DRL
 - simulated robot inventing walking... \Rightarrow single general algorithms to learn

12.1.3 Ian Goodfellow

Generative Adversarial Networks

- Generative Models
 - Resembling
 - trained to optimized the distribution behind training data
(then sampled from that distribution to get more imaginary training data)
 - \Rightarrow produce data to resemble the training data
 - Usage
 - semi-supervised learning
 - data augmentation
 - simulating scientific experiment
 - Previous Ways
 - Boltzmann Machine
 - Sparse Coding
 - Now: Generative Adversarial Networks (GANs)
 - Future
 - increase reliability of GANs (stabilizing)

12.1.4 Yoshua Bengio

Thoughts

- Fallacy
 - Smoothness in Nonlinearity
 - to ensure non-zero gradients every where
- Surprising Fact
 - ReLU in Deep Net
 - inspired initially by biological connection
- **Distribution v.s. Symbolic Representation**
 - Distributed Representation
 - distributed in lots of units, instead of a symbolic representation in a single cell
(agree on Geoffrey Hinton)
 - Curse of Dimensionality
 - neural net's distributed representation for joint distribution over random variables
 - \Rightarrow Word Embedding
 - generalized to joint distribution over sequence of words

Works

- Piecewise Linear Activation (PLU)
- Unsupervised Learning
 - Focus
 - Denoising auto-encoder
 - GANs
 - Importance
 - human ability: self-teaching, building world-model from perception
 - Unsupervised Learning + Reinforcement Learning
 - underlying concept across two fields: machine can learn through interactions
⇒ learning "good" representation (yet, what is "good")
 - Possible Directions
 - loss function: not even defined for each task
(not knowing which is good for what?)
- Attention
 - Machine Translation (Founder)
 - Generalized into Other Fields
- Back-prop in Brains (Neural Science)
 - Reasons for Efficiency of Backprop
 - Larger Family behind Credit Assignment

12.1.5 Yuanqing Lin**National Deep Learning Lab**

- Paddle Paddle
- Baidu Lab

12.1.6 Andrej Karpathy**Human Benchmark**

- Programming by Showing
 - Requirement
 - input + output as specification
 - metric as goal
 - Writer
 - the optimizer
- Understanding Importance of Benchmark
 - importance to do better given the current performance on the dataset
(as important increase after passing human error)
- Understanding Network Behavior
 - compared to the process of human decision

Transfer Learning

- Image Task
 - feature extractor + fine tune/modification onto various task

12.1.7 Ruslan Salakhutdinov**Restricted Boltzmann Machine**

- Auto Encoder
 - Encoding All Kind of Data
 - from digit to face, document, etc...
 - deeper and deeper structure
- Training Boltzmann Machine
 - Pretraining
 - increase the low boundary by training the previous layer
 - then add another layer to train, ...
 - Direct Training (with GPU)
 - similar, or better result
- Boltzmann Machine Ability
 - Generative Model
 - model coupling distributions in data
⇒ scalable (more scalable than current model&operation)
 - only way to train the model in the early age
- Progress on Generative Model
 - probabilistic max pooling
 - variational encoder
 - deep energy model
 - semi-supervised Model

12.1.8 Research**Topics**

- Point Cloud
 - Operations on Points: how to embed location in operation
 - select fixed number of points via coord?: then take weighted average (conv) / max (pooling) on them
 - need a "select input points" op: like deformable conv?
 - Bounding Box Directly from Points: no voxel
 - clustering + regression on each cluster?
- Unsupervised Learning
 - Deep Belief Nets

- Reinforcement Learning
 - Deep Reinforcement Learning
 - scalable system
 - communicative& cooperating agents
- One-shot / Transfer Learning
 - Learning the Ability to Learn
- General AI
 - Structure for General Task
 - neural network or other structure, shared for multiple tasks
(instead of breaking down to different parts like segmentation, detection, etc.)
(instead of the split of cv, nlp, planning, etc.)
 - \Rightarrow a full agent (instead of decomposed function)
 \Rightarrow optimization method/objective need to be carefully defined
 - Attempt for General AI
 - scaling up supervised learning: imitating human
 - unsupervised learning: AIXI, artificial evolution, etc.
- AI Security
 - Anti Inducing
 - NOT to be fooled/induced to do unappropriated things
(even if algorithm is right)
 - Built-in Security
- Fairness in AI
 - Dealing Societal Issue
 - Reflecting Preferred Bias
- Auto Optimization (Hyperparameter Tunning)
 - Swarm Optimization
 - Expectation Maximization
 - target variable θ = hyperparameters
 - hidden variable Z = weights of network
 - data X = dataset

\Rightarrow

 - E-step: evaluate $\mathbb{E}_{Z|\theta_n, X}(\ln P(Z, X|\theta))$
 - $\ln P(Z, X|\theta)$: log likelihood of hyperparam θ (for weights & data to be observed)
 - $P(Z|\theta_n, X)$: posterior of weights Z
 - \Rightarrow evaluate (approximate) the expectation of the log likelihood of hyperparam θ
(from a functional view, train with $\theta_0 - \theta_N$, evaluate model M times in training, thus with weights $Z_{00} - Z_{NM}$)
 - \Rightarrow a matrix with n as row entry, m as column entry, mapping to both $\ln P(Z, X|\theta)$, $P(Z|\theta_n, X)$
 - \Rightarrow then marginalize (taking the expectation) over Z , to get a (sampled) function over θ
 - M-step: maximize the result function from E-step

· fit a curve & maximize w.r.t hyperparams θ

- World Understanding: after perception
 - Unsupervised Learning + Reinforcement Learning
 - machine learns from interactions
 - machine builds a representation of world (like human ability, without fine label)
 - ⇒ building world-model from perception
 - Causality Mining
- Model Interpretation
 - Logical Formalization
 - deep learning can be understood logically
 - e.g. what make deep net training harder? understand the limit of current algorithm/model and **why**

Advises

- Learning Direction
 - Math
 - statistic
 - linear algebra
 - calculus
 - optimization
- Reading
 - read a little bit & find somewhere intuitively not right
 - good intuition: eventually work;
 - bad intuition: not working no matter what it is doing
 - if other doubts your idea as bullshit ⇒ a sign for real good result
 - a supervisor with similar belief
 - PhD vs. Company
 - amount of mentoring
 - faster if dedicated supervisor available
 - resource
- Practice
 - open-source learning resource
 - open source contribution
 - contribute to open source framework (e.g. conv on sparse matrix in TF)
 - implement the paper, the open source it (as a tool for other)
 - work on a projected and open source it
 - ⇒ the stage (e.g. github) will bring people to you
 - implement the tools: to find out how & why it works
 - ⇒ derive theories from the
 - full stack of understanding
 - ⇒ understand the implementation under the deep learning framework

Direction - Detection

- Faster One-stage
 - Segmentation as Detection
 - for output mask, giving classification (p_e) + localization (regression) (instead of giving class probability as per-pixel classification)
 - at most as many objects as pixel number
⇒ not possible to lose detection due to gridding
- No Non-Maximum Suppression
 - Auto Filtering
 - regress threshold of objectness as well
⇒ objectness under threshold not considered by encoder-decoder (under the GAN framework ?)
(under multi-tasking ?)
 - CRF
 - still not end2end, yet trainable
 - RNN Encoder-Decoder
 - encode all bbox / spatial feature, then decode (generate) till end-of-sequence (as YOLOv1 still use dense layer...)

Direction - Tracking

- Current State - 2019
 - NN in Data Association
 - pred association prob of prediction-detection in JPDA
 - pred prob of detection being true / false alarm (the existence prob for bbox), used in MHT
 - Tracking Initiation & Deletion
 - currently, by rules...
e.g. 3 observations in 4 consecutive frames; observation missed for 4-7 frames
 - Prediction Encoding
 - crop the image according the predicted area (gating)
- NN in Gating
 - Region Proposal Network
 - propose the search region (gating) by RPN
(instead of using rule based on current track location)
⇒ predict a large region with RPN based on track info
- End-to-end / unified framework
 - End-to-End Regression for Data Association in MOT
 - cite: Online Multi-Target Tracking Using Recurrent Neural Networks
 - directly output association decision
 - input: raw image (semantic), current detection, track prediction, track history
 - track-level info: extended Siamese for relation of: track - current detection
(instead of only det-det in 2 frames)

- fully NN approach for JPDA
 - ⇒ track - all det: concat (tile) track+pred to each det ⇒ use fcn, then pooling to regress to a bbox for update
 - (direct regress the box, instead of Siamese net + weighted sum)
 - (cite: Data-Driven Approximations to NP-Hard Problems)
- ⇒ fast JPDA: single RNN encode all detections as D , then regress for each track- D for update
 - (instead of num of track \times num of det, now only num of track + num of det)
- fully NN approach for MHT ???
 - trace-back available ... encoding past detection as well? (in the tracker?)
 - N-scan pruning?
- Track Initiation/Deletion
 - directly output decision result
 - ⇒ regress the decision boundary as well (instead of universal 0.5)
- Interesting Predicted Bounding Box Encoding
 - direct encoding: x,y,w,h, o, + encoding of full image
 - mask: bounding box plot onto a separate mask, concat to the (encoding of) full image
 - further, attention mechanism
 - (soft-crop ???)
 - ⇒ naturally develop into instance tracking, even tracking in point cloud data
 - two masks: one for all bbox from detector, one for current track

⇒ enable global track for multiple obj track (single RNN for all track in an image)
- SOT
 - directly use cnn + conv-lstm
- single Obj Tracking for MOT
 - common cnn encodes image t as f_t , detected bbox as D_t (given or trained)
 - (prefer to train, as auxiliary loss for multi-tasking ?)
 - encode each detection $d_t \in D_t$ as $c(d_t) \in c(D_t)$
 - empty detection always $\in D_t$
 - ⇒ track may always have NO compatible det
 - one rnn tracker for each $c(d_t)$
 - rnn takes in encoded previous prediction $c(p_{t-1})$ (produced by itself), with f_t , and each encoded det $c(d_t) \in c(D_t)$
 - ⇒ produce as many track as D_t , each with a probability
 - (control the track's death)
 - ⇒ only the highest remains, others deleted
 - remaining track rnn regress pred bbox p_{t+1} (an encoded bbox as well)
 - (a decoder at t trying to induce the d_{t+1})
 - (GAN ?)
- MOT as SOT
 - attention mask contain multiple interest
 - ⇒ not able to handle birth/death of independent object
 - (as modeled all as a whole)
- End-to-End MOT
 - design requirement
 - arbitrary start of track
 - arbitrary num of track

- arbitrary end of track
- End-to-end MOT with States (Markov Decision Process / Deterministic Finite Automaton)
 - action at each state given by NN
 - history of track in RNN \Rightarrow provided when making decision
 \Rightarrow RNN modeling all the state ???
 (update with different set of weights on different chosen actions ?)
 (design transition instead of states: states as the closure of all actions ?)
 - design state for each challenge scene separately (out-of scene, occlusion, long-term lost, etc.)
 \Rightarrow directly tackle each scene
 (state growing: auto-discover state ???)
 - RNN for actions ?
 output score for all actions, only legal actions (given current state) considered&selected, then transfer state accordingly
 - trained with RL
- Training
 - MDP as Hard-Negative Mining
 - Sparse Detection
 - feed only a few detection, demand NN to fill up using track history&prediction
 - Other NLP Training
 - training of n-gram model used on rnn tracker ?
- Fixable Track
 - Fixing After Lost
 - backward rnn for re-associated track after lost
 \Rightarrow to fix the previous prediction given current observation
 (than recompute the forward rnn for consecutive tracking)

Directoin - Referring Seg

- Integrating Encoder-Decoder Architecture
 - Upsampling
 - similar to Unet, concat low-level spatial info
 - introduce language info as well
 (e.g. early combination, explicit introducing, ...)
- Info Early Fusion
 - Tiling at First Conv
 - as siamese net for joint input
 - downsampling more responsible for language info processing
 \Rightarrow hopefully get more fine-tuning alone with conv filters
 - can be used with pre-trained net:

$$ReLU(conv_1 * X_1 + conv_2 * X_2) = ReLU([conv_1, conv_2] * [X_1, X_2])$$
 - Multiple Entries
 - combining info at different stages of downsampling / upsampling


- Attention
 - Attention from Combined Info
 - as key-word-aware net
 - Attention on Language Info
 - 1-D spatial pyramid pooling / attention mask on the sentence encoding
- Language Info Throughout Network
 - Encoder-Decoder for Language Info
 - network asked to recover language info after processing combined info (potentially via a separate branch only at training time)
⇒ auxiliary loss
 - Language as Conv Filter
 - Language Info, through a subnet, becoming a set of conv filters
⇒ then imposed in downsampling, tunnel, upsampling or bridge stage(s)
- Data Augmentation
 - Translation Module
 - using the same image
 - expression translated to a middle language and then back to English
⇒ language info trained more finely

12.2 Basic Neural Network

12.2.1 Advantages

Large/Big Data

- Larger Maximum Capability
 - Curve given Amount of Data



- Reasons
 - the scale of data (labeled)
 - the scale of neural network (computability)
 - the scale of efficiency: e.g. ReLu, faster parallel algorithm

Flexibility

- Different Structures for Different Tasks
 - Same Data & Task
 - changing settings/structures of deep learning model can make a difference (v.s. SVM, etc.)
- Ability to Choose Basis Functions
 - Functional View
 - $y(\mathbf{x}, \mathbf{w}) = f(\mathbf{w}^T \phi(\mathbf{x}))$, where ϕ is basis function, $f(\cdot)$ is net as a function
 - Learning ϕ : choose embedding \Rightarrow choose basis function
 - Learning \mathbf{w} : choose which feature / basis functions more useful
- Solving Bias-Variance Trade-off
 - Complexity + Data/Regularization
 - easy complexity via depth, size
 - \Rightarrow reduce bias, without hurting variance by utilizing big data
 - easy regularization via L2 and etc.
 - \Rightarrow prevent high variance without hurting bias much in a deep/big net

Power of Depth

- Deep Representation
 - Low-level \rightarrow High-level
 - multiple layers to choose & combine useful information (creating new feature/basis)
 - \Rightarrow next layer use chosen/combined simple basis to build more complex one
 - \Rightarrow an hierarchy from low-level information to high-level information
- Circuit Theory
 - Power of Combination
 - functions that can be compactly represented by a depth k architecture might require an exponential number of computational nodes using a depth $k - 1$ architecture (from the perspective of factorization)

Yet, start from the SHALLOW (logistic regression) before trying the deep

12.2.2 Problem

(n units in one hidden layer)

Weight-space Symmetries

- Symmetries in Activation Function
 - $\mathcal{O}(2^n)$, e.g. $\arctan(-x) = -\arctan(x) \Rightarrow$ changing signs of all input & output has the same mapping (reduce effective data)
- Positional Combination in One Layer
 - $\mathcal{O}(n!)$ exchange unit with each other (together with their input output weights) \Rightarrow mapping stay the same

$\Rightarrow \mathcal{O}(n!2^n)$ overall weight-space symmetries

High-Dimension Search Space

- Multiple Critical Points
 - Symmetries
 - at least $\mathcal{O}(n!2^n)$ critical points ($\nabla E(w) = 0$), where $E(w)$ is error function due to weight-space symmetries
 - Saddle Points
 - both the bottom (in one dimension) and the top for another
 - due to high-dimension weight space
 \Rightarrow more likely to have functions being convey/convex in different dimensions
 - Local Optima
 - less then saddle points in amount, due to high-dimension weight space
 e.g. usually $\geq 10^4$ -D for modern deep nets
- Plateaus
 - a large flat region where gradient $\rightarrow 0$
 \Rightarrow gradient descent slowly down the flat surface (before exiting)
 - \Rightarrow slow down gradient descent significantly
- Expensive in Finding Critical Point
 - expensive for even local optima with gradient decent
 - as expensive as $\mathcal{O}(n^3)$ if using Laplace approximation

Gradient Vanishing/Exploding

- Gradient Vanishing
 - Saturated Function
 - sigmoid/tanh function: gradient $\rightarrow 0$ when input $\rightarrow \pm\infty$
 - Exponential Effect
 - with depth L , each activation (e.g. tanh) output $a^l < 1$ and weight $\mathbf{w}^l < 1$
 $\Rightarrow y(\mathbf{x}, W) \approx w^L \mathbf{w}'^{L-1} \mathbf{x}$, with $\mathbf{w}' < 1$
 \Rightarrow all the gradient along the way get multiplied by number < 1
 \Rightarrow gradient exponentially decayed in back-prop
- Gradient Exploding
 - Exponential Effect
 - similarly, each activation (e.g. ReLU) output $a^l > 1$ and weight $\mathbf{w}^l > 1$
 $\Rightarrow y(\mathbf{x}, W) \approx w^L \mathbf{w}'^{L-1} \mathbf{x}$, with $\mathbf{w}' > 1$
 \Rightarrow all the gradient along the way get multiplied by number > 1
 \Rightarrow gradient exponentially augmented in back-prop
- Possible Solutions
 - Random Initialization
 - Xavier Initialization: for gradient vanishing & exploding
 - Activation
 - ReLU: for gradient vanishing
 - Skip/Concat Connection
 - residual block
 -

12.2.3 Learning

Forward-Backward Propagation

- Representation
 - Layers
 - input layer
 - hidden layer(s): layer with NO ground truth (for the associated weights) available
note: input & hidden layers have associated biases as well (usually)
 - output layer
 - Neuron (Unit)
 - s_l : num of units in layer l
 - w^l : weight matrix of mapping from layer l to $l + 1$, with shape of $(s_{l+1}, s_l + 1)$
 - $h(\cdot)$: activation function (usually shared)
 - a_j^l : activation output of unit j at layer l
 - z_j^l : output of unit j at layer l
(represent parameterized basis, also the input for layer $l + 1$)
 - Intuition
 - all stacked vertically (vertical vector)
⇒ horizontally for different examples; vertically for different units
- Forward Propagation (Inference)
 - Activation $a^{j+1} = w^j \cdot [z_0^j, \dots, z_{s_j}^j]^T$, with $z_0 = 1$
 - Unit Output $z^{j+1} = h(a^{j+1}) = [z_1^{j+1}, \dots, z_{s_{j+1}}^{j+1}]^T$
- Backward Propagation
 - Loss $\mathcal{L}(W) =$
- Practice of Back Prop
 - Caching Intermediate Result
 - naturally cached: input $a^0 = x$, weights matrix w and bias b
 - activation input/output a/z
(since will be used in back-prop)
 - Auto Difference
 - achievement: calculate the derivatives along the forward prop !

12.3 Operations & Layers Structure

12.3.1 Operations in Network

Activations

- Sigmoid $a = \sigma(z)$
 - Pros
 - mapping to $(0, 1)$, with $\sigma(0) = 0.5$
 - Cons

- gradient vanishing: $\sigma(z)' = \sigma(z)(1 - \sigma(z)) \Rightarrow \lim_{z \rightarrow \pm\infty} \sigma(z)' \rightarrow 0$
(as the gradient passed through (via chain rule) $= \frac{a}{z} \frac{z}{w}$)
- Tangent $a = \tanh(z)$
 - Pros
 - empirically, almost always better than sigmoid (in hidden layers)
 - maps to $(-1, 1)$, with $\tanh(0) = 0 \Rightarrow$ help centering data (0-mean)
 \Rightarrow make the learning of next layer easier
 - Cons
 - still, gradient vanishing when $z \rightarrow \pm\infty$
- Rectified Linear Unit (ReLU) $\max(0, z)$
 - Derivation: approximated by a stack of sigmoid
 -
 - Pros
 - mitigate gradient vanishing: $\forall z > 0, a = z \Rightarrow$ learn much faster
 \Rightarrow the default choice!
 - Cons
 - undefined behavior at $x = 0$ (actually, gradient becomes the sub-gradient)
 - gradient totally vanished for $x < 0$
 - \Rightarrow dead units: weights learned/initialized to always output negatives
 \Rightarrow activation always output 0
 \Rightarrow the unit always output 0
- Leaky Relu $a = \max(\alpha z, z), \alpha \rightarrow 0^+$ (e.g. $\alpha = 0.01$)
 - Pros
 - mitigate the gradient vanishing problem for $(-\infty, +\infty)$
 - avoid dead units problem
 (yet not that popular as ReLU)
- Piecewise Linear Unit (PLU) $a = \max(\alpha(z + \beta) - \beta, \min(\alpha(z - \beta) + \beta, z))$
 - Pros
 - hybrid of tanh & ReLU: three linear pieces approximating tanh in a given range
 - more expressive than ReLU: more nonlinear, better to fit smooth nonlinear function
 - mitigate gradient vanishing problem: due to linearity
 - Cons
- Linear (Identity) Activation $a = z$
 - Pros
 - used in regression to output real number $\in (-\infty, +\infty)$
 - used in compression net
 - Cons
 - stacked units with linear activation \Leftrightarrow single linear transformation
 - logistic regression with linear activation in hidden layer is NO more expressive than logistic regression with no hidden layer !

Normalization in Network

- Batch Normalization
 - Definition
 - for an activation in hidden layer with input z , a batch with size N_b
 - calculate the mean of current batch $\mu = \frac{1}{N_b} \sum_n z_n$, where z_n for the n^{th} example
 - calculate the deviation of current batch $\sigma = \sqrt{\frac{1}{N_b} \sum_n (z_n - \mu)^2}$
 - normalize to be $z'_n = \frac{z_n - \mu}{\sigma}$
 - allow model to recover/manipulate original distribution: $\hat{z}_n = \gamma z'_n + \beta$, where γ, β being trainable (updated by optimizer using gradients)
 - Implementation
 - preferred to apply batch norm on z (before activation), instead of after it
 - for math stability, $z'_n = \frac{z_n - \mu}{\sigma + \epsilon}$, with $\epsilon \rightarrow 0+$
 - (usually) with mini-batch, calculate the mean & variance from only the mini-batch
 - with batch norm, original bias b in calculating $z = wx + b$ becomes pointless \Rightarrow integrated into the β in batch norm
 - at test time (1 example a time): need an estimation for μ, σ \Rightarrow exponentially weighted average over β, σ in training time
 - Understanding
 - normalize the intermediate data to have 0 mean, unit variance \Rightarrow to speed up the training from some hidden layers (as normalization does)
 - remain the ability to transfer the data to have other mean & variance (controlled by γ, β)
 - control the distribution of data in hidden layer \Rightarrow suppress the change of input data distribution for the layer after it \Rightarrow increase robustness for later layers, against covariate shift (from both the weight update in early layers and the input data change)
 - regularize the net by adding noise to the input data of hidden layer (due to computing mean/variance only on mini-batch) \Rightarrow enforce robustness against noise, hence unintended slight regularization effect

12.3.2 Operations on Network

Initialization


- Random Initialization for Weights
 - Practice
 - weights initialized to a random variable in a small range e.g. $(-0.03, 0.03)$
 - Pros
 - avoid symmetry problem: if identical initialization for weights \Rightarrow units in same layer computing exactly same function \Rightarrow get the same learning step propagated back \Rightarrow then always compute exactly the same function (by induction)
 - avoid gradient vanishing: especially for gradient of sigmoid/tanh activation

- Cons
 - NOT concern various nets: sampling in a fixed range may not work for all nets
- Xavier Initialization for Weights
 - Practice
 - set $\forall l \in [1, L], \text{Var}(w^l) = \frac{1}{n_l}$ for tanh, $\frac{2}{n_l}$ for ReLU, where n_l is the number of unit in layer l
 - Implementation
 - draw random variable $r \sim \mathcal{N}(0, 1)$
 - set each of $w^l = r \cdot \sqrt{\frac{2}{n_l}}$ for ReLU, $r \cdot \sqrt{\frac{2}{n_l}}$ for tanh or $r \cdot \sqrt{\frac{2}{n_{l-1} + n_l}}$ proposed by
 - Pros
 - theoretically justified to initialize weights to be around ± 1
 \Rightarrow mitigate gradient vanishing& exploding problem statistically
- Zero Initialization for Bias
 - Reason
 - default to use 0 bias
 (can NOT used for weights as explained)

Regularization

- L2 Regularization
 - Definition
 - \Rightarrow also called "weight decay"
 (as in gradient decent, weight is multiplied by a < 1 number due to L2 term)
 - Understanding
 - forcing weights to be smaller
 - single node has smaller effect
 - input of activation closer to 0
 \Rightarrow activation becomes more linear-alike (e.g. sigmoid, tanh)
 \Rightarrow layers perform more linear-alike transformation
 - \Rightarrow simpler network, less able to fit extreme curly decision boundary
 (hence less able to overfit)
- L1 Regularization
 - Definition
 - for each weight w we add the term $\lambda|w|$ to the objective. It is possible to combine the L1 regularization with the L2 regularization: $\lambda_1|w| + \lambda_2 w^2$ (this is called Elastic net regularization). The L1 regularization has the intriguing property that it leads the weight vectors to become sparse during optimization (i.e. very close to exactly zero). In other words, neurons with L1 regularization end up using only a sparse subset of their most important inputs and become nearly invariant to the "noisy" inputs. In comparison, final weight vectors from L2 regularization are usually diffuse, small numbers. In practice, if you are not concerned with explicit feature selection, L2 regularization can be expected to give superior performance over L1.

- Practice
 -
- Understanding
 -
- Cons
 -
- Dropout Regularization
 - Definition
 - for each of selected units, set a drop probability
i.e. for each forward/back-prop, nodes are "dropped" according to the probability
⇒ for each time, a randomly reduced net is trained
 - Implementation: Inverted Dropout
 - set a keep prob k instead of drop prob, for a selected layer
 - generate random numbers for all units & turned into a boolean "keep" vector \mathbf{k}
 - dropped activation $\mathbf{d} = \mathbf{a} \times \mathbf{k}$ (element-wise),
where \mathbf{a} is original activation output vector from the layer
 - ⇒ activation becomes 0 for dropped units in \mathbf{d}
 - scaling up by dividing the keep prob: \mathbf{d}/k
⇒ so that expected output value of each activation remains the same
 - test time: no dropout ⇒ no random output & consider all robust features learned
(randomness in training, mitigated by big data)
 - Understanding
 - can NOT rely on any one feature ⇒ have to spread out weights
⇒ results in shrinking the squared norm of weights (as $L2$)
 - used on layers with enormous features as input (e.g. computer vision)
⇒ reduce the chance of relying on small set of features



- Cons
 - training loss may have bigger glitch ⇒ harder to debug
(make sure loss decreasing before introduced dropout)
- Max norm constraints
 - Definition
 - enforce an absolute upper bound on the magnitude of the weight vector for every neuron and use projected gradient descent to enforce the constraint. In practice, this corresponds to performing the parameter update as normal, and then enforcing the constraint by clamping the weight vector w_n^l of every neuron to satisfy (\cdot) .

Typical values of α are on orders of 3 or 4. Some people report improvements when using this form of regularization. One of its appealing properties is that network cannot “explode” even when the learning rates are set too high because the updates are always bounded.

- Practice
 -
- Understanding
 -
- Cons
 -
- Early Stopping
 - Definition
 - stop the training at lowest validation loss (with training loss decreasing)
 \Rightarrow at the start point of overfitting
 - Practice
 - evaluate both train & val loss, saving models along the way
 \Rightarrow use the model corresponding to the start of overfitting
 - Understanding
 - at relatively early stage, weights are still relatively small
 (due to random initialization in $[0^-, 0^+]$)
 - Cons
 - couples task of optimizing loss and task of not overfitting
 \Rightarrow no longer one task at a time

Optimization

- Batch Descent
 - Practice
 - evaluate on entire training set; then update weights
 - Pros
 - largest optimization every time
 - Cons
 - greedy optimizing
 - slow & memory demanding on large dataset
- Stochastic Gradient Descent
 - Practice
 - shuffle data to have training set X_{train} , further split into $X_{\text{train}}^1, \dots, X_{\text{train}}^T$
 - train the net iteratively with $\forall t \in [1, T], X_{\text{train}}^t$
 i.e. one mini-batch for a gradient descent (weights update)
 - after training through all T batches, an epoch of training is finished
 $\Rightarrow 1 \text{ epoch} = 1 \text{ full scan of training set}$
 - Pros
 - faster: more weight upgrade over the same amount of data

- better chance to reach global change: not greedy anymore
 - more affordable for training in GPU memory
- ⇒ preferred choice
- Cons
 - observing noisy loss: not monotonically decreasing (but overall decreasing)
- Gradient Descent with Momentum
 - Definition: exponentially weighted average
 - calculate the gradient for weight update: $dW'_t = \beta dW'_{t-1} + (1 - \beta)dW_t$, where dW the original gradient
 - ⇒ average over past gradients with exponentially decaying weight, ⇒ for past $k \in [0, K]$ gradient, coefficient becomes $\beta(1 - \beta)^k$ (with $k = 0$ denoting current gradient)
 - bias correction: avoid slow start (due to: gradient dW_0 initialized to 0 & not enough gradients for averaging) ⇒ set $dW_t = \frac{dW_t}{1 - \beta^t}$ in the early stage (after starting stage, bias correction $\rightarrow 0$ for large t)
 - Implementation
 - approximation: weighted average over past $K = \frac{1}{1 - \beta}$ gradients due to $(1 - \epsilon)^{1/\epsilon} \approx \frac{1}{e}$, recognized as small enough ⇒ discard gradients with further exponentially small weights
 - apply element-wise multiplication on gradients and pre-calculated coefficient
 - sum up to be the gradient for weight update (include bias correction term if necessary, yet often omitted)
 - note: $dW'_t = \beta dW'_{t-1} + dW_t$ is another version, yet discouraged (coupling momentum β with learning rate α , as α needs to cooperate)
 - Understanding
 - averaging/smoothing out the regular oscillation in stochastic gradient descent ⇒ β popularly chosen to be 0.9 (averaging over last 10 gradients)
 - Pros
 - avoid some regular oscillation (slowing down the training & not true randomness)
- Root Mean Square Propagation (RMS prop)
 - Definition
 - compute $S_t = \beta S_{t-1} + (1 - \beta)dW_t^2$ (S_0 initialized to 0), where dW^2 the original gradient being element-wisely squared ⇒ exponentially weighted square of gradients
 - calculate the gradient for weight update $dW'_t = \frac{dW_t}{\sqrt{S_t}}$
 - Implementation
 - calculate S_t similarly (as an exponentially weighted average)
 - $\sqrt{S_t}$ becomes $\sqrt{S_t + \epsilon}$, where $\epsilon \rightarrow 0^+$ for mathematical stability
 - Understanding
 - for gradients with large variance in training ⇒ S_t large ⇒ $\frac{1}{\sqrt{S_t}}$ small ⇒ weighted less, hence stabilized (as it should be noisy & taking smaller step)
 - for gradients with small variance ⇒ weighted more, encouraged (as it should be on the "trend" towards optimum)

- Pros
 - recognize trend from noise via variance of their gradient \Rightarrow speedup training
 - auto-fixing learning rate for each weight given the recorded behavior (protect learning process from a too large learning rate)
- Adaptive Momentum (Adam) Optimization
 - Definition
 - compute $M_t = \beta_1 M_{t-1} + (1 - \beta_1) M_t$ as momentum
 - compute $S_t = \beta_2 S_{t-1} + (1 - \beta_2) dW_t^2$ as root mean square
 - apply bias correction on both: $M'_t = \frac{M_t}{1 - \beta_1^t}, S'_t = \frac{S_t}{1 - \beta_2^t}$
 - \Rightarrow calculate gradient for update $dW'_t = \frac{M'_t}{\sqrt{S'_t + \epsilon}}$, where $\epsilon \rightarrow 0^+$
 - Implementation
 - implement M_t, S_t as momentum and root mean square (popular choice: $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$)
 - do implement bias correction
 - Understanding
 - combine momentum with root mean square \Rightarrow for each weight
 - smooth out regular oscillation
 - encourage the trend & adapt learning rate given history record
 - Pros
 - effective for a large range of problem
- Learning Rate Decay
 - Definition
 - update learning rate $\alpha = \frac{1}{1+r \cdot e}$, where r the decay rate, e the epoch number
 - other decay formula:
 - exponential decay: $\alpha = r^e \cdot \alpha_0$, where α_0 the base learning rate
 - $\alpha = \frac{k}{\sqrt{e}} * \alpha_0$, where k a constant
 - Implementation
 - set learning rate for each epoch, or after some global steps
 - Understanding
 - fast learning at the beginning, more cautious when approaching the optimum \Rightarrow in order to finally converge

12.3.3 Cost

Probabilistic Cost

- Log Maximum Likelihood / Posterior
 - Definition
 - convert the logits into probability-alike prediction \Rightarrow then interpreted as predicted likelihood $p(\mathbf{y}|\mathbf{w}, \mathbf{x})$
 - bayesian regression $L = -\frac{1}{2} \sum_{\mathbf{y} \in \mathbf{Y}} (\mathbf{y} - \hat{\mathbf{y}})^2$
(for \mathbf{y} real number vector label, $\hat{\mathbf{y}}$ real number vector prediction)

- classification with logistic assumption $L = - \sum_{\mathbf{y} \in \mathbf{Y}} (\mathbf{y}^T \cdot \log \hat{\mathbf{y}})$
(t one-hot encoded label, $\hat{\mathbf{y}}$ one-hot encoded prediction)
- to use posterior with Gaussian distribution: add $L2$ regularization term

12.3.4 Layers


Prediction

- Sigmoid
- Softmax
 - Input
 - arbitrary input \mathbf{z}^L being logits, containing multiple multi-class predictions z^L
⇒ each prediction being the same dimension as one-hot encoded label
 - Output
 - probabilistic-alike prediction \mathbf{a}^L , with the same shape as the input (logits)
 - Operation
 - for K classes to predict ⇒ $\dim(z^L) = K$
 - for each dimension $k \in [1, K]$, compute $a_k^L = \frac{e^{(z_k^L)}}{\sum_{k=1}^K e^{(z_k^L)}}$
 - Implementation
 - vectorize the exponential computation $\hat{\mathbf{z}}^L = \exp(\mathbf{z}^L)$
 - compute normalization $N = \sum_{k=1}^K \hat{z}_k^L$
 - normalize as $\mathbf{a}^L = \frac{1}{N} \hat{\mathbf{z}}^L$
 - maximum likelihood with softmax: $L = \frac{1}{N} \sum_{\mathbf{y}} -\mathbf{y}^T \cdot \log \hat{\mathbf{y}}$,
where \mathbf{y} the one-hot encoded label, $\hat{\mathbf{y}}$ the prediction
⇒ easy gradients: $d\mathbf{z}^L = \hat{\mathbf{y}} - \mathbf{y}$, where \mathbf{z}^L the logits (vector)
 - Understanding
 - contrasting the hard-max function (non differentiable): $a_k = 1$ if $\arg \max_k(z)$; else 0
 - exponentially normalizing the output of arbitrary net into probabilistic form
(reduced to logistic for binary class i.e. $K = 2$)
⇒ generalize logistic prediction to K -class prediction
 - for maximum likelihood loss, only the gap with true class generate gradients
(due to one-hot encoding)
⇒ trying to predict the class true with higher probability
- Hierarchical Softmax
 - Input
 - arbitrary input \mathbf{z}^L being logits, considered as a flatten tree
 - Output
 - a tree structure, flatten into 1-D array
 - Operation

- for each node (if not leaf), a softmax to predict prob for its direct children
 \Rightarrow multiple softmax connected to various input nodes in \mathbf{z}^L
- Implementation
 - create a Bayesian network with a single root node (e.g. "obj")
 \Rightarrow each node being a conditional probability conditioned on its direct parent
 - for absolute probability of each class (represented by a node)
 \Rightarrow apply sum rule & product rule accordingly
- Understanding
 - NOT assuming mutual exclusion between class
 - \Rightarrow enable graceful degrade in classification
 e.g. can still recognize by $p(\text{animal})$, if failed with $p(\text{cat})$, $p(\text{dog})$, etc.
 - enable joint training with multiple datasets (involving classification)
 e.g. YOLOv2(YOLO9000)
- Normalization

Convolution Layer

- Convolution 2D
 - Input
 - spatially 2D feature maps, usually with multiple channels
 - Operation
 - given hyperparameter: kernel/filter size, stride, padding
 - kernel (weights) structured as a matrix (for each input channel)
 - input maps padded if required
 - an element-wise weighted sum on the spatially corresponding position
 - sum across channels: sum over kernel output from each channel + optional bias
 - kernel strides spatially across the image, with stride along each axis defined
 \Rightarrow to calculate **1** channel in the output feature maps
 \Rightarrow for multi-channels output: multiple sets of kernels



- activation then taken after convolution operation, element-wisely
- Padding
 - reason
 - prevent output feature maps from spatially shrinking
 - prevent info lost on the edge&corner of image
 (compared to the central part of feature map, multiplied less with the kernel)
 - convention: 0-padding on both directions of an axis

- valid conv: no padding
 - same conv: pad so that output size same as input size
- Kernel Size
 - odd square matrix: avoid asymmetric padding & a central pixel for filter location (even becomes a convention)
- Stride
 - the step for kernel to move its location as striding over feature maps
 - kernel striding over the edge (after padding): NOT convoluted
- Output
 - a feature maps with defined output channel & size (output channel depends on the number of sets of kernels)
 - on a $2D$ square feature map, with padding at each edge p , stride on all axes s kernel size $k \times k$, input size $i \times i$, output size $o \times o$

$$\Rightarrow o = \left\lfloor \frac{i + 2p - k}{s} \right\rfloor + 1$$
- Understanding
 - computer vision filter: learn the weights in filters, instead of hand design
 - \Rightarrow guided by data statistics
 - weights in the l^{th} conv layer: $n_c^{l-1} \times k \times k \times n_c^l$, where n_c the channel number (to output n_c^l channels with n_c^{l-1} input channels from previous layer)
 - \Rightarrow invariant to the input size (number of trainable variables fixed on design)
 - \Rightarrow less weights (then dense layer), more generalizability, hence less overfitting
 - sharing weights spatially: apply same weights over the whole space
 - \Rightarrow NO need for special design at each location
 - \Rightarrow as need to handle spatial variance in processing images
 - sparse connection: output connected only to the local input \Rightarrow as a high-pass bandwidth
 - \Rightarrow robust to spatial variance
- Back Propagation
- Implementation
 - implement cross-correlation instead of convolution (skip the flipping operation: as their results are symmetric) yet, convolution is associative, due to the flipping
- 1×1 Convolution
 - Input
 - multi-dimension feature maps with multiple channels
 - Operation
 - integrate channels at each spatial location together (a weighted sum with bias, as conv definition)
 - Output
 - feature maps with same dimensions, but different channels
 - Understanding
 - shrink the number of channels
 - add more non-linearity & info combination (more representability)
 -


- Atrous Convolution
- Deconvolution
- Convolution Implementation
 - Understanding
 - span the original input in the extra dimension, by duplicating the input
 \Rightarrow each input location generate its contribution at corresponding output location
 (collected by the sum op in matrix multiplication)

Pooling Layer

- Normal Pooling
 - Input Data
 - feature maps
 - Operation
 - given hyperparameter: kernel size, stride (usually no padding)
 - compute the max/average of the elements covered by kernel
 - kernel strides along each axis over the feature maps (like conv)
 \Rightarrow does NOT change the channel
 \Rightarrow one kernel per channel (compared to conv)
 - Output
 - a downsampled feature maps
 - given a 2D feature map with kernel size $k \times k$, strides along each axis $s \times s$
 input size $i \times i$, output size $o \times o$
 $\Rightarrow o = \lfloor \frac{i-k}{s} \rfloor + 1$ (same as conv)
 - Understanding
 - downsampling the feature maps (NO weights to learn)
 \Rightarrow if desired features detected anywhere, represent it by local max/average
- Unpooling
- Spatial Pyramid Pooling (SPP)
- Region of Interest Pooling (RoI Pooling)
 - Input
 - feature maps from CNN
 - RoIs i.e. proposal region (from selective search etc.) projected on feature map
 - Operation
 - divide each RoI with grid of desired size (proportional to the RoI size)
 - max pooling from each cell \Rightarrow single-size SPP for each RoI
 - Output
 - a fixed size feature maps for each RoI
- Probabilistic Max Pooling


RNN Layer

- Overview
 - Input Data
 - sequence data: includes time/precedence (conventionally arrived from left to right)
 - for each time step, data can be vector, feature maps, etc...
 - RNN Cell
 - consume the input of current time step & the hidden state from last time step (hidden state usually initialized to $\mathbf{0}$)
 - calculate a hidden state at each time step
 - Operation
 - RNN cell at time t , calculate hidden state (activation) $h^t = g_h(w_h[h^{t-1}, x^t] + b_h)$, where $g_h(\cdot)$ the activation function, $g_h = \tanh$ by convention $[h^{t-1}, x^t]$ the concat of h^{t-1} (hidden state of time $t-1$), x^t (input at time t)
 - expose its hidden state at each time steps
 - calculate its output $y^t = g_y(w_y h^t + b_y)$, where $g_y = \sigma, \text{softmax}$ or *identity*



- Types of RNN Mapping
 - many-to-one: encoder scans through the input, only the last output considered
 - one-to-many: decoder with single input x^1 , take $x^t = y^{t-1}$, till $y^{t'} = \text{stop}$
 - many-to-many: a many-to-one encoder, followed by a one-to-many decoder
 \Rightarrow able to map between various length
- Back Propagation through Time
 - unroll the recurrent operation into a sequential network with length T
 - given the loss for each time step $L^1, \dots, L^T \Rightarrow L = \sum_{t=1}^T L^t$
 - for time $t = 1, \dots, T-1$,
$$\frac{\partial}{\partial a^t} L = \frac{\partial}{\partial a^t} L^t + \sum_{t'=t+1}^T \frac{\partial L^{t'}}{\partial a^{t+1}} \frac{\partial a^{t+1}}{\partial a^t} = \sum_{t'=t}^T \frac{\partial}{\partial a^t} L^{t'}$$
- Truncated Back Propagation through Time
- Challenge
 - bad at modeling longterm dependency due to gradient vanishing problem
 \Rightarrow loss at late time needs to go through multiple activations to the early time (similar to the deep plain net, after unrolled)
 \Rightarrow loss at late time are hard to affect weights when evaluated at early time (i.e. larger the t' , smaller the $\frac{\partial}{\partial a^t} L^{t'}$)
 \Rightarrow hard to find out error in late time due to observation in early time
 \Rightarrow hard to represent longterm dependency (e.g. Car...is fast vs. Cars...are fast)
 - easily affect by local dependency (as longterm dependency lost)
 - gradient exploding, due to multiple / too many updates on the same weights (solved by gradient clipping)
 - hard to converge, due to fluctuating gradient in an unroll (noisy intermediate stage)
 \Rightarrow initial short sequence to overcome plateaus, then long sequence for dependency

- Understanding
 - sharing weight across time: same weights used on each time step
 - ⇒ solve various length input by applying weights recurrently on each of them
 - information early in the sequence reserved & passed through in the hidden state
- Long Short Term Memory (LSTM)
 - Memory Cell
 - c^t the memory maintained by LSTM cell at time t
 - Gates
 - forget gate $G_f = \sigma(w_f[h^{t-1}, x^t] + b_f)$
 - input gate $G_i = \sigma(w_i[h^{t-1}, x^t] + b_i)$
 - ⇒ together control the memory update (how past¤t info fused)
 - output gate $G_o = \sigma(w_o[h^{t-1}, x^t] + b_o)$
 - ⇒ control the generation of hidden state
 - Fusing Info
 - propose candidate $\hat{c}^t = \tanh(w_c[c^{t-1}, x^t] + b_c)$ for memory update
 - update memory as $c^t = G_f c^{t-1} + G_i \hat{c}^t$
 - Hidden State (Activation)
 - generate as $h^t = G_o \cdot \tanh(c^t)$



- Understanding
 - easy to learn an identity mapping $c^{t-1} \rightarrow c^t$
 - ⇒ memory (info) generated early can last for long term
 - ⇒ better model the long-term dependency
- Gated Recurrent Unit (GRU)
 - Memory Cell
 - c^t the memory maintained by GRU cell at time t
 - Gates
 - relevance gate $G_r = \sigma(w_r[c^{t-1}, x^t] + b_r)$, a mask $\in [0, 1]$
 - ⇒ control how memory candidate proposed (fusion of last memory & input)
 - update gate $G_u = \sigma(w_u[c^{t-1}, x^t] + b_u)$, a mask $\in [0, 1]$
 - ⇒ control how memory update happen (fusion of last memory & candidate)


(two gates computed with the same input, though with different weights)

◦ Fusing Info

- propose memory candidate $\hat{c}^t = \tanh(w_c[G_r c^{t-1}, x^t] + b_c)$ for the update
 \Rightarrow decide whether the previous memory useful (relevant) with the context x^t
- update memory as: $c^{t+1} = G_u \hat{c}^{t+1} + (1 - G_u)c^{t-1}$
 \Rightarrow decide how the memory updated & remained

◦ Hidden State (Activation)

- $h^t = c^t$



where,
 joint arrows denotes a concatenation
 $h_{t-1}, \tilde{h}_t, h_t$ the previous/candidate/current memory;
 r_t, z_t the relevance/update gate

◦ Understanding

- single gate control the generation of current memory
- memory directly as hidden state
- \Rightarrow less weights, simpler structure \Rightarrow faster
 (basic logic inherent from LSTM \Rightarrow similar performance)


• Bidirectional RNN (BRNN/BiRNN)

◦ Structure

- one RNN layer scanning as $t = 1 \rightarrow T$, hidden state exposed as h_1^t
- another RNN layer scanning from $t = T \rightarrow 1$, hidden state exposed as h_2^t
- generate output $y^t = g(w_y[h_1^t, h_2^t] + b_y)$
 (concat all hidden states at the same time step from each RNN)

◦ Understanding

- account for the info from both previous & latter time step
 \Rightarrow global context info acquired
- cons: need the entire input sequence before processing
 \Rightarrow NOT the case in real-time speech recognition etc.



(using LSTM cell)

• Convolutional LSTM (ConvLSTM)

12.3.5 Blocks

Residual Block

- Structure
 - Definition
 - for the $l + 2$ layer, $a^{l+2} = g(z^{l+2} + a^l)$,
where $g(\cdot)$ the activation, $z^{l+2} = W^{l+2}a^{l+1} + b^{l+2}$ from layer $l + 1$
 - Main Path
 - the usual passing of a^l to layer $l + 1$, then the z^{l+2} at layer $l + 2$
 - Skip/Passthrough Connection (Shortcut)
 - $s(x^l)$: the passing of x^l , from layer l directly to the layer $l + 2$
 - different spatial size in $l, l + 2$: adjust x^l by
 - padding: pad to the size in layer $l + 2$
 - spatial stack: merge local channels into single channel by stacking
(e.g. $26 \times 26 \times 512 \rightarrow 13 \times 13 \times 2048$)
 - different channel size $l, l + 2$: adjust x^l by linear linear transformation
(a 1×1 conv / matrix with trainable or fixed weights)
 - Join
 - add the result from two paths $\Rightarrow a^{l+2} = g(z^{l+2} + s(a^l))$
 - or, concat the input $\Rightarrow a^{l+2} = g([z^{l+2}, s(z^l)])$
- Understanding
 - Guaranteed Baseline
 - with ReLU as activation, easy to learn identity function
 \Rightarrow layer $l + 2$ only need to make $z^{l+2} = 0$ (as weights & bias initialized near 0)
 - \Rightarrow deeper net can easily guarantee to be at least as good as its shallow version
(then search for luck to surpass baseline)
 - Ensemble
 - contains much more shallow paths than real deep paths
 \Rightarrow ensemble those shallow paths at various level
 - More Gradient
 - gradient more easily passed to the early layers
(as shortcuts not attenuating gradients)
 - \Rightarrow early layers settle down faster \Rightarrow late layers get a more consistent input

Inception Block

- Basic Inception
 - Input
 - feature maps with multiple channels
 - Operation
 - op1 = 1×1 conv
 - op2 = 1×1 conv, 3×3 conv with same padding
 - op3 = 1×1 conv, 5×5 conv with same padding
 - op4 = max pooling with same padding and stride 1, 1×1 conv

- channel concat: concatenate the output channel from each op
(as output size ensured to be the same in each op)
- Output
 - feature maps with the same spatial size as input
- Understanding
 - 1×1 conv to shrink
 - less computation for afterwards larger kernel (e.g. $3 \times 3, 5 \times 5$)
 - prevent output of other ops from being overwhelmed by pooling
 - enable network to learn the desired combination of info
(instead of predefined hyperparameter)
⇒ more representability
- Xception
 -

12.4 Architectures

12.4.1 Convolutional Networks

Classic Convolutional Networks

- Overview
 - Convolution Part
 - one or multiple "same" conv layer(s) with stride $s = 1$
 - followed by a max (rarely average) pooling
 - apply on input & repeat on feature maps afterwards
⇒ usually decreasing spacial size (width, height), increasing channel
 - Fully Connected (Dense) Part
 - apply flattening / average pooling on last feature maps
⇒ generate a single vector as input
 - apply fully connected layers & repeat for 1 2 times
⇒ usually with decreasing size
 - finally output prediction probability
 - Understanding
 - trainable weights are mostly from dense layer
⇒ conv layers, though large in number, contains far less weights
 - conv stride $s = 1$ (compare to $s > 1$)
⇒ decouple downsampling into the pooling
⇒ account for more info/possibility before downsampling
(also trade-off between required computability)
- Receptive Fields
 -
- VGG-16
 - Building Blocks
 - fixed conv: 3×3 kernel, stride $s = 1$, same padding, ReLU activation
 - fixed pooling: 2×2 kernel, stride $s = 2$, max pooling

- Notation
 - [conv64] to denote a conv layer with 64 output channels
 - pool to denote max pooling
 - [fc4096] to denote a fully connected layer with output vector of length 4096
- Structure
 - input image: $224 \times 224 \times 3$
 - $([\text{conv}64] \times 2, \text{pool}) \rightarrow ([\text{conv}128] \times 2, \text{pool}) \rightarrow ([\text{conv}256] \times 3, \text{pool}) \rightarrow ([\text{conv}512] \times 3 + \text{pool}) \times 2$
 - last feature maps (7×512) flatten into a input vector of length 4096
 - $[\text{fc}4096] \rightarrow [\text{fc}4096] \rightarrow \text{logits} \rightarrow \text{softmax prediction for 1000 categories}$
- Understanding
 - fixed conv & pooling operation
 - \Rightarrow few hyperparameters, yet large in parameters (~ 138 million)
 - decrease spatial size by 2 & increase channel by 2 on each step (till small or deep enough e.g. 7×7 size, 512 channels)
- ResNets
 - Building Blocks
 - residual block with addition
 - batchnorm applied in the conv(s) inside residual block
 - Structure
 - stack of residual block, for hundreds, even thousands of layers
 - Understanding
 - ensemble nets with shallow & medium depth \Rightarrow essentially a wide net (as the real deep net takes only few paths)
 - \Rightarrow avoid the training of real deep net (instead of solving it)
- Inception Net
 - Building Blocks
 - inception block
 - Structure
 - stack of inception blocks
 - two auxiliary loss from hidden layers

Fully Convolutional Networks (FCN)

- FCN for Classification & Regression
 - Convolution as Flattening
 - given input feature maps $i \times i \times c$, with c channels
 - perform valid conv with kernel $i \times i \times o$
 - \Rightarrow output size $1 \times 1 \times o$
 - 1×1 Convolution as Fully Connected Layer
 - for flatten vector $1 \times 1 \times i$, perform 1×1 conv with kernel size $1 \times 1 \times o$
 - \Rightarrow output size $1 \times 1 \times o$, with same computation as fc layer
 - Use Case
 - object detection

12.4.2 Recurrent Neural Network

Classic RNN

- Overview
 - Stacked RNN Layers
 - the output of layer l becomes the input for layer $l + 1$
 \Rightarrow RNN scanning the output of previous RNN layer
 - 3 ~ 5 layer considered deep: as RNN can be unrolled into a deep plain net
 - Encoder-Decoder RNN
 - encoder RNN scans the input sequence, last hidden layer as sequence encoding
 - decoder RNN takes encoding as initial hidden state, unrolled the output sequence

Attention

- Overview
 - Goal
 - overcome long-term dependency, via considering all input with different attention
 - Practice
 - a bi-RNN encode rich input context at each hidden state: h^1, \dots, h^{T_x}
 - global context after attention at time t : $c^t = \sum_{i=1}^{T_x} \alpha_i^t h^i = [h^1, \dots, h^{T_x}] \alpha^t$,
 where hidden states h^1, \dots, h^{T_x} and attention α^t column vectors
 \Rightarrow a weighted sum over all hidden states as context
 - a small (1-layer) net mapping $[c^{t-1}, h^i] \xrightarrow{\text{dense}} \text{logits} \xrightarrow{\text{softmax}} \alpha_i^t$
 \Rightarrow attention on h^i depends on previous global context c^{t-1} & h^i itself
 (softmax to ensure $\sum \alpha^t = 1$)
 - a decoder RNN taking c^t as input, until stop sign generated
 (as c^t can be calculated infinite times)
 - Understanding
 - use the sequence of hidden states as encoding
 (as sentence encoded in last hidden states can be bias & have info lost)
 \Rightarrow account input sequence (with different attention at different places) when generating output sequence
 - quadratic cost: attention α a $T_x \times T_y$ matrix to be calculated
- Attention on Image
 - Goal
 - focus on correct spatial context when generating corresponding caption/sequence
 - Practice
 - calculate weights for a sum over all spatial location of the image encoding
 \Rightarrow calculate a 2-D mask over a 3-D tensor

Transformer

12.4.3 Encoder-Decoder Architecture

Basic

- Encoder
 - Functionality
 - downsample/encode input into rich feature maps/vectors
 - Implementation
 - visual input: CNN backbone
 - natural expression input: RNN backbone
- Decoder
 - Functionality
 - upsample/decode rich feature maps back to the original size
 - actually, impose requirement onto the encoder
 - Implementation
 - visual output: CNN backbone
 - natural expression output: RNN backbone
- Connection
 - Functionality
 - combine high level information with low level information
 - image \rightarrow image: outline refinement ...
 - language \rightarrow language: sentence style capturing
 - Implementation
 - concatenation

Extension

- Siamese Network
 - Input
 - image pair (p_1, p_2)
 - Structure
 - CNN as feature encoder
 - \Rightarrow same CNN process each image, yielding (f_1, f_2) as extracted feature
 - jointly process concat $[f_1, f_2]$
 - Variants
 - joint process from the beginning
 - \Rightarrow concat $[p_1, p_2]$, then CNN till the end
 - \Rightarrow generally **better performance**, especially when detailed structure are vital
- Multiple Encoder (Pseudo Siamese Network)
 - Understanding
 - project different information into the same space
(jointly process those information after feature projection/extraction)

- Multiple Decoder
 - Functionality
 - impose multiple requirements to the encoder (via auxiliary loss)
 - ⇒ multi-task training (multiple output branch)

12.4.4 Generative Network

Generative Adversarial Network

12.5 Computer Vision

12.5.1 Objects Detection

Problem Formulation

- Input Data
 - Spatial Information
 - image maps as 3-D tensor: width, height, channel (rgb, etc.)
 - multi/stereo -images for 3-D detection
- Goal
 - Bounding Box Prediction
 - localization as regression task on box attributes (x, y, h, w)
 - classification on object classes & object existence
 - Landmark (Key Point) Prediction
 - landmarks: the coordinates of key points (of each type) in image
 - label of landmark should be consistent across image
 - output real number as regression task
 - ⇒ used in pose detection, bounding box detection, etc.
- Understanding
 - multi-object localization & classification

Common Postprocessing

- Non-max Suppression
 - Input
 - multiple predicted bounding boxes, with probability of existence (p_e) relatively large
 - Goal
 - clean up & account for duplicate bounding boxes on the same object
 - ⇒ for each (predicted) object, finalize prediction to be a single bounding box
 - Naive Operation
 - choose the bounding boxes with highest (maximal) p_e (for each object classes)
 - remove those bounding boxes whose IoU with it are large
 - ⇒ remove (suppress) bounding boxes with large overlap
 - repeat until all bounding boxes have a low IoU with each other
 - ⇒ remaining boxes are the final predictions

Region Proposal Based Approaches

- Sliding Window Detection
 - Bounding Box Proposal
 - slide the window with various size, across the image
⇒ propose bounding boxes with predefined size & location
 - feed the window into CNN for classification
⇒ CNN as classifier, window as bounding box
 - Fast Implementation: Sliding Window as Convolution
 - implement CNN classification as FCN
⇒ as conv independent from input size
 - run directly the CNN on the input image (instead of on each sliding window)
⇒ output size $m \times n \times k$,
where $m \times n$ the total number of sliding windows on the image, k the class number
 - ⇒ one times running for all sliding windows
(as sliding window essentially is a crop from image)
 - Understanding
 - though with fast implementation, still not promising regarding result
⇒ sliding window propose a fixed set of window
⇒ fail if not matching window size; or missing location/rotation
- Region Proposal
 - Goal
 - effectively propose a various bounding boxes on the image for potential focused object
⇒ recall all focused object with preferably fewer boxes (e.g. $\sim 1k$)
 - Input
 - segmentation mask by a segmentation network
⇒ to place bboxes on high-probability blobs
- R-CNN
 - Overview
 - region proposal by selective search
 - CNN to predict class & a bounding box (to refine)
- Fast R-CNN
 - Overview
 - based on r-cnn
 - convolution implementation to predict all proposed regions
(similar to fast sliding window)
- Faster R-CNN
 - Overview
 - based on fast r-cnn
 - use CNN for segmentation & region proposal (faster than selective search)


One-stage Detection

- You Only Look Once (YOLO)
 - Preprocessing
 - grid the input image into disjoint $S \times S$ cells
(apply usually fine grid, e.g. $S = 19$)
 - assign each object to a cell by its central point
⇒ try ensuring maximal 1 object in a grid
 - Bounding Box Encoding
 - $conf = p(obj) \times \text{IoU}_{\text{pred}}^{\text{label}}$ confidence of current box containing an object
⇒ 0 desired if none, IoU between pred and label desired if exist
where, $p(obj)$ the probability of object existence in the CELL
 - $x, y \in [0, 1]$ the box center in the cell, normalized by cell size
⇒ easier to learn, as dense layer not confused by varying locations
 - $h, w \in [0, 1]$ the box size normalized by image size
⇒ easier to learn for dense layer
 - Classification Encoding
 - p_1, \dots, p_C the conditional probability $p(\text{class}_c | obj)$ for object inside the CELL
i.e. the class probability of object, given object existing & assigned to the cell
 - Inference
 - for each cell, predict both $[p_1, \dots, p_C]$ and B bounding boxes $[conf, x, y, w, h]$
⇒ final prediction as a $S \times S \times (C + 5B)$ tensor
 - $p_c \cdot conf = p(\text{class}_c | obj) \cdot p(obj) \cdot \text{IoU}_{\text{pred}}^{\text{label}} = p(\text{class}_c, obj) \cdot \text{IoU}_{\text{pred}}^{\text{label}}$
where p_c from the cell, $conf$ from each box in the cell
⇒ combine the decouple pred for various class-specific boxes prediction
 - non-max suppression to fix multiple detection box (at test time)
 - Structure
 - 24 conv layer downsampling from $448 \times 448 \times 3$ to $S \times S \times 1024$
 - followed by: →dense with dropout→ 4096 →dense→ $S \times S \times (C + 5B)$
(dropout to prevent co-adaption, i.e. dependent units between 2 dense layers)
 - leaky ReLu $a = \max(z, 0.1z)$ used, except for linear activation in final layer
 - Optimization Goal
 - $$\begin{aligned} \text{loss} = & \lambda_{\text{obj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbf{1}_{ij}^{\text{obj}} [(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2] + \\ & \lambda_{\text{obj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbf{1}_{ij}^{\text{obj}} [(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2] + \\ & \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbf{1}_{ij}^{\text{obj}} (conf_{ij} - \hat{conf}_{ij})^2 + \lambda_{\text{none}} \cdot \mathbf{1}_{ij}^{\text{none}} (conf_{ij} - \hat{conf}_{ij})^2 + \\ & \sum_{i=0}^{S^2} \mathbf{1}_i^{\text{obj}} \sum_{c=0}^C (p_{ci} - \hat{p}_{ci})^2 \end{aligned}$$
 - $\lambda_{\text{obj}} = 5$ to up-weight box with obj, $\lambda_{\text{none}} = 0.5$ for empty box
 - balance the dataset: avoid large num of empty box pushing $conf$ pred to 0
 - emphasize on predicting box with object: adjust the box pred x, y, w, h
 - $\mathbf{1}_{ij}^{\text{obj}}$ the indicator: 1 if obj exists in j box of i cell, similar for $\mathbf{1}_{ij}^{\text{none}}$

- minimize $\sqrt{w_i}, \sqrt{h_i}$: mitigate various box size into similar scale
(as same error occurred in small box matter more than that in big box)
- p_{ci} the classification prob for i cell (only trained if it contains obj)
- Training
 - optimize the sum-squared error (due to its simplicity)
 - for each cell: correct its classification if it contains obj
 - for each box:
 - increase *conf* & adjust box pred if it overlaps the most with label of its cell
 - decrease *conf* for other box (not responsible for having obj)
 - pre-training: pretrained first 20 conv layers on classification dataset
 - data augmentation: random scaling size, adjusting exposure and saturation
- Limitation
 - use only coarse feature to predict box
(due to multiple downsampling & produced by dense layer)
 - NOT able to handle: multiple objects of the same type in a cell
(NOR when same-type objects more than total pred boxes)
⇒ suffer from small objects in group (e.g. birds)
 - struggle in box localization
- Understanding
 - bounding box localization & classification as regression problem
(predicted totally by a dense layer)
 - end-to-end optimizing directly the detection score
(vs. separate modules: region proposal, bounding box localization, classification)
 - global information for each bounding box localization & classification
(instead of taking only local info to classify boxes)
⇒ less false alarm (mistake background as object)
 - generalizable representation for detection
(as feature extraction embedded in the network)
⇒ able to generalized to art work
 - much less predicted box (compared with 2000 boxes in R-CNN framework)
 - fast ⇒ ensemble with (fast) R-CNN framework with no overhead
⇒ mitigate the false alarm in R-CNN framework
- YOLOv2, YOLO9000
 - Preprocessing
 - no more explicit grid, implicitly set by downsampling factor
(similar to spatial pyramid pooling)
 - select N_A anchor boxes: covering most of the interested objects
(e.g. tall-thin box for pedestrian, low-wide box for car)
 - assign the objects to a tuple (cell, anchor/prior box)
 - assign to cell by its central point ⇒ try ensuring maximal N_A object in a cell
 - assign to one from N_A anchor boxes depending on its IoU with all N_A boxes
(NOT assigned, if all $\text{IoU}_{\text{anchor}}^{\text{label}} < 0.5$)
 - Anchor Box Selection (Prior)
 - run K-means on all label box in train set
with distance $\text{dist}(\text{box}, \text{centroid}) = 1 - \text{IoU}(\text{box}, \text{centroid})$
 - use cluster centroid as the anchor box
⇒ for anchor box $a \in A, a = [w_a, h_a]$

- num of anchor box (centroid) N_A reflect recall \leftrightarrow model complexity trade-off
- Bounding Box Encoding
 - $t_x, t_y \in (0, 1)$ the box center in the cell, normalized by cell size
 \Rightarrow true pred center $x = c_w \sigma(t_x) + c_x, y = c_h \sigma(t_y) + c_y$ where c_x, c_y, c_w, c_h the cell
 \Rightarrow easier to learn, as YOLO
 - t_w, t_h coefficient to adjust anchor box size
 \Rightarrow true pred size $w = w_a e^{t_w}, h = h_a e^{t_h}$ where w_a, h_w the anchor box size
 \Rightarrow more robust, avoid influence from anchor prior
 - p_1, \dots, p_C classification pred \Rightarrow decouple from cell, each box a classification
- Classification with Hierarchical Encoding
 - tree-structure for multi-label: a hierarchical softmax
 (hence, NOT assuming mutual exclusion between class)
 - lookup WordNet for class label & its path to the selected root "physical object"
 (if multi-path, choose the one with least new word involved)
 \Rightarrow build a WordTree
 - multiple softmax: one for each set of direct children of a node (if not leaf)
 e.g. at node "dog": predict $p(\text{malamute dog}|\text{dog}), p(\text{terrier dog}|\text{dog}), \dots$
 \Rightarrow pred a Bayesian network, according to WordTree
 - $\Rightarrow t_o$ objectiveness by summing conditional prob accordingly
 \Rightarrow true objectiveness: $\text{conf} = p(\text{obj}) \cdot \text{IoU}_{\text{pred}}^{\text{label}} = \sigma(t_o)$
- Inference
 - for each cell, N_A bounding boxes: [spatial info[t_x, t_y, t_w, t_h], a flatten WordTree[p_1, \dots, p_C]]
 \Rightarrow final prediction as a $\frac{I}{32} \times \frac{I}{32} \times N_A(4 + C)$ tensor (each box a WordTree!)
 - objectiveness of each box: from WordTree of each box
 - box classification: choose highest confidence path at every node
 (till a specified threshold or a leaf node reached)
 - non-max suppression to fix multiple detection box (at test time)
- Structure
 - fully conv net: downsample by 32, directly predict $\frac{I}{32} \times \frac{I}{32} \times (C + 4N_A)$ tensor
 - passthrough layer with concat (1% \uparrow)
 e.g. spatially stack $26 \times 26 \times 512$ into $13 \times 13 \times 1024$ before concat
 - batch norm in ALL conv layer ($> 2\%$ mAP \uparrow)
- Training
 - pre-training on classification with data augmentation
 (e.g. random crop, rotation, hue, saturation and shifting shift)
 - more time for model fine-tuning on detection dataset ($\sim 4\%$ mAP \uparrow)
 as model takes time to adjust to resolution change
 - multi-scale training s.t. model able to cope with varying resolution
 - joint training: mix detection & classification dataset ($\sim 5\%$ \uparrow)
 \Rightarrow for detection label, full loss available
 \Rightarrow for classification: only classification & objectiveness loss considered
 - select responsible pred box by the largest objectiveness
 - objectiveness loss: considered if t_o of responsible box $<$ threshold (0.3)
 - classification loss: only nodes in its WordTree above current label considered
 with $p(\text{physical object}) = 1$ for label
 - balancing joint dataset: oversampling detection dataset
 \Rightarrow s.t. detection data : classification data = 1:4
- Limitation

- each box a classification: too much output channel
- still, detectable objects number restricted by anchor box num in a cell (e.g. group of small objects)
- Understanding
 - still, fast & global context for box detection
 - auto-select prior, better than handpicked (examined by using anchor box directly as prediction)
 - adjustable speed↔accuracy trade-off via input resolution (enabled by being fully conv & multi-scale trained)
 - ⇒ larger resolution, slower, more accurate
 - joint training with WordTree: multi-task (detection + classification) learning
 - ⇒ enrich classes to detect & more robust in class-detection
 - ⇒ graceful degrade on detecting new/unknown data
 - decoupled localization & classification
 - ⇒ able to handle some multiple same-type objects in same cell
- YOLOv3
 - Preprocessing
 - same as YOLOv2
 - Bounding Box Encoding
 - t_x, t_y, t_w, t_h : same as YOLOv2
 - t_o the objectiveness, modeling directly $p(\text{obj})$, instead of $p(\text{obj}) \cdot \text{IoU}$
 - p_1, \dots, p_C the classification of obj: conditional prob $p(\text{class}_c | \text{obj})$, as YOLO
 - Classification Encoding
 - independent logistic regression for each class (instead of hierarchical softmax)
 - ⇒ multiple binary classification (NOT assuming mutual exclusive classes)
 - ⇒ multi-label approach (box may contain more than one class)
 - Inference
 - for each cell, N_A bounding boxes $[t_x, t_y, t_w, t_h, t_o, p_1, \dots, p_C]$ (N_A the num of anchor box)
 - ⇒ final prediction as a $\frac{I}{N^l} \times \frac{I}{N^l} \times N_A(C + 5)$ tensor where N^l the size of downsampled feature map (at pred layer l)
 - Structure
 - residual blocks & batch norm as YOLOv2
 - 2D encoder-decoder structure with concat connection, with 3 pred branches
 - ⇒ final num of pred boxes $\sum_{l \in \text{pred layer}} (N^l \times N^l \times N_A)$
 - Optimization
 - sum of the binary cross-entropy for classification loss
 - Training
 - pretrain backbone (feature extractor) on classification
 - multi-scale training & data augmentation as YOLOv2
 - Failed Approach
 - box x, y as an offset to anchor box position x_a, y_a (stability decreased)
 - linear regression for x, y (worse than bounded by logistic)
 - focal loss (already solved by decoupled $p(\text{obj})$ & $p(\text{class}_c | \text{obj})$)



- dual IoU threshold in assigning object (need more tuning for stabilized model)
- Understanding
 - real-time applicable detector
 - detection at multiple scale, with encoder-decoder structure
 - decent localization + great box classification
 - ⇒ emphasize on box classification, since human insensible to IoU change

12.5.2 Object Tracking

Problem Formulation

- Input Data
 - Raw Data
 - time-series observation from sensor ⇒ need to embed detection
 - Object Representation
 - bounding box: 2/3-D size, with pose, motion, velocity, etc.
 - point model: centroid point with attributes denoting an object (including bbox)
 - silhouette: for non-rigid object
 - articulate: for articulated model, e.g. human skeletal model
 - Object Existence
 - probability density for existence
 - probability of existence as a feature in representation info (objectiveness in detection)
- Goal
 - Trajectory
 - recover the true trajectory of the object (including current location, potentially future prediction)
- Overview
 - Online Tracking

- trade off speed - model complexity, as train on arriving frames
 - adaptive, as accounting the history info of a track
 - ⇒ may provide more info (e.g. covariance matrix from Kalman filter)
- Near-Online Tracking
 - look ahead a fixed number of frames, before consolidating a method
- Offline Tracking
 - learn similarity functions between frames offline
 - fast, as no online training needed
 - NOT adaptive ⇒ failed in large appearance change, e.g. occlusion etc.
 - batch method: generate the track after all frames examined
- Tracking-by-Detection
 - detect target(s) in each frame; link target into track
 - ⇒ as a two-stage problem
- Challenge
 - Appearance Change
 - shape&color can be trap: changing clothes etc.
 - Long Term Dependency
 - re-identification: people leaving and re-entering the scene
 - hot standby surveillance
- Evaluation Metrics
 - Percentage of Time Step with Success
 - overlap success: IoU of pred-label bounding box \geq threshold = 0.25
 - orientation success: diff of pred-label yaw angle \leq threshold = 10°
 - other success...
 - ⇒ review the overall estimated quality
 - Normalized Cumulative Sum of Success vs Normalized Time
 - a plot ⇒ review tracking quality over time
 - Trajectory Difference
 - compare the similarity of pred-label trajectory
 - (may consider abrupt change, slowly drift, and etc.)

Detection in Tracking

- Single-frame Detection
 - Bounding Box Detection
 - YOLO, R-CNN, etc.
 - Point Detection
 - detect landmarks
 - Background Modeling
 - segmentation, ...
- Temporal Detection
 - Optic Flow

- able to represent non-rigid, deformable object
- yet, may failed in moving foreground (e.g. birds, fog, smoke...)
- Motion Detection
- Orientation
 - frame differencing of location, ...
- Background Modeling
 - adaptive background, ...

Single-Object Tracking

- Kalman Filter
 - Understanding
 - predict covariance \Rightarrow better gating in association
 - cooperate with noise
 - \Rightarrow able to modified to account association uncertainty
- Particle Filter
- Interacting Multiple Model
 -
- GOTURN:
 - Detector
 - CNN
 - Tracker
 - no explicit tracker (not modeling temporal info)
 - Inference
 - crop t^{th} image: centered at current bounding box with context
 - crop $t + 1^{\text{th}}$ image: centered at bbox in t , doubled size (gating)
 - encode both image via CNN detector; concat encoding & feed to dense layer
 - regress the bbox in $t + 1^{\text{th}}$ crop
 - Tracking
 - arbitrary target object selected for tracking (treated as ground truth)
 - 1-step prediction of location, then predict further on...
 - Training
 - train with consecutive image pair from sequences
 - $L1$ loss for exact match
 - data augment: motion noise applied on the 2^{th} image; random crop on the 1^{st} (distribution of motion noise chosen from cross-validation)
 - Understanding
 - 2-frame model, 1-step prediction, no explicit temporal tracking
 - \Rightarrow net may learn to locate the nearest similar object
 - real-time timing (100fps)
- Re³: Real-time Recurrent Regression Network
 - Detector

- CNN to extract multi-scale representation from image (more descriptive info e.g. human in red/blue shirt)
 - ⇒ a siamese net as encoding
- Tracker
 - two-layer, factored LSTM, taking tracker input at both layer
 - ⇒ longer dependency & more complex object transformation with 2 layers
 - hidden state as tracker state
 - ⇒ forward-prop to update (no training)
- Inference
 - crop t^{th} image: centered at current bounding box, extended to twice of box size
 - crop $t + 1^{\text{th}}$ image: centered at box of t , doubled size (gating)
 - late fusion: concat CNN output from $t, t + 1$ & fed into dense layer for fusion
 - LSTM tracker updates on fused info
 - dense layer regress the tracker (LSTM) output for bounding box at $t + 1$
- Training
 - bounding box defined by up-left, right-bottom coordinates in the crop (hence as a ratio of bounding box size)
 - $L1$ loss to encourage exact match
 - short sequence (2 unrolls) & multi-batch (64) to overcome plateaus
 - ⇒ slowly increase to 32 unrolls & batch size 4
 - use ground-truth crop when training with short sequence
 - ⇒ slowly increase probability to use predicted crop, when increasing unrolls
 - ⇒ to prevent from accumulating drifts
- Data Augmentation
 - utilize detection dataset: crop the object with bounding box
 - random crop a patch from the same image as background
 - occluders randomly taken from the same image
 - object with box initialized with velocity with Gaussian noise
- Tracking
 - an initial bounding box over arbitrary object given at the start
 - crops fed into net at each time step
 - LSTM state reset after each 32 frames (as maximally trained with 32 unrolls)
 - ⇒ hot-start by using the state in 1st forward pass (instead of $\mathbf{0}$)
 - ⇒ preserve the initial encoding & recover from drift
- Understanding
 - end-to-end training for both detector and tracker (improved GOTURN)
 - model regress changes to the box ratio ⇒ easier
 - observation usefulness modeled by LSTM tracker
 - LSTM tracker needs specialized training, and single-layer LSTM NOT enough (or, can hurt performance due to instability)
 - LSTM state reset: prevent drift, yet can fail if initial box overlaps other object
- C-RPN: Siamese Cascaded Region Proposal Network
 - Tracking Initialization
 - bounding box at first frame as ground truth / feature template
 - Detector

- fully conv siamese net for encoding (multi-scale info utilized)
 - ⇒ extracts info from target template x and search region z (gating)
 - ⇒ $\phi_n(\cdot)$ for features at n^{th} layer, backwards
 - feature fusion $\Phi_l(\cdot) = f(\Phi_{l-1}, \phi_l)$, with $\Phi_1 = \phi_1$
 - ⇒ recursively fuse semantic info with lower-level spatial info
 - $f(\Phi_{l-1}, \phi_l)$ = feature transform block
 - ⇒ upsample Φ_{l-1} by deconv, further 2 convs on ϕ_l for channel matching
 - ⇒ element-wise sum, then interpolation as downsampling
- Tracker
 - no explicit tracker (not modeling temporal info)
- Inference
 - siamese net extracts features from target template & search region
 - l^{st} RPN takes input $\Phi_l(x), \Phi_l(z)$, regresses based on anchor boxes A_l
 - discard any boxes $\in A_l$ with confidence/objectness lower than a threshold
 - ⇒ produce anchor boxes A_{l+1}
 - fuse semantic info with lower-level info, for both branch for x, z
 - t^{th} RPN takes
- Loss
 - loss for l^{th} RPN $L_l = \sum_{a \in A_l} L_{\text{cls}}(c_a^l, \hat{c}_a^l) + \lambda \sum_{a \in A_l} \hat{c}_a^l \cdot L_{\text{loc}}(r_a^l, \hat{r}_a^l)$,
 where c_a^l/\hat{c}_a^l the predict/label objectness for anchor a ;
 with r_a^l/\hat{r}_a^l the predict/label location for anchor a (encoded as YOLOv2)
 note: label based on current anchor $a \in A_l$ and ground truth \hat{r}_a
 - ⇒ total loss $L = \sum_l L_l$
- Training
 - random sample image from a video
 - ⇒ forming image pair (target template image, image with/without target)
- Tracking
 - arbitrary target object selected for tracking (treated as ground truth)
 - ⇒ extract the feature template
 - gating on frame t based on $t - 1$ result
 - C-RPN detect target within gated region of frame t
 - (the best remained, adjusted, anchor as the target to track)
- Understanding
 - multi-stage tracking: each RPN sequentially refine bbox (size & location)
 - hard negative mining by filtering out box proposal at each RPN stage
 - ⇒ training samples sequentially more balanced
 - ⇒ RPNs sequentially more discriminative
 - fusion of multi-level feature (spatial + semantic info) for RPN
- Fast Online Object Tracking and Segmentation
 - Detector
 -
 -

Multi-Object Tracking & Data Association

- Gated Association
 - Procedure
 - given prediction, its variance and detection noise, filter out an interested area
 - consider only detection inside the interested area (satisfying requirements)
 - score each detection & associate detection (detection result) with tracker \Rightarrow which detection belongs to which trajectory
 - Global Nearest Neighbor
 - choose the best / most probable / nearest
(under the constraint that an detection can associate with at most one track) \Rightarrow assume one detection is produced by single object
 - require accurate and sparse detection, with few false alarm \Rightarrow sensible to noise (easily fail in crowded scene)
 - Nearest Neighbor
 - choose the best / most probable / nearest (thought one detection may be used by multiple tracks)
 - Understanding
 -
- Joint Probability Data Association (JPDA)
 - Procedure
 - Understanding
 - probabilistic perspective for prediction-detection relation
 \Rightarrow cooperate with uncertain association: weight all detections by probability
 - hence, crowded detections tends to pull multiple tracks together
 \Rightarrow coalescence problem
- Multiple Hypothesis Tracking
 - Assumption
 - Gaussian models for target dynamics and noise in detection
 - uniform distribution for false alarm (false-positive detection)
 - Track Hypothesis
 - association result given track initiation, prediction and detection
 \Rightarrow a sequence of selected detection
 - compatibility: tracks are compatible, if they do NOT share any detection
 \Rightarrow any track update using the same detection are INcompatible
 - Track Tree (Clustering)
 - use incompatibility as edge, track as vertice, sort tracks by time
 \Rightarrow each connected tree becomes a cluster (track family)
(the tree level denotes time sequence)
 - each tree shares a common root node (the initial detection)
 - growing: whenever a new detection can be accounted for a track hypothesis (node)
 \Rightarrow the node generates 2 children nodes (tracks): update / not update
 - Global Hypothesis

- a global hypothesis contains only compatible track(s)
 \Rightarrow the collection of track, with ≤ 1 track from each tree/family
- Track Score
 - posterior ratio $r = \frac{p(D|T)p(T)}{p(D|F)p(F)} \triangleq \frac{p_T}{p_F}$,
 where $p(D|T), p(D|F)$ the likelihood given detection is true, false alarm
 with D the detections in current track
 \Rightarrow log ratio $lr = \ln \frac{p_T}{p_F}$
 - use log ratio as score, at time $t, L(t) = L(t-1) + \Delta L(t)$,
 where $\Delta L(t) = \begin{cases} \ln(1 - \hat{P}(D)) & \text{no update} \\ \Delta L_u(t) & \text{update} \end{cases}$,
 with $\hat{P}(D)$ the expected probability of detection;
 and $\Delta L_u(t)$ the residual error between prediction and detection
 ($\Delta L_u(t)$ may include covariance, density, \hat{P}_D , etc.)
- Global Hypothesis Score
 - $s_H = \sum_{k \in K} L_k(t)$,
 where K_H all (compatible) tracks in hypothesis H , L_k the score for track k
- Global Hypothesis Probability
 - computed from hypothesis score (a maximum weighted independent set problem)
- Track Probability
 - the sum of probability of all hypothesis that contains the track
- N -scan Pruning
 - given detection at time t , eliminate IMplausible tracks originated at $t - N$
 \Rightarrow suppress tree from exponentially growing;
 - N the time step buffer before decision (scan = time), usually $N \geq 5$
- Understanding
 - probabilistic perspective towards the result of decisions for data association
 (a larger scope than JPDA)
 - defer critical decision into the future
 \Rightarrow make decisions for the past after their observation available
 - model track alternatives, each with a probability, by track tree and hypothesis
 (for all possible tracks, model joint prob over all detections in a track)
 - model global joint probability of all tracks, by global hypothesis
 - similar to DP-longest substring: maintain a set of candidates
 - essentially, a bread-first search \Rightarrow real-time ability constrained by tree size
- Maximum Net Flow
 - Input
 - detection
 - Output
 - association result
 - Inference
 - detection as node, possible association as edge
 \Rightarrow construct a graph, with detection time as layer
 - \Rightarrow solve as maximum flow / minimal cost
 - Understanding

- another global optimization, vs. probabilistic perspective in MHT
- Inversed Reinforced Learning for Data Association with Markov Decision Process
 - Input
 - current detected bounding boxes, with objectness
 - previously predicted bounding boxes
 - Output
 - decision of data association between track & detection
 - Markov Decision Process (MDP) for Track Management
 - states: active, tracked, lost, inactive
⇒ model the state of a track
 - probability at each state given by trained model for each state
 - Inference
 - binary classification for track start (given a detection)
 - optic flow to track & association
(rule model to decide if lost)
 - regression model to associate lost track & current detection
(measure similarity)
 - rule for track death: lost for consecutive 6 frames
 - Training: Inversed Reinforcement Learning for Lost Recovery
 - classifier for track start: trained offline
 - regressor for lost track association: trained only when MDP make wrong decision
(similar to hard-example mining?)
 - Tracking ‘
 - one MDP for a track in multi-object tracking
⇒ multiple tracks may update with same detection (need to tune optic flow)
 - Understanding
 - explicit expression for track state
⇒ can design state for hard scenario
⇒ enable explicit control over optimization
 - ugly crashed model for each state ⇒ can be all unified to NN(s)
- Siamese CNN for Association
 - Input
 - image pair (I_1, I_2) , with optic flow $I_1 \rightarrow I_2$ as (O_1, O_2)
⇒ $D = [I_1, I_2, O_1, O_2]$ (resized & channel concat)
 - Output
 - probability of data association between two detections
 - Inference & Structure
 - 3 conv layers, max pooling, 4 dense layers
 - examine spatio-temporal info: position change & relative velocity by difference
 - NN feature vector concat with handcraft feature
 - feed into gradient boosting classifier (with 400 trees)
⇒ output as binary classification of (match, no match)
 - Training

- true positive: associated ground truth detection in 2 frames (time gap ≤ 15 frames)
- negative: wrong association to true detection of other track / false detection
- data augmentation: false alarm, distortion on image
- Tracking
 - given current data association probability for all track-detection pair \Rightarrow construct a linear program problem (with constraints) \Rightarrow a global optimization for association given probability
 - online tracker (e.g. kalman filter)
- Understanding
 - NN approach for association probability (with fusion of NN & rule-model via GB classifier) \Rightarrow fusion much better than pure NN \Rightarrow spatio-temporal info important
- Online Multi-Target Tracking Using RNN
 - Input
 - D the dimension of bbox encoding (e.g. x, y, w, h, objectiveness, etc.)
 - $x_t \in \mathbb{R}^{ND}$ all the N track state (bbox) at time t
 - $x_t^* \in \mathbb{R}^{ND}$ all the N predicted bboxes for time t , from time $t-1$
 - $z_t \in \mathbb{R}^{MD}$ all the $M-1$ detected bbox at time t , with an empty detection
 - $\varepsilon_t \in (0, 1)^N \in \mathbb{R}^N$ the existence probability (liveness) for all tracks
 - $A_t \in \mathbb{R}_{N \times M}$ the probability matrix for data association between track-detection
 - h_t hidden state of track RNN at time t
 - $C_t \in \mathbb{R}_{N \times M}$ the distance matrix between x_t^* and z_t
i.e. $C_t[i, j] = \text{dist}(x_t^*[i] - z_t[j])$
 - available ground truth: $\tilde{x}_t, \tilde{A}_t, \tilde{\varepsilon}_t$
 - Inference
 - given x_t, h_t , track RNN outputs state prediction x_{t+1}^* , compute h_{t+1}
 - based on pred x_{t+1}^* and detection z_{t+1} , compute C_{t+1}
 - given C_{t+1} and hidden state for i^{th} track $h_{t+1}[i]$
 \Rightarrow association LSTM scans over all detection z_{t+1}
 \Rightarrow regress $A_{t+1}[i, :]$, the association prob for i^{th} track and each bbox in z_{t+1} (as part of track RNN process)
 - given detection z_{t+1} , association prob A_{t+1} , latest liveness ε_t , with h_{t+1}
 \Rightarrow update state to be x_{t+1} , estimate liveness ε_{t+1}
 - Structure
 - track RNN consists of a 2-layer association LSTM
 - Loss
 - prediction $L_{\text{pred}} = \frac{\lambda}{ND} \sum (x_{t+1}^* - \tilde{x}_{t+1})^2$
 - updated state $L_{\text{update}} = \frac{\kappa}{ND} \sum (x_t - \tilde{x}_{t+1})^2$
 - liveness $L_{\varepsilon} = \tilde{\varepsilon}_t \log \varepsilon_t + (1 - \tilde{\varepsilon}_t) \log(1 - \varepsilon_t) + |\varepsilon_t - \varepsilon_{t-1}|$
 \Rightarrow minimize the diff between consecutive liveness estimation \Rightarrow smoothness (prevent track from termination for only a single detection lost)
 - association $L_a = -\log(A_{t+1}[i, \tilde{j}])$, where \tilde{j} the true association for i^{th} track
 - Training
 - data augmentation: sample synthetic trajectories from each labeled video (Gaussian distribution)

- Tracking
 - forward unroll track RNN, if liveness ≤ 0.6 , corresponding track ignored
 - liveness ≥ 0.6 again, a new track initiated
- Understanding
 - specialized RNN cell accounting for prediction, update, birth-death of all tracks
 - another LSTM cell designed for data association
 \Rightarrow able to learn 1 – 1 association by scanning
 (yet, unnecessary, since N, M fixed i.e. a fixed size mapping)
 - utilize given detector \Rightarrow no appearance model (but only location & size)
 - able to maintain at most N track with maximally M detection at a time
- Collaborative Deep Reinforcement Learning for MOT
 - Notation
 - I_t the t^{th} image frame
 - $b_{i,t}^*$ a bbox for i^{th} ground truth object p_i at frame t
 - $B_{i,t}^*$ a set of bboxes sampled around $b_{i,t}^*$
 - $g(a, b)$ cal the IoU between bbox a, b
 - $p_{i,t} = \{b, f\}$, the i^{th} detected object at t ,
 where $b = \{x, y, w, h\}$ the bbox; f the appearance model
 \Rightarrow distance $d(p_1, p_2) = \alpha(1 - g(b_1, b_2)) + (1 - \underbrace{\frac{f_1^T f_2}{\|f_1\| \|f_2\|}}_{(\cos \text{ dist})})$
 - $H = \{b_1, \dots, b_t\}$ the history trajectory of an object
 $\Rightarrow H^K = \{b_{t-K+1}, \dots, b_t\}$ the history of past K frames
 - an agent for each object $g = \{H, p\}$
 - detections as environment $\hat{P}_t = \{\hat{p}_1, \dots, \hat{p}_{n_t}\}$
 - state at frame t , $s_t = \{G_t, \hat{P}_t\}$, where $A_t = \{g_1, \dots, g_m\}$
 - set of actions $\mathcal{A} = \{\text{update, ignore, block, delete}\}$
 - Prediction Net Inference
 - crop the frame $t + 1$ at the location of estimated bbox $b_{i,t}$ (of frame t)
 - 3 conv layers, then dense, then concat with $H^{K=10}$ (fuse with temporal info)
 - 2 dense layers to regress $b_{i,t+1}$, the bbox for object i of frame $t + 1$
 - Prediction Net Training
 - regression loss $L = \sum_{i,t} \sum_{b \in B_{i,t}} g(b_{i,t+1}^*, \phi(I_t, b, H_i^{K=10}))$,
 where I_t the image frame at time t , ϕ the mapping of pred net
 - Action
 - update: $f_{t+1} = (1 - \rho_f) \cdot f_t + \rho_f \cdot \hat{f}_{t+1}$, where $\hat{f}_{t+1} \in$ selected detection \hat{p}_{t+1} ;
 $b_{t+1} = (1 - \rho_b) \cdot b_t + \rho_b \cdot b'_{t+1}$, where b'_{t+1} predicted position;
 (ρ_f, ρ_b pre-selected)
 - ignore: no detection suitable, use only prediction for update ($\rho_f = 0, \rho_b = 1$)
 - block: same as ignore, no detection due to occlusion
 - delete: remove the agent
 - Reward
 - for agent g at time t , with pred and ground-truth box b'_{t+1}, b_{t+1}^*

- reward for agent g at time t : $r_t^* = r_t + \beta r_{j,t}$, where
 r_t for itself; $r_{j,t}$ for its nearest neighbor, β a balance factor
 $(r_t, r_{j,t}$ calculated in the same manner)
 \Rightarrow agents need to collaborate for better reward
- for action $a \in \{\text{update, ignore, block}\}$
 $\Rightarrow r_t = \begin{cases} 1 & \text{if } IoU \geq 0.7 \\ 0 & \text{if } 0.5 \leq IoU \leq 0.7 \text{ (IoU calculated between } b'_{t+1} \text{ and } b^*_{t+1}) \\ -1 & \text{otherwise} \end{cases}$
for action $a = \text{delete} \Rightarrow r_t = 1$ if object disappear; else -1
- $\Rightarrow Q(s_t, a_t) = r_t^* + \gamma r_{t+1}^* + \gamma^2 r_{t+2}^* + \dots$, where γ decaying param
- Decision Net Inference
 - for each $g \in G_t$ with its current & predicted location b_t, b'_{t+1}
 - select a neighbor agent $g_j \in G_t - \{g\}$, that is nearest to b_t
 - select the detection $\hat{p} \in \hat{P}_{t+1}$ that is nearest to b'_{t+1}
 - 3 feature maps ($p \in g, p_j \in g_j, \hat{p}$), flatten as 1-D vector input
 $\Rightarrow 3 \times$ dense layer to output prob over actions $\pi(a|s, \theta)$, where θ the weights
- Decision Net Training
 - goal $\arg \max_{\theta} L(\theta) = \mathbb{E}_{s,a} \log(\pi(a|s, \theta)) \cdot Q(s, a)$
 $\Rightarrow \frac{\partial}{\partial \theta} L = \mathbb{E}_{s,a} \frac{\partial}{\partial \theta} [\log(\pi(a|s, \theta)) \cdot Q(s, a)]$
 $= \mathbb{E}_{s,a} \left[\frac{Q(s, a)}{\pi(a|s, \theta)} \cdot \frac{\partial}{\partial \theta} \pi(a|s, \theta) \right]$
 \Rightarrow increase probability for actions with $Q > 0$; decrease for those with $Q < 0$
 - to speed up converge: value for state $s, V(s) = \frac{\sum_a p(a|s) Q(s, a)}{\sum_a p(a|s)}$
 \Rightarrow advantage value $A(s, a) = Q(s, a) - V(s)$
(in case all $Q > 0$, or all $Q < 0$ at the beginning: use expectation as zero-line)
 - policy gradient as $L(\theta) = \mathbb{E}_{s,a} \log(\pi(a|s, \theta)) A(s, a)$
 - pre-training: set $\gamma, \beta = 0$ before searching hyper-parameter β, γ
 - training set: detection bbox = $\{\hat{b} \in \text{detection } \hat{P} \mid \text{IoU}_{\text{label}}^b > 0.5\} + \{\text{label } b^*\}$
- Tracking
 - initiate a track for each initial detection
 - for each agent, predict its location at $t+1$, by pred net
 - for each agent g , select its closest detection $\hat{p} \in \hat{P}_{t+1}$, $p_j \in$ closest agent g_j
 \Rightarrow decision net: $(p \in g, \hat{p}, p_j) \rightarrow \text{action } \mathcal{A}$
 - track terminates when decision net decides to "delete"
- Understanding
 - prediction net as a detector to give more precise location
(an implicit self-trained detector)
 - decision net to eliminate false alarm & combine prediction
 \Rightarrow robust to different detector/predictor
 - decision net trained to collaboratively maximize utility
 \Rightarrow mitigate the false negative
 - may trapped in false appearance feature \Rightarrow ID switch
(e.g. blue box handed from one person to another person)

Instance Tracking

- Fast Online Object Tracking and Segmentation : A Unifying Approach
 - Input
 - initial bbox z as target template at first frame (exemplar)
 - cropped search region x_t , centered at target location of $t - 1$
 - Output
 - response map for object localization
 - bbox regression for target (with resizing & rotation)
 - binary segmentation mask i.e. pixel \in (target, not target)
 - Inference
 - given z, x_t backbone CNN extract feature as $f(z), f(x_t)$
 - cross-correlation response map $g(z, x_t) = f(z) \star f(x_t)$,
where $f(z)$ used as kernel
 $\Rightarrow g$ as response of candidate window (RoW)
 \Rightarrow encode similarity between z and each candidate window/bbox in $f(x_t)$
 - Structure
 - backbone CNN ResNet-50 as feature extractor, with adjustment (dilated conv for better resolution)
 - cross-relation between two extracted feature maps
 \Rightarrow response map, each location an RoW
 - bbox regression: $2 \times (1\text{-by-}1 \text{ conv})$ for k anchors at each RoW
 - bbox classification: $2 \times (1\text{-by-}1 \text{ conv})$ for score of k anchors at each RoW
 - segmentation: $2 \times (1\text{-by-}1 \text{ conv}) +$ upsampling with skip + per-pixel sigmoid
 \Rightarrow upsample into a mask for each RoW location
 - Training
 - loss $L = \lambda_1 L_{\text{mask}} + \lambda_2 L_{\text{score}} + \lambda_3 L_{\text{box}}$
($L_{\text{mask}}, L_{\text{box}}$ considered only at location for ground truth)
 - trained different branch using corresponding dataset
 - data augmentation: random jitter, translation, rescaling
 - Tracking
 - one-step update, may use mask to produce
 - multi-object: multiple initialization, each with a net as tracker
 - Understanding
 - multi-task training improves all branches
(semi-supervised video object segmentation, bbox tracking)
 - mask branch: output the mask in the context of x for that obj in z

12.5.3 Face Recognition**Problem Formulation**

- Input Data
 - Image
 - image from camera
- Goal

- Identity Recognition
 - recognize the identity of the face in image
 - refuse to recognize if the face does NOT belongs to any stored identity
- Liveness Detection
 - make sure the face in image are from a live human (instead of picture etc.)
- Face Verification
 - Input
 - image from camera & identity
 - Goal
 - true/false, regarding whether the image content belongs to the identity
- Challenge
 - One-shot Learning
 - given only single (at most, few) face-identity pair for each identity
 - still, need to build a robust system for recognition task

Classic Approach


- Siamese Network as Encoder
 - Structure
 - CNN + dense layer to encode the input image x^i as a vector $f(x^i)$
 - Learning Goal
 - minimize $\|f(x^i) - f(x^j)\|^2$ if x^i, x^j from same identity
 - maximize $\|f(x^i) - f(x^j)\|^2$ if x^i, x^j from different identity
 - \Rightarrow learning encoding given a fixed distance function $d(x_1, x_2) \geq 0$ (here, $d(x_1, x_2) = \|x_1 - x_2\|^2$)
 - Triplet Loss
 - given an anchor image A representing the identity I
 - take a positive image $P \in \text{identity } I$; an negative image $N \notin \text{identity } I$
 - $\Rightarrow L(A, P, N) = \max(d(f(A), f(P)) - d(f(A), f(N)) + \alpha, 0)$, where α an hyperparamter to make sure the net differentiate them by a margin; $\max()$ to make the loss = 0 as long as the requirement satisfied
 - Training: Hard Negative Mining
 - due to large variance in the dataset $\Rightarrow d(A, P) \ll d(A, N)$ in most case
 - due to large number of identities \Rightarrow permutation explosion
 - \Rightarrow evaluate current net on dataset, use mistakes for the next epoch
 - Understanding
 - learn a encoder towards a selected distance function \Rightarrow use permutation to have more training examples
 - able to precomputing the encoded vector for fast recognition
- Encoding + Binary Classification
 - Structure
 - still, CNN + dense layer to encode input image

- Learning Goal
 - given two encoded vectors, another net (or logistic regression) to perform binary classification
 - ⇒ 1 for two image has same identity; 0 for different identities
- Understanding
 - still, utilize permutation for larger training set (use pair, instead of triplet)
 - learn the similarity function as well: output directly the result of comparison
 - pre-compute the encoding of Siamese net
 - ⇒ enable flexible deployment (device performs only bi-classification)


12.5.4 Recognition at a Distance

Problem Formulation

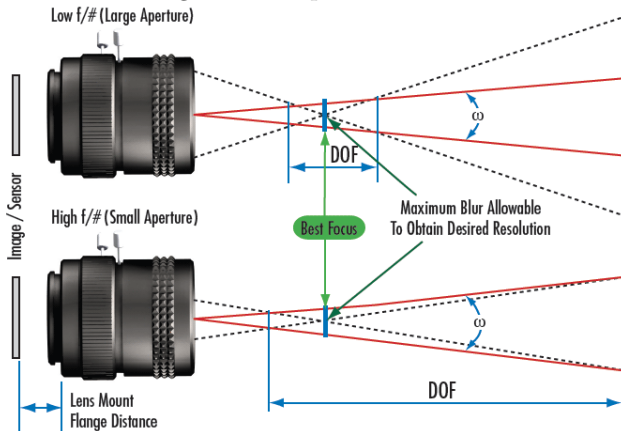
- Basic Image-forming Theory
 - Camera Coordinate
 - image axes: same as axes for matrix index
 - x, y align with the image axes, z align with focal line
 - Resolution
 - the minimum feature size of the object under inspection
 - ⇒ finally, represented as the pixel in the image
 - (hence, jointly controlled by pixel size, lens, object distance and sensor size)
 - pixel count: total number of pixels in the row/col of image
 - pixel density: dots/pixels per inch (dpi)
 - sensor resolution: pixel count on sensor and their physical size and spreading
 - lens resolution: the impact of optics (diffraction, etc) on light info towards sensor
 - ISO: Light Sensitivity
 - used to be the light sensitivity of the film
 - ⇒ in digital camera: the amplification/attenuation for raw sensor value
 - ⇒ can add noise if amplified too much, due to dark current
 - base ISO:
 - Exposure Time
 - control incoming light (denoted by shutter speed)
 - Focal Length F
 - the distance from last lens central point to focal point, determined by lens physics
 - given a fixed object distance, larger F , larger formed object (object distance need to be larger than focal length)



- Aperture
 - diameter of entrance pupil D : control the size of aperture
 \Rightarrow aperture stop: aperture setting to restrict input pupil size and hence brightness
 - larger size (D), more incoming light
- F-Number $N = \frac{F}{D}$
 - reveal the signal to noise ratio, contrast, image brightness
 - larger N , smaller D , less incoming light, larger DOF, larger depth of focus
- Depth of Field (DOF) & Depth of Focus
 - definition: the range on z axis where resolution can be maintained
 - before the lens: depth of field (in the field)
 - after the lens: depth of focus (inside camera)




- Depth of Field
 - DOF as the range where 1 pixel can hold all info from 1 position



where, dashed line the incoming light;
 red line the system resolution (the pixel size, changing with object dist)
 \Rightarrow DOF defined by intersection of two type of lines
 \Rightarrow with large enough DOF, resolution can be bad at the end of DOF (pixel size too large)

- resolution with DOF




where black boxes a series of pixels, beams in colors as light info from an object
 \Rightarrow the placement of box on beams shows the information contained in a pixel

- ⇒ larger DOF, large depth range for acceptable resolution
(as pixel contain more separable info)
- ⇒ out of DOF: pixel can not distinguish beams from diff location

- Depth of Focus

- depth of focus as the range where 1 pixel can hold all info from 1 position
- resulted from tilted sensor ⇒ may not in the ideal position to receive light info



- where right-most line as ideal focal plane, left lines as sensor tilted by $12.5/25\mu m$
 ⇒ tilt not influence pixel in the center, but the pixel in the corner
 ⇒ out of depth of focus: pixel can not distinguish beams from diff location

- Input Data

- Video

- hot standby camera with stationary & active vision
- stationary vision: wide field of view with low resolution for detection
- active vision: narrow field of view with high resolution for detail analysis

- Goal

- Recognition in the Wild

- large coverage areas: $> 100m$ range
⇒ scale beyond the theoretical analysis and practical design advice
- with NO subject cooperation

- Challenge

- None Cooperative Subject

- not cooperating, may even be evading the system

- Resolution

- low resolution due to the very long object distance
- restricted by lens resolution, which is then restricted by price
- trade-off between wideness (coverage) & depth (zooming)

- Illumination

- dynamic illumination in the wild scene
- maximum light intensity restricted by aperture size
(given a fixed exposure time)

- Distortion and Blur

- amplified noise, due to:
low brightness ⇒ low signal-to-noise ratio ⇒ ISO amplification
- motion blur (if trade-off between ISO & exposure time)
- blur from sensor tilt, as a hot standby system
- fog, haze & atmosphere blur for very long distance recognition

- Pose
 - the view angle due the camera position
(usually overhead for less occlusion \Rightarrow downward tilt)
- Multi-Object
 - schedule the limited high-resolution vision resource for multiple candidates
 \Rightarrow time window prediction, scheduling & resource allocation
- Physical Coupling
 - expensive field test (mitigated by virtual environment)
- Application
 - Watch-list Recognition
 - an alert when person of interest appear/approach
 - Re-recognition
 - cross-camera tracking, long-range persistent tracking
 - Logging
 - catalog best recognized feature (e.g. face) for each person entering a region
 - marketing: understand customer activities and behavior

Stationary Vision

•


3D Imaging

• ◦

12.5.5 Image Style Transfer


Problem Formulation

- Input Data
 - Content Image C
 - the image containing the spatial info (content)
 - Style Image S
 - the image containing the style of presenting the content
- Goal
 - Generated Image G
 - a image with content from C drawn in style of S



Classic Approach

- Neural Style Transfer
 - Learning Goal
 - given input C, S with output G , loss $L = \sum_l [\alpha L_{\text{content}}(C, G) + \beta L_{\text{style}}(C, S)]$, where l is sum over chosen hidden layers of the CNN
 - \Rightarrow minimize content & style difference
 - Content
 - given input, the activations from a set of (hidden) layers of the net
 - \Rightarrow similarity of C, G measured as $\sum_l d(a^{l(C)}, a^{l(G)})$, where $a^{l(\cdot)}$ the feature maps at layer l given the input, $d(\cdot)$ a distance function (e.g. $d(x_1, x_2) = \|x_1 - x_2\|^2$)
 - Style
 - given input, the correlation between activations across channels, for chosen layers \Rightarrow correlation matrix across feature map at each channel as style matrix (actually, gram matrix)
 - let $a_{i,j,k}^l$ the activation at a $h \times w \times c$ conv kernel location i, j, k in layer l \Rightarrow style (gram) matrix $M_{k,k'}^l = \sum_{i,j} a_{ijk}^l \cdot a_{ijk'}^l$, for all $k, k' \in \{1, \dots, c\}$
 - \Rightarrow similarity of S, G measured as $\sum_l \left[\frac{1}{(2h^l w^l c^l)^2} d(M^{l(S)}, M^{l(G)}) \right]$ where $M^{l(\cdot)}$ the gram matrix at layer l given the input, $d(\cdot)$ a distance function, with normalization term $\frac{1}{(2h^l w^l c^l)^2}$ (e.g. $d(x_1, x_2) = \|x_1 - x_2\|_F^2$, the euclidean norm between matrices)



12.5.6 Point Cloud Data Processing

Problem Formulation

- Input Data
 - Point Cloud
 - from lidar \Rightarrow 3-D position x, y, z & reflection intensity
 - from radar \Rightarrow 2-D bird-view x, y & intensity (rcs) & speed (along radial direction)
- Goal
 - 3-D Environment Modeling
 - bounding box
 - segmentation (per-point classification)
 - instance segmentation, etc.
- Challenge
 - Sparsity
 - point gets much sparser in distance (e.g. $> 40\text{m}$)
 - Varying Density
 - due to occlusion, distance etc.

Common Preprocessing

- Voxelization
 - Goal
 - create 3-D pixel, the voxel
 - Procedure
 - apply 3-D grid on the space
⇒ each point resides in a spatial cell, the voxel
- Bird-view Projection
 - Goal
 - project onto 2-D map of a bird-view perspective
⇒ better resolution due to reduced dimension
 - Procedure
 - apply 2-D grid on the ground
⇒ each point resides in a cell, or, each cell consists of several point
 - extract height information from points in each cell
⇒ each pixel with a height h
- Cylindrical Projection (Frontal View)
 - Goal
 - project onto 2-D map of pilot perspective
⇒ better align with camera perspective
 - Procedure
 - given 3-D cartesian coord x, y, z
⇒ spherical coord $r = \sqrt{x^2 + y^2 + z^2}, \theta = \arctan \frac{y}{x}, \phi = \arcsin \frac{z}{r}$
 - slicing on horizontal & vertical angle: each point resides in a 3-D slice
⇒ normalized by angle resolution $\theta' = \lfloor \frac{\theta}{\delta\theta} \rfloor, \phi' = \lfloor \frac{\phi}{\delta\phi} \rfloor$
 - each pixel (θ', ϕ') extract depth $d = \sqrt{x^2 + y^2}$ & height z from its point(s)

Point Cloud

- PointNet
 -
- VoxelNet
 - Preprocessing
 - voxelization
 - random sampling points in each voxel to be at most T points a voxel
 - Bounding Box Encoding
 - x, y, z for center position; l, w, h for size; θ for orientation, the yaw rate
⇒ $x_g, y_g, z_g, l_g, w_g, h_g, \theta_g$ the ground truth box
⇒ $x_a, y_a, z_a, l_a, w_a, h_a, \theta_a$ an anchor box
 - normalized residual position: $\Delta x = \frac{x_g - x_a}{\sqrt{l_a^2 + w_a^2}}, \Delta y = \frac{y_g - y_a}{\sqrt{l_a^2 + w_a^2}}, \Delta z = \frac{z_g - z_a}{h_a}$
 - normalized size ratio: $\Delta l = \log(\frac{l_g}{l_a}), \Delta w = \log(\frac{w_g}{w_a}), \Delta h = \log(\frac{h_g}{h_a})$
 - residual orientation: $\Delta\theta = \theta_g - \theta_a$

- Inference
 - voxel map with raw points as input, output 3-D classification & regression map
- Voxel Feature Encoding Layer
 - voxel $V = \{p_i = [x_i, y_i, z_i, r_i] \in \mathbb{R}^4\}_{i=1, \dots, t}, t < T$
 - augment each point with its offset to the centroid (v_x, v_y, v_z) , the mean of $p \in V$
 $\Rightarrow p'_i = [x_i, y_i, z_i, r_i, x_i - v_x, y_i - v_y, z_i - v_z]$
 - 1×1 conv (with batch norm, ReLu) on each feature point: $p'_i \rightarrow f_i$
 - element-wise max pooling on $f_i \in V$: voxel-wise feature \hat{f}
 - augment each processed feature point with voxel feature: $f_i^{\text{out}} = [f_i, \hat{f}]$
- Structure
 - stacked voxel feature encoding layer: deep net to extract feature for each voxel
 - detection model: RPN modified for 3-D detection
- Training
 - per-voxel weighting: balanced according to num of positive example
 - smooth L1 for regression loss
 - further up-/down-sample positive/negative voxel in classification loss
 - predicted box considered positive, if its IoU with label box > 0.6 & is the highest
 (negative, if its IoU with any label box < 0.45)
 (not-care, if its IoU with any label box $\in [0.45, 0.6]$)
 - data augmentation: box (collision-free) perturbation in location, size, rotation
- Understanding
 - range: $[-3, 1] \times [-20, 20] \times [0, 48]$ in height-width-length
 - voxel feature extraction: voxel-level PointNet
 \Rightarrow able to concat with various detection models (e.g. YOLO, etc.)
- LMNet
 - Preprocessing
 - cylindrical projection
 - Bounding Box Encoding
 - given a pixel (corresponding to a lidar point) p in a bounding box
 \Rightarrow encode box under coord originated at p
 \Rightarrow axes: x along radial direction; y parallel with horizontal plane; z accordingly
 - encode 8 corner \Rightarrow 24 channel encoding for each box
 - Classification Encoding
 - per-pixel classification (car, pedestrian, cyclist, none)
 - Inference
 - per-pixel regression & classification \Rightarrow standard CV detection
 - non-max suppression as postprocessing
 (score modified to be the num of nearby-box)
 - Structure
 - convs - max pooling - dilated convs - unpooling - branch for regression/classification
 - Training
 - point-wise weighting
 \Rightarrow regression: consider box size for each class
 \Rightarrow classification: downsample background pixel (standard)

- Understanding
 - real-time timing due to simple structure with dilated conv
(though performance hurt...) \Rightarrow enable further fusion with image, pertaining real-time timing

Point Cloud + Image Fusion

12.6 Natural Language Processing

12.6.1 Language Representation

Problem Formulation

- Input
 - Language Token/Corpus
 - words, sentences, paragraphs, ... \Rightarrow can be language at various level
- Goal
 - Distributed Vector Representations as Embedding Matrix $E_{M \times N} = [e_1, \dots, e_N]$
 - e the column vectors, M the desired embedding length, N the total tokens num
 \Rightarrow look up for the desired embedding
(NOT using matrix multiplication due to sparsity from one-hot encoding)
 - distributed representation: decomposed yet meaningful
 \Rightarrow fight the curse of dimensionality
 - \Rightarrow Meaningful Vector
 - able to measure the (dis-)similarity of between tokens (words)
 \Rightarrow semantic meaning: "Germany"- "Berlin" & "France"- "Paris"
 \Rightarrow syntactic meaning: "quick"- "quickly" & "slow"- "slowly"
(e.g. $e_{\text{man}} - e_{\text{woman}} \approx e_{\text{king}} - e_{\text{queen}}$, where e_{text} the embedding for word "text")
 - \Rightarrow allow NLP model to be more robust & generalize better
- Challenge
 - Problem of Bias
 - word embedding reflect biases of text used to train the model
e.g. "father-doctor" as "mother-nurse" \Rightarrow gender bias
 - \Rightarrow can cause discrimination when making decision

Overview

- Character Embedding
 - One-hot Encoding
 - a one-hot vector with length 26
- Word Embedding
 - Word Dictionary
 - a collection of high-frequency word, embedded as one-hot vector
 - special token <UNK> for unknown word
 - Features from Rule Model

- number at each vector location denotes the score for the word matching a rule (e.g. location for "is.food" contains score $s \rightarrow 1$ for "apple", $s \rightarrow 0$ for "man")
- Part-of-Speech (POS) Tag
 -
- Word2Vec Embedding
 - construct supervised learning from UNlabeled corpus
- Global Vector for Word Embedding (GloVe)
 - linear model with simple optimization goal
- RNN Encoder
 - apply RNN model as encoder on characters in the word
 \Rightarrow no more `UNK` or unknown word
- Sentence/paragraph Embedding
 - RNN Encoder
 - apply RNN model as encoder on words in the sentence
 (last hidden layer as encoding)

Word2Vec Embedding

- N-gram Model
 - Learning Objective Setup
 - simple network to predict the $N + 1^{\text{th}}$ word given previous N words as input (e.g. using single softmax layer)
 - Practice
 - E randomly initialized & all words in corpus encoded in one-hot vector
 - forward prop: word in one-hot \rightarrow lookup $E \rightarrow$ linear layer \rightarrow softmax to predict
 - training: update linear layer parameters & matrix E as weights (with cross-entropy loss)
 - Understanding
 - learn $P(t|c)$, where t the target word, c the previous N context words
 - setup a even larger training set from a large corpus
 - Generalization
 - more context: take input from both previous and after words
 - less & close context: take only the last word as input \Rightarrow 1-gram model
- Skip-gram Model
 - Learning Objective Setup
 - choose only 1 single word as context word
 \Rightarrow balance sampling w.r.t. word frequency (e.g. prevent tons of "the", "a", ...)
 - randomly choose other word(s) in the sentence as target word(s)
 - \Rightarrow to predict target word(s) given only context word as input
 \Rightarrow learn word vector representations that are good at predicting the nearby words
 - Practice
 - same as N-gram model
 - Understanding

- harder supervised learning task, yet goal is to learn E
- better reflect the statistic: similar word appear in similar context
(e.g. "soviet"- "union" appears much more often than "soviet"- "sasquatch")
⇒ embedding for similar target word adjusted with similar gradients
⇒ lie closer in vector space
- cons: softmax over large word dict ⇒ low computation
⇒ mitigated by hierarchical softmax, noise contrastive estimation (NCE)
- Skip-gram with Negative Sampling
 - Learning Objective Setup
 - choose a pair of context and target word (c, t) as positive example
 - generate k negative examples by: same context word c & random word t' as target
 - given a pair of words, binary classification: is a (context, target) pair?
⇒ distinguish valid target word from k draws from noise distribution
 - Practice
 - detect meaningful (c, t) pair/phrase by heuristic method
e.g. if c, t co-appear within 10-words distance more than a threshold, etc.
 - $k = 5 - 20$ for small train set; $k = 2 - 5$ for large train set
(larger noise to avoid overfitting)
 - sample random word t' from modified uniform distribution $\frac{1}{Z}U(t)^{3/4}$ over words
 - subsample frequent words: sampled t' discarded by probability $P(t') = 1 - \sqrt{\frac{thr}{f(t')}}$
where thr a threshold, $f(t')$ the frequency of t' in corpus
(to avoid meaningless words like "the", "a", etc.)
 - E randomly initialized & all words in corpus encoded in one-hot vector
(forward prop similarly)
 - Understanding
 - learn $P(y = 1|c, t)$ via logistic regression
⇒ much computationally affordable compared to giant softmax (less weights)
⇒ much simpler approach than hierarchical softmax & NCE
 - non-linear model (logistic reg) also prefers linear structure of word embedding
⇒ cosine distance still measures (dis-)similarity

Global Vector for Word Embedding (GloVe)

- Overview
 - Context-Target Matrix X
 - x_{ij} : the count of times word w_i appear in the context of word w_j
(context definition can be non-symmetric)
 - Learning Objective Setup
 - given embedding matrix E , minimize $\sum_{i,j} f(x_{ij})(\theta_i^T e_j - \log x_{ij})^2$,
where e_j the embedding for w_j , θ_i the weights associated with w_i
 - $f(x_{ij})$ a weighting term to balance infrequent-frequent words
($f(x_{ij})$ for $x_{i,j} = 0$, preventing $-\log 0$)
 - Practice
 - gradient decent directly optimize the simple objective
 - final embedding for word w , $w_e = \frac{1}{2}(e_w + \theta_w)$
⇒ as θ_w, e_w in objective interchangeable

- Understanding
 - directly model the linear structure in word representation (project input e directly to output $\theta^T e$)
 - final linear structure probably NOT align with human interpretable axis \Rightarrow yet probably a combination of them (from a higher view)

Addressing Bias in Word Embedding

- Overview
 - Input Data
 - a trained word embedding
 - Goal
 - identify the bias in embedding
 - eliminate the bias if it appears in undesired places
 - Identify Bias Direction
 - singular value decomposition to identify the axes where biases lie (similar to a PCA)
 - e.g. principle component of $e_{\text{man}} - e_{\text{woman}}, e_{\text{male}} - e_{\text{female}}, \dots$
 - Neutralize
 - for all NOT definitional word (where bias should NOT appear) \Rightarrow project to axes orthogonal to bias axes (to get rid of bias)
 - e.g. project e_{doctor} to the axes to reduce component in bias axes
 - Equalize Pairs
 - for all definitional word (where bias should appear) \Rightarrow adjust their distance towards non-definitional word to be the same (may train/handpick all definitional words, which is only a small set)
 - e.g. make sure $d(e_{\text{boy}}, e_{\text{doctor}}) = d(e_{\text{girl}}, e_{\text{doctor}})$

12.6.2 Language Modeling

Problem Formulation

- Input Data
 - Sequence
 - a word with sequence of characters
 - a sentence with sequence of words/characters
 - a paragraph with sequence of sentences/words/characters
- Goal
 - Probability Distribution
 - model the appearance probability of the sequence $P(z^1, \dots, z^t)$, where z^t the token at time t

Classic Approach

- RNN
 - Practice
 - at each time, output token distribution conditional on previous token(s)
 $\Rightarrow y^t = p(z^t | z^1, \dots, z^{t-1})$, where z^t the token at time t
 - \Rightarrow sequence probability $P(z^1, \dots, z^T) = P(z^1)P(z^2 | z^1) \dots P(z^T | z^1, \dots, z^{T-1}) = \prod_t^T y^t$
 - Inference
 - $\mathbf{0}$ vector as both (initial) hidden state & input at time 0
 \Rightarrow estimate $y^1 = p(z^1)$, the distribution for being the 1st token
 - for time $t = 2, \dots, T$, take input x^2, \dots, x^T with hidden state h^1, \dots, h^{T-1}
 \Rightarrow estimate each conditional distribution (conditioning by passing hidden state)
 - Training
 - for time 1, input $x^1 = \mathbf{0}$, previous hidden state $h^0 = \mathbf{0}$
 - for time $t = 2, \dots, T$, input $x^t = z^{*t-1}$ the true token of $t-1$ in the given sequence
 - Generative Model: Sampling New Sequence
 - sampling the first token \hat{z}^1 according to the distribution y^1
 - for $t = 2, \dots$, take input $x^t = \hat{z}^{t-1}$, the token sampled from y^{t-1}
 - until $t > T$ or the end signal sampled (e.g. the period "." in a sentence)

Masked Language Model**12.6.3 Name-Entity Recognition**

•

12.6.4 Sentiment Classification**Problem Formulation**

- Input Data
 - a sentence / paragraph
- Goal
 - predict the degree of positive/negative attitude
- Challenge
 - small training set

Classic Approach

- Many-to-one RNN
 - Practice
 - each word encoded by word embedding
 - RNN scanning through paragraphs
 - last hidden layer as paragraph representation & used to classify/regress

12.6.5 Neural Machine Translation

Problem Formulation

- Input Data
 - Sequence
 - typically sentence, can be also multiple sentences (paragraph)
- Goal
 - Sequence
 - generated sentences in desired language

Classic Approach

- RNN Encoder-Decoder
 - Overview
 - an RNN encodes the input sequence by its last hidden layer
 - input encoding used as initial hidden state for decoder RNN
 - decoder RNN generates (conditional) distribution over words at each step
 \Rightarrow for time t , $y^t = p(z^t | x^1, \dots, x^{T_x}, \hat{z}^1, \dots, \hat{z}^{t-1})$,
 where z^t the token at time t , $\hat{z}^1, \dots, \hat{z}^{t-1}$ the tokens chosen from y^1, \dots, y^{t-1}
 (z^t a random variable, \hat{z}^t a concrete assignment, y^t a conditional distribution)
 - unroll until stop sign generated
 - Understanding: Conditional Language Model
 - decoder functions like language modeling, only different in its initial hidden state
 - \Rightarrow measure the conditional distribution $p(z^1, \dots, z^{T_y} | x^1, \dots, x^{T_x}) = \prod_{t=1}^{T_y} y_t$,
 where z^1, \dots, z^{T_y} the generated sequence, x^1, \dots, x^{T_x} the input sequence
 - Improvement
 - combined with attention model

Choosing Output Sequence

- Greedy Search
 - Practice
 - choose the words of highest (conditional) probability at each time step
- Beam Search
 - Practice
 - with a vocabulary size of N , a beam with size b , input sequence $\mathbf{x} = x^1, \dots, x^{T_x}$
 - to start, choose top- b tokens (among N tokens) at the 1st step
 - for step t , input each previous stored b tokens to have b conditional distributions
 - choose the top- b token pairs (among $b \times N$ pairs) regarding joint probability
 $\Rightarrow P(z^1, \dots, z^t | \mathbf{x}) = P(z^t | z^1, \dots, z^{t-1}, \mathbf{x}) P(z^1, \dots, z^{t-1} | \mathbf{x})$
 - Normalization by Length
 - reason: short sequence with less $y^t \in [0, 1] \Rightarrow$ larger in general
 - choose t^{th} pair regarding the normalized probability $\frac{1}{t} P(z^1, \dots, z^t | \mathbf{x})$

- \Rightarrow more numerically stable
- Understanding
 - approximately search the sequence with highest joint (conditional) probability
 \Rightarrow try to maximize $P(z^1, \dots, z^{T_y} | x^1, \dots, x^{T_x})$
 - similar to viterbi algorithm in HMM $\Rightarrow b = 1$ reduce to greedy search
 - B usually chosen in 10 in research, > 1000 in commercial system
 (still faster than BFS/DFS, yet no guarantee on finding best result)

•

12.6.6 Speech Recognition

Problem Formulation

- Input Data
 - Sequence
 - an audio sample, with each frame as a time step
- Goal
 - Sequence
 - text (words/sentences) corresponding to the audio
- Challenge
 - Variable Timing
 - output (letter/words) usually has much less time steps than input (audio frames)
 \Rightarrow multiple input time steps corresponding to same output time step

Learning Objective

- Connectionist Temporal Classification (CTC) Loss
 - Goal
 - avoid learning boundaries and timings
 - Practice
 - two sequences considered equivalent if they differ only in alignment, ignoring blanks
 \Rightarrow remove duplicate token (e.g. letters) from both sequence before comparison

Classic Approach

- RNN Encoder-Decoder
 - Overview
 - encoder scan through audio frames & decoder output letter/punctuation/"blank"/"space"
 - "blank": no symbol v.s. "space": delimiter for letters \rightarrow words

Trigger Word Detection

- Goal
 - trigger word: a specific predefined audio signal to invoke system (e.g. xiaodu xiaodu)
 - detect where trigger word included in an audio (if any)
- Train Set Setup
 - Basic
 - 0 for frames not corresponding to trigger word; 1 for frames consisting trigger words
 - upsampling
 - upsampling positive example: extends 1 label a few frames after the trigger words (as trigger word often appears once in an interaction with system)
- Classic Approach
 - RNN Encoder-Decoder
 - encoder scan through audio & decoder output 0-1 classification at each step
 - Conv RNN
 - a fixed window to better capture context for detecting trigger word

12.6.7 Machine Reading Comprehension**RNN with Attention****Convolution with Self-attention - QAnet****12.6.8 Image Caption****Problem Formulation**

- Input Data
 - Image
 - visual input as the target of description
- Goal
 - Natural Expression
 - description of the image in natural language, e.g. English

Baseline Approach & Previous Work

- Neural Image Caption
 - Visual Information
 - encoded by CNN backbone into a 1-D vector
 - Word Information
 - a set of word selected beforehand
 - word embedding performed
 - Language Generation
 - generated by an LSTM decoder
 - combining info: visual encoding as initial state of LSTM

- process: LSTM gives each word a to-be-selected probability at each time step
 - Inference
 - sampling: sample each word according to the distribution given by LSTM
 - beam search: iteratively consider extending k best sentence of length t to $t + 1$
 \Rightarrow select k best sentence of length $t + 1$ from all resulted sentences
- (beam search selected in the paper)

12.6.9 Referring Segmentation

Problem Formulation

- Input
 - Image
 - visual input for segmentation
 - Natural Language Expression
 - expression to denote the interested object(s)/stuff(s)
- Goal
 - Segmentation Mask of Referred Object(s)
 - currently (till early 2019), mostly binary segmentation
- Related Area
 - NLP + CV
 - referring localization
 - image caption

Baseline Approach & Previous Work

- Segmentation from Natural Language Expressions
 - Spatial Info
 - FCN-32s to encode the image into 2-D feature maps (the last conv layer)
 - Language Info
 - LSTM to encode the sentence into 1-D vector (the last hidden state)
 - Combining Info and Output
 - per-pixel info: concat [coordinates of current pixel (coord info), language info]
 - tile the per-pixel info into a feature map, then concat to the spatial info
 (per-pixel info concatenated at every pixel of spatial info)
 - followed by a series of conv and finally a deconv for upsampling
 - Training
 - per-pixel cross-entropy loss
 - Pros
 - special spatial info: coord of each pixel
 - standard info combination: concatenation
 - Cons
 - no powerful spatial info encoder: FCN-32s instead of Resnet/Unet...
 - weak upsampler, compared to encoder-decoder architecture

- language info comes late: after downsampling
- weak language info: only integrated once
- Recurrent Multimodal Interaction for Referring Image Segmentation
 - Spatial Info
 - DeepLab-101 as encoder (Resnet as backbone, with atrous conv)
 - then tiled (concat at every pixel) by coord info (coordinate of current pixel)
 - Language Info
 - word embedding w_t for $t = 1, \dots, T$
 - LSTM scanning the sentence, with hidden state h_t at time t
 - language info $l_t = \text{concat}[h_t, w_t]$
 - Combining Info
 - l_t tiled to spatial info, at each time step
 \Rightarrow creating combined feature maps F_t (of shape [height, wide, channel])
 - combined feature maps F_1, \dots, F_T fed to an convolutional LSTM, where the ConvLSTM shares weight over both space and time
 \Rightarrow feature vector of $F_t[i, j]$ is the input of the ConvLSTM at time t
 \Rightarrow conv in ConvLSTM implemented as 1×1 conv
 - a series of conv following the last hidden state of the ConvLSTM
 - Output
 - bilinear interpolated to original input size
 - optionally post-processed by dense CRF, using pydensecrf (hence inference only)
 - Pros
 - more powerful spatial info extractor: DeepLab-101
 - better language info: integrated at every time step, maintained by an ConvLSTM
 - Cons
 - weak architecture for spatial info: still no upsampling (blur segmentation)
 - no spatial relation considered in ConvLSTM (?)
 - weak language representation
 (better with pos tag, word2vec, word dict, biLSTM, and maybe even attention)
 - language info still comes late: still after downsampling

Current State-of-The-Art (early 2019)

- Key-Word-Aware Network for Referring Expression Image Segmentation
 - Spatial Info
 - DeepLab-101 as encoder for comparability
 - then tiled by coord info (coordinate of each pixel)
 - Language Info
 - LSTM scanning sentence, each hidden state as word info
 - Combining Info
 - attention mask from combined info (spatial info with language info tiled)
 (at each time step)
 - attention weighting over spatial info at each time step
 \Rightarrow an 1-D global encoding for each time step (via weighted mean over space)
 \Rightarrow filling feature maps: global encoding if attention here $>$ threshold; else $\mathbf{0}$
 \Rightarrow summing filled feature maps over time for the global spatial maps c

- attention weighting over tiled language info at each time step, correspondingly
 \Rightarrow tiled language info maps summed over time for the global language maps q
 - concat [spatial info, c , q], followed by 1×1 conv
 - Output
 - upsampling performed
 - Pros
 - attention introduced: from combined info
 - better combination: attention masked interact with both spatial & language info
 - Cons
 - blur segmentation: no encoder-decoder architecture
 - attention mask obtained sequentially: only last mask has complete language info
 - language info comes late: after downsampling
- Referring Image Segmentation via Recurrent Refinement Networks
 - Spatial Info
 - DeepLabv1 ResNet-101 as encoder
 - last feature maps tiled (concat at each pixel) with coord info
 - Language Info
 - LSTM scanning sentence, generating word info at each time step
 - last hidden layer as language info
 - Combining Info
 - combined info = spatial info tiled with language info
 - selecting set of feature maps from downsampling stages
 - all selected feature maps resized and fed to 1×1 conv
 \Rightarrow to match the dimensions of combined info
 - convolutional LSTM applied to refine the combined info
 (with matched selected feature maps as input at each time step)
 - Output
 - a conv after final hidden state of ConvLSTM for segmentation
 - upsampled to original image size
 - Pros
 - ConvLSTM integrating info at downsampling stage \Rightarrow segmentation refined
 - Cons
 - no upsampling: blur segmentation, mitigated by ConvLSTM though
 (yet no language info introduced in refinement)
 - CNN fixed during training: relying on ConvLSTM
 - single info combination: only by tiling
 (though, currently performing the best in all dataset)

12.7 Special Learning

12.7.1 Transfer Learning

Problem Formulation

- Input Data

- Source Data
 - a large amount of labeled data
 - having different distribution than the desired target data
- Target Data
 - a small amount of labeled data, with a large amount of unlabeled data (due to hardness of labeling, etc.)
 - from the distribution where model need to handle
- Goal
 - Model Performance
 - good performance on val&test set (containing target data)
 - good generalization ability on the target distribution

Standard Baseline

- Pre-training & Fine-tuning
 - Assumption
 - distribution of source & target data share some common features
⇒ different task shares some common knowledge
 - Practice
 - train network on source data only
 - swap/modify the last few layers (including prediction layer)
 - retrain the last layer (limited target data) / all net (enough target data)
 - Guideline
 - small dataset: freeze pretrained network & use it as fixed feature extractor
⇒ only train the last prediction layer
 - medium dataset: freeze fewer layers, design some own last layers
 - exceptionally large dataset with large computation budget: train from scratch
⇒ pretrained weights as initialization (nor preferred in most cases)
 - Understanding
 - sharing weights/structure: low level feature extraction useful for both
⇒ based on model ability
- Transfer Ada Boost (trAdaBoost)
 - Assumption
 - distribution of target & source overlap more or less
⇒ able to extract helpful guides from source data
 - Practice
 - setup train set with mixed target & source data
 - weighting example from target & source differently:
for source data weight = $\frac{1}{N_{\text{source}}}$; target data weight = $\frac{1}{N_{\text{target}}}$
⇒ target data more important (as smaller in number)
 - for each weight-update iteration (may contain multiple epochs), update the weight:
⇒ shift the weight (importance) towards target data & normalize all the weight
⇒ based on data distribution

- Understanding
 - learn the shared feature/knowledge with the help of source data
 - focus more on target as making progress
- Feature Projection
 - Assumption
 - few or NO overlap between source & target (as data examples directly)
 - source & target can be mapped onto a shared feature space, where overlap can be discovered
 - Practice
 - project/map the source & target data onto the same feature space
 - transfer learning in the shared space
 ⇒ based on distribution transformation
 - Understanding
 - try discover common feature through transformation
(may need a decoder to map back to desired output space)
 - Example
 - word embedding: learn word relation as unsupervised learning
⇒ a shared feature space to represent word

12.7.2 Multi-task Learning

Problem Formulation

- Input Data
 - Multi-labeled Data
 - one input data corresponds to multiple desired outputs
 - ⇒ require similarity/common knowledge in different tasks
(e.g. object detection for multiple object types)
 - Partial-labeled Data
 - desired outputs may not be all labeled in the input
(i.e. some may be missed)
- Goal
 - General Solution to Multi-task
 - give all desired outputs from a single network

Standard Baseline

- Single Networks with Multiple Predictions
 - Sharing
 - shared low-level layers to extract features from the input
 - shared loss as a sum over all prediction for corresponding label
 - shared training as back-prop computed as a single network
 - shared input data as trained together
⇒ shared knowledge discovered when training on data for other tasks
 - Understanding
 - each task help each other, by contributing to the common knowledge
 - overcome data shortage: augmented by data for other tasks
 - partial labeled still useful: help train the shared layers

12.7.3 K-shot Learning

Problem Formulation

- Training Set
 -

12.8 Generative Adversarial Nets

12.8.1 Temporal Generative Adversarial Nets

Video Generation